SoK: Reusing Defense Evidence in LLM and Agent Side Channels

***Abstract*—Side-channel defenses for LLM serving and agents are evaluated across caches, traffic traces and state-changing actions. This SoK examines what those evaluations report together and what can be reused at a new deployment. We map 82 study families and extract 33 mitigation records from 20 families, alongside six contextual records. The inspected sources show a reporting asymmetry: an inference outcome is explicitly bound in 16 of 17 applicable families, utility in 4 of 17, and matched adaptive evaluation in 2 of 11. Accepting partial evidence raises utility to 6 of 17 but leaves matched adaptation at 2 of 11. Among eight families with explicit inference and cost evidence, five place them in different records. Agent privacy and attribution studies supply partial usefulness evidence but no explicitly bound task-utility result in the four-family subset. We organize these findings by observed signal, executed intervention and evaluation setting. A record atlas preserves positive examples, including CachePrune's shared quality-and-cost package, while exposing when costs or outcomes belong to another mechanism. Comparison rules, complete query memberships and filled reporting cards translate the synthesis into candidate selection and deployment measurements. These findings describe reporting in the inspected corpus, with applicability and status sensitivities made explicit.**

1. Introduction

Encryption conceals message content, yet LLM serving and agents expose packet sizes, cache timing, destinations and tool sequences [1], [2], [3], [4], [5], [6]. Padding, cache isolation and decoy actions address different parts of this problem. Choosing among them requires both a protection result and evidence of what the protected system still does, at what cost.

What the literature reports together. The inspected mitigation literature emphasizes inference outcomes more consistently than utility or matched adaptive evaluation. Explicit bindings occur in 16/17 applicable families for inference, 4/17 for utility and 2/11 for matched adaptation. This contrast survives accepting partial evidence: 17/17, 6/17 and 2/11, respectively. Cost is explicitly available in 11/19 families, yet five of the eight families with both inference and cost place the two in different records. These are observable patterns of reporting in the extracted sources. They motivate the SoK's central question: *which evidence describes the same protected system, and what additional measurement does a deployment change require?*

One example, three decisions. Consider a serving operator choosing a cache-reuse defense. Early Bird evaluates a prefix-recovery oracle and separately measures semantic-cache deidentification overhead [3, V-A–B]. CachePrune evaluates direct and contextual recovery and links its default configuration to generation quality and serving cost [7, 5.1–5.6]. Prefix and direct-recovery results inform a shortlist for a shared token-recovery threat. Early Bird's semantic-cache overhead prices a different mechanism, whereas CachePrune's shared package informs feasibility at its tested configuration. A changed workload then identifies the next workload and cost measurements. This example connects the reporting pattern to a concrete selection decision.

The 82-family map supplies context. The defense synthesis covers 33 mitigation records from 20 families, with six contextual records locating observation and execution boundaries. We organize changes into the *observed signal*, *executed intervention* and *evaluation setting*. This taxonomy connects familiar traffic-shaping and isolation questions with agent-specific issues: a cover action can change the environment, and component output quality can differ from task success.

Contributions.

- **A synthesis of reporting asymmetries.** Source-linked family and record censuses identify where inference, utility, adaptation and cost evidence are available, including the agent task-utility gap. Status relaxation, common denominators and family-removal checks distinguish stable patterns from strict-binding-dependent findings (§5).
- **Configuration-level comparison and positive reuse.** The record atlas separates costs measured for another mechanism and retains outcomes measured together, including the CachePrune package and SparLeak's intensity-specific inference/quality pairs. Exact record identities make each combination reconstructable in this PDF (§5).
- **Deployment measurement plans.** The holdout synthesis separates attack generalization from tested defense transfer. A reusable reporting card and filled examples translate a changed signal, intervention or setting into retained evidence and a specific next measurement (§§6–8.1).

The synthesis builds on website-fingerprinting and adaptive-attack evaluation practice [8], [9], [10], [11]. It connects that practice to the measured objects in LLM/agent systems, where serving primitives, tool state and planner utility coexist. Worked applications illustrate the resulting decisions.

TABLE 1. RESEARCH QUESTIONS, SOURCE-BASED FINDINGS AND DEPLOYMENT CONSEQUENCES.

Question	Synthesis finding	Concrete use / evidential basis
RQ1: landscape	Evaluated mitigation covers 10/22 pure service-layer families, 1/8 pure hardware and 2/13 attribution families.	Target underrepresented settings within the empirical core map; preserve its C1–C5 eligibility.
RQ2: comparability	Five of eight families with explicit inference and cost place them in different records.	Keep Early Bird prefix recovery on the shortlist while obtaining a cost for that primitive.
RQ3: joint evidence	Explicit inference 16/17; utility 4/17; matched adaptation 2/11. With partial evidence: 17/17, 6/17, 2/11.	Request a task outcome for agent privacy cover; retain configuration-linked packages such as CachePrune.
RQ4: transfer	Three model and two topology holdouts concern attack generalization.	Reuse the tested attacker coverage; separately establish source-to-target defense-calibration outcomes.

Reading the findings. Counts describe bindings located in the inspected sources, under the stated applicability rules. An unreported binding is an open measurement for the proposed reuse; it does not establish that the source authors never performed the experiment.

2. Scope and Systematization Method

2.1. Research questions, eligibility and threat polarity

We organize the systematization around four questions. **RQ1 (landscape):** which hidden targets, threat polarities and auxiliary observables appear across LLM and agent systems, and where is each signal generated versus observed? **RQ2 (comparability):** when do two reported defense results measure the same target, receiver, intervention and cost object? **RQ3 (joint evidence):** which results bind protection, utility and cost at one evaluated operating point? **RQ4 (transfer):** which evidence distinguishes attack generalization, defended-condition evaluation and defense transfer across tasks, models, topologies and hardware? RQ1 supplies the field context; RQ2–RQ4 correspond to the three contributions.

Core inclusion requires five conjunctive criteria: (C1) LLM or agent execution, (C2) a hidden target property, (C3) an execution-dependent auxiliary observable rather than ordinary output semantics alone, (C4) specified observer permissions, and (C5) an empirical inference evaluation. Mixed-feature papers qualify only if they isolate an auxiliary-only C3 component; thus coding fingerprints, WildAuth, ZeroPrint, endpoint-content fingerprints and StepFinder remain adjacent boundary comparators [12], [13], [14], [15], [16]. Table 2 states the observable eligibility rules. For C3, the extraction must name an execution mechanism and an evaluated observable separate from ordinary linguistic meaning. A filter qualifies when its accept/reject behavior or differential response is the measured oracle for a hidden property; merely analyzing a reply's wording or style does not. If a mixed paper does not isolate that mechanism, its membership stays adjacent. This operational distinction concerns the measured signal, not whether the output is text.

We use four narrative groupings of objectives and observer roles, not mutually exclusive classes. An unauthorized identity inference can be both attribution and a confidentiality concern. Each record separately specifies the inference target, observer, authorization relation and protected interest: **Type I** concerns confidentiality/privacy; **Type II** concerns model, harness or identity attribution; **Type III** concerns evasion of a platform's human–agent detector; **Type IV** concerns authorized operator diagnostics. Types can overlap. WRA protects research intent (I); UI Traces attributes agents (II); AHB evades a defensive detector (III); ClawGuard diagnoses loops (IV) [6], [5], [17], [18]. These roles are mapped per core family, alongside generating and receiving layers, in Table 7.

TABLE 2. OBSERVABLE-BASED ELIGIBILITY. A HIDDEN TARGET ALONE DOES NOT MAKE A SEMANTIC ATTACK A SIDE CHANNEL.

Observed input	Decision and required distinction
Length, timing, cache response	Core when measured execution dependence supports hidden-target inference.
Own-site actions	Core attribution with site permission; no automatic access to other sites or private tool logs.
Reply states the secret	Direct disclosure; adjacent, not auxiliary-channel evidence.
Reply style/content identifies a user	Adjacent unless a separately evaluated auxiliary mechanism is established. WildAuth stays a mixed-view comparator.
Output-filter membership oracle	Core only for an isolated accept/reject or differential-response oracle tied to filter execution; semantic content alone remains adjacent.
Privileged loop/fault logs	Separate diagnostic component when C1–C5 hold; does not increase confidentiality evidence.

2.2. Corpus construction and reconciliation

Forty seeds, venue-targeted proceedings screening and one-hop citation snowballing produced an initial map of 60 families (43 core, 17 adjacent). A supplementary database and bulk keyword search returned 615 occurrences and 565 unique records; auditing its 204 exclusions rerouted four records, and the search added 15 core and seven adjacent families. The final collection has 58 core and 24 adjacent families. The completed search frames, closed on October 21, 2026, combine Semantic Scholar and native arXiv queries with venue-targeted screening of S&P, USENIX Security, NDSS, CCS and PMLR. C1–C5 screening and

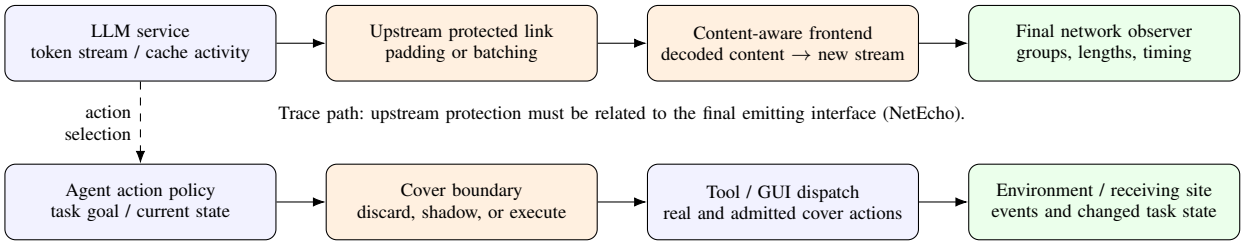


Figure 1. Receiving views and execution boundaries. A content-aware frontend can emit a new trace; admitted cover actions can change state. Evidence must describe the object a deployment will expose or execute.

version deduplication operate on records before family reconciliation; the supplement retains queries, dispositions and excluded incomplete search frames.

2.3. Source-bound result records

Each record links its target, observation boundary and operating point to a source version and section/table/figure locator. The authors report reviewing the classification of all 82 families using titles, abstracts and experimental sections, and checking all 39 result records and 180 family-audit cells against source locators. The subsequent 351-cell record audit and 133 within-question pair decisions retain author-assigned reasons; executable checks establish their accounting consistency. Candidate fields produced with AI assistance were subject to this author review. The historical material does not provide a complete item-linked draft-to-final error log, so it supplies no calibrated AI extraction-error rate. Thirteen later changes corrected the family aggregation/display while leaving record judgments fixed; they are not 13 measured extraction errors.

Appendices B and F contain the record index, comparison summary and query memberships in this PDF. The ledger also separates cited publication status, inspected version, measurement mode and remaining binding uncertainty. Published status increases traceability of venue review but does not establish the correctness of a result; a preprint's observations remain provisional at their inspected scope.

SoK Takeaway 1 (Map evidence to the receiving deployment). A family-level label does not identify which signal, observer permission or protected objective a result supports. The complete core map separates signal origin from receiving access; the result ledger then binds each comparison to a source version and operating point.

3. Result Records and Comparability Rules

The atomic unit of our systematization is one *evaluated result record*, rather than a whole paper or a single headline score. A result record binds a source version and locator to seven structured fields: Source, Target (and threat polarity Types I–IV), Receiver projection, Intervention boundary, Attacker training/calibration (D_0: clean-condition fit; D_2: protected-condition fit), Workload/holdout, and Outcomes (inference metric, task/component utility, and cost object;

TABLE 3. A SINGLE RESULT RECORD. ONLY FIELDS RELEVANT TO THE REQUESTED COMPARISON MUST ALIGN; UNAVAILABLE REQUIRED FIELDS REMAIN UNRESOLVED.

Record field	Required meaning
Source	Family/result ID, version, source locator, reading scope and hash.
Target	Hidden property, inference goal and protection relation: confidentiality, attribution, human/agent distinction or authorized diagnosis.
Receiver	Actual permission and released representation, including session/document linkage and any auxiliary inputs.
Intervention	Where protection ends; whether edits affect stored traces, transmission or executed actions.
Attacker	Fit/calibration condition, test population, training access and adaptation budget.
Workload	Task, system, topology/hardware and grouping/holdout.
Outcomes	Inference metric / denominator; objective or proxy utility; applicable cost object; uncertainty and field states.

Table 3). A record is a comparison unit with a source-indexed parameter grid, not an independent trial. We split when a change of primitive, execution mode, target/receiver contract or separately measured cost object would alter what a result can support. We retain models, datasets, parameter values and repetitions as coordinates within that unit; numerical joins must preserve those coordinates. Source versions are reconciled within a family before extraction.

Indexing schema and comparison contracts. The index names the hidden property, measured channel, observer permissions, intervention and evaluation provenance. Generating layers span hardware (P1), serving (P2), transport (P3), application/API (P4), orchestration (P5), tools/MCP (P6), environment (P7) and multi-agent communication (P8). We apply three comparison contracts (Table 4): a comparison is *compatible* when required fields align; *incompatible* when a positive mechanism or boundary conflict is established (such as substituting token-payload ratios for wire bytes or offline timestamp edits for live action delays); and *unresolved* when a required binding is unreported (U), unavailable (A) or unverified (?). Table 5 illustrates a complete record-to-judgment lineage.

TABLE 4. COMPARISON RULES. A KNOWN MISMATCH AND A MISSING BINDING HAVE DIFFERENT CONSEQUENCES.

Requirement	Condition and resulting license
Descriptive comparison	Align target, receiver permission, metric meaning and cost object. Different systems may be described together; token payload and wire bytes cannot be ranked as the same cost.
Paired effect	Compare the same defended test units and outcome. A known test-unit mismatch makes that subtraction incompatible; an unverified fit binding leaves the effect unresolved.
Deployment transfer	Establish a bridge for the changed receiver, intervention, workload, topology or hardware. A demonstrated violation blocks the proposed transfer; an untested bridge leaves it unresolved. Permit only the range and endpoints actually covered.

TABLE 5. FILLED RECORD AND ITS LICENSED JUDGMENT. VALUES ARE SOURCE-REPORTED; NO CROSS-MODE EFFECT IS COMPUTED.

Join key / field	Entry
Family / result	Dangol workflow / `DANGOL-livebinding`.
Version / locator	arXiv:2606.07150v3, §9.7, live-protocol paragraph after Table 5 [19].
Target / view	Workflow classification from live Oblivious HTTP (OHTTP)/Agent-to-Agent (A2A) traffic; not GUI-state observation.
Intervention	Live fixed-rate cover with fresh per-session identifiers, distinct from transformed stored captures.
Workload / testbed	Workflow classes with client, relay, gateway and agent co-located on one loopback host.
Outcome / cost	Accuracy 0.30; about $13\times$ bytes and $11\times$ latency in that condition.
Unresolved fields	Objective terminal task utility U; precise paired live-refit comparison ?.
Licensed judgment	These costs describe this live protocol condition. They neither price dummy GUI actions nor prove task preservation. A matched objective endpoint is needed for the latter.

Record boundaries and coverage. The extraction frame comprises the evaluated-mitigation sections and their linked methods, figures and cost/utility analyses for the 20 mitigation families, plus six explicitly named contextual controls. A defense discussion without a scoped evaluated result stays in the map. The 39-record ledger inventories these comparison units, rather than every plotted number in the 82-family corpus. S34 publishes all record scopes and split decisions. Three examples fix the granularity:

- **Split:** Early Bird's K-prefix recovery oracle and semantic-cache processing benchmark change both primitive and measured object; their numbers cannot form one protection–cost point.
- **Retain coordinates:** SparLeak's model-by-replacement-intensity ASR/ROUGE grid remains one mitigation record. Each intensity's inference/quality pair must be retained together; selecting the lowest ASR and highest ROUGE from different rows is disallowed.
- **Split outcome, share package:** CachePrune's direct recovery and contextual guessing require two target contracts. They reference one shared default quality/cost package, counted once when joining configurations.

TABLE 6. TAXONOMY OF CHANGES RELEVANT TO EVIDENCE REUSE.

Changed object	Typical transition	Evidence to carry or renew
Observed signal	Receiver, released representation, hidden target	Keep the target and permissions; measure the actual receiving trace.
Executed intervention	Prefix/semantic cache; transformed trace/live scheduler; read/write cover	Keep mechanism and state boundary; bind function and cost to the executed configuration.
Evaluation setting	Attacker fit, workload, model, hardware or network	Keep training/test provenance; establish the proposed holdout and relevant baseline.

Question cohorts precede compatibility judgments. Records enter one declared comparison question by the intervention and receiving object, not by score or verdict: delivered traffic, service-cache reuse, local memory/resource behavior, behavioral humanization, generator fingerprinting or fixed instructions. Three contextual singletons concern access gates, contention and filter-created views. S34 lists operational membership rules and all 39 assignments. When a mechanism spans several questions, retain the primary extracted question and document the alternative; a new cross-cohort question requires a new inventory. A shared cohort admits a pair for inspection; it does not make different targets or costs compatible. These retrospective rules define the comparison index. Appendix F varies record granularity and retains the original judgments. The reporting-pattern synthesis in Section 5 uses item availability and within-record bindings directly.

A taxonomy of reuse changes. The categories below describe what changes between a source result and a deployment question; several can apply together. They organize the synthesis, while the nine audit items are the evidence needed to answer the resulting questions.

From contracts to an audit. Section 5 operationalizes these contracts as nine reuse items with a shared evidence-status alphabet (Table 11) and applies them to every mitigation family (Table 12).

4. Corpus Map Across 82 Study Families

4.1. Generating layers, receiving projections and threat polarities

Table 7 maps all 58 core families by signal-generating layers, receiving projections and threat polarity. Multi-layer families retain every reported layer in their generating and observing sets.

TABLE 7. COMPLETE MAP OF THE 58 CORE STUDY FAMILIES. EACH FAMILY APPEARS ONCE; LAYER AND TYPE LABELS MAY BE MULTIPLE AND DO NOT IMPLY ONE-TO-ONE CAUSAL PATHS. THE BLOCKS GROUP SOURCE LAYERS, NOT EVIDENTIAL QUALITY.

Core family	$G \to O$	Type
hardware generation, P1 (8)		
Edge EM leakage [20]	$1 \to 1$	I
Energon [21]	$1 \to 1$	II
MoEcho [22]	$1 \to 1$	I
I Know What You Said [23]	$1 \to 1$	I
Behind Bars [24]	$1 \to 1$	I
PAGE404 [25]	$1 \to 1$	I
Spill the Beans [26]	$1 \to 1$	I
ArchSnoop [27]	$1 \to 1$	II
hardware and service, P1/P2 (8)		
SparLeak [28]	$1,2 \to 1$	I
Shape and Substance [29]	$1,2 \to 1$	I
TDXRay [30]	$1,2 \to 1$	I
SparSEEty [31]	$1,2 \to 1$	I
OTRO [32]	$1,2 \to 1$	I
Detokenization Leaks [33]	$1,2 \to 1$	I
LLM-WFIN [34]	$1,2 \to 1,3$	I
LLM-serving RCA [35]	$1,2 \to 2$	IV
service generation, P2 (13 of 22)		
SplitDecoy [36]	$2 \to 2$	I
KVGov [37]	$2 \to 2,4$	I
LLMs Have Rhythm [38]	$2 \to 3$	II
Remote timing [39]	$2 \to 3$	I
NetEcho [40]	$2 \to 3$	I
Prompt Echo [41]	$2 \to 3$	I
Speculative decoding [42]	$2 \to 3$	I
Remote keylogging [1]	$2 \to 3$	I
Whisper Leak [43]	$2 \to 3$	I
Cache contention [44]	$2 \to 4$	I
Routing metadata [45]	$2 \to 4$	I
SafeKV [46]	$2 \to 4$	I
ML privacy channels [47]	$2 \to 4$	I

Core family	$G \to O$	Type
service generation, P2 (9)		
Early Bird [3]	$2 \to 4$	I
Prompt-cache audit [48]	$2 \to 4$	I
InputSnatch [49]	$2 \to 4$	I
Leaky Language Models [50]	$2 \to 4$	I
PromptPeek [2]	$2 \to 4$	I
SpliceLeak [51]	$2 \to 4$	I
Time Will Tell [52]	$2 \to 4$	I
CachePrune [7]	$2 \to 4$	I
MoE prompt stealing [53]	$2 \to 4$	I
cross-layer sources (5)		
ClawGuard [18]	$1,5 \to 1$	IV
MIRAGE [54]	$2,5 \to 2$	IV
Cloud traffic benchmark [55]	$2,5,8 \to 3$	I, II
Web-agent unmasking [56]	$3,7 \to 3,7$	II, III
LumiMAS [57]	$4,5 \to 4,6$	IV
orchestration/workflow sources (15)		
AgentPrint [4]	$5 \to 3$	I, II
Cycle detection [58]	$5 \to 4,6$	IV
SIGLEAK [59]	$5,6 \to 4,6$	I
ProcGrep [60]	$5,6 \to 4,6$	II
AgentProv [61]	$5,6 \to 6$	II
Lidar [62]	$5,6 \to 6$	II
Dangol workflow [19]	$5,6,8 \to 3,6,8$	I
Campaign attribution [63]	$5,6,8 \to 6,8$	II
WRA [6]	$5,7 \to 3$	I
PersonaFingerprint [64]	$5,7 \to 3$	I
Whose Agent? [65]	$5,7 \to 7$	II
AgentTell [66]	$5,7 \to 7$	I
UI Traces [5]	$5,7 \to 7$	II
FP-Agent [67]	$5,7 \to 7$	II
AHB [17]	$5,7 \to 7$	III

G denotes signal-generation layers and O the receiving observation layers. Layer indices: 1 hardware; 2 serving/runtime; 3 transport; 4 application/API; 5 orchestration; 6 tools; 7 browser/environment; 8 multi-agent communication. Types: I confidentiality/privacy; II attribution; III human–agent distinction/evasion; IV authorized diagnostics. These are objectives and observer roles, not mutually exclusive security guarantees. Privileged operator receivers and physical receivers are marked by the layer they inspect. Source versions, locators and normalization of prose-only layer descriptions are in the artifact.

Cross-layer synthesis. A remote network observer, a same-site script and a privileged diagnostic hook expose different views even when they share an attribution target. WRA's read-only privacy decoys, AHB's stateful evasion and SplitDecoy's exposed return path illustrate why receiving permissions and protected objectives must accompany the generating layer [6], [17], [36].

Layer incidence highlights where the systematization's comparison problem arises. In the 58-family map, signal-generation labels cluster at P2 serving/runtime (32 families) and P5 orchestration (19); receiving views cluster at P4 application/API (18), P1 hardware (16), and P3 transport (14). Generation at P3/P4 is rare (one family each), despite frequent observation at those layers. This source/receiver asymmetry makes the transfer question concrete: protection at a frequent source layer cannot be presumed to cover a different receiving projection. The map records layer sets, so it does not assign a causal edge to each source/receiver pair.

Purpose labels further show why a single "agent side channel" category is insufficient: Type I appears 41 times, Type II 13, Type III 2, and Type IV 5. The 32 map-only families extend coverage beyond mitigation comparisons: 15 list serving origins, 11 hardware and 11 orchestration; 19 concern confidentiality, nine attribution, one evasion and all five authorized-diagnostic families. Mapping these families locates privacy targets, operator capabilities and candidate validation environments even when no scoped mitigation record enters the ledger. Counts are overlapping family-label incidences; the audit then applies its reuse rules to the 39 scoped result records.

Mitigation coverage is uneven across the map. Table 8 crosses the map with the evaluated-mitigation ledger: 10/22 pure service-layer families have a mitigation record, compared with 1/8 pure hardware and 4/15 orchestration/workflow families. The mixed hardware/service block contributes 4/8 and other cross-layer sources 1/5. Type II attribution has two mitigation families among 13 mapped

TABLE 8. CORE-MAP COVERAGE BY SOURCE-LAYER BLOCK. NUMERATORS ARE FAMILIES WITH AN EVALUATED MITIGATION RECORD; DENOMINATORS ARE THE EMPIRICAL CORE MAP, NOT ALL DEFENSE RESEARCH.

Generating-layer block	Mitigation	Core
P1 only	1	8
P1/P2	4	8
P2 only	10	22
Orchestration/workflow	4	15
Other cross-layer	1	5

Across overlapping purpose labels, Type II has 2/13 mitigation families (UI Traces and the cloud benchmark). C5 and the scoped evaluation rule define this coverage frame.

families, UI Traces and the cloud benchmark. These counts describe the empirical C1–C5 collection: a defense-only proof without a qualifying inference study can lie outside it. The contrast directs deployment work toward hardware and agent-attribution settings while preserving the distinct scopes of the existing mitigation results.

4.2. Relationship to prior systematizations

Prior surveys map compound-system threats, TEE boundaries and privacy contracts [68], [69], [70]. WF already supplies feature leakage analysis, cost trade-offs and environment-shift tests [71], [8], [72], [73]. Mathews et al. re-evaluate nine defenses with stronger attacks and an on-network DynaFlow implementation [9]. Adaptive attacks and online-cost evaluation are established foundations.

Adaptive evaluation and benchmarking. Tramèr et al. show why defense-specific adaptive attacks require more than applying an unchanged attack; Carlini et al. set out robustness-evaluation practices; RobustBench standardizes the evaluated threat and test interface [10], [11], [74]. These works already motivate protected-data evaluation and reproducible benchmarks. Our R4/R5 items instantiate that discipline. The LLM/agent-specific question is whether a refit, cost or functional result belongs to the same prefix policy, content-aware emitter or state-changing action under discussion.

What this synthesis adds. Our unit is a *scoped result*, so execution mode, protected state and cost remain attached when one paper reports several mechanisms. The 39-record atlas documents these distinct measurement objects; the 82-family holdout matrix separates evidence for each transfer axis. This connects evidence that earlier taxonomies organize by attack or mechanism: one family can contain several execution modes, a receiving projection can cross system layers, and utility/cost can belong to a different operating point from inference. The synthesis preserves these distinctions when translating a reported result into a deployment decision.

Formal guarantees and runtime measurements. The mitigation inventory covers evaluated runtime results, not the complete literature on privacy proofs. Differential privacy constrains releases for a specified adjacency relation and privacy budget [75]; information-flow measures distinguish mutual information from single-guess vulnerability [76]. Constant-time verification and oblivious memory protect stated traces under their own leakage models [77], [78]. These guarantees can supply stronger source-scoped evidence than a low empirical attack score. Their scope still matters: a guarantee for model weights need not include a serving trace, and hiding addresses need not hide access counts. TDXRay and OTRO connect these foundations to the indexed LLM components. For reuse, retain the stated adjacency, leakage model and execution assumptions. Attack-refitting items are inapplicable to a purely formal proof; absence of a classifier experiment does not count against that guarantee.

5. Reporting Patterns and Joint Defense Evidence

5.1. Evidence frame and reporting rules

The corpus partitions as $82 = 58$ core $+24$ adjacent families. Within the core, $58 = 20$ mitigation $+6$ contextual $+32$ map-only families. The ledger holds 33 mitigation records and six contextual records: SIGLEAK, AgentTell, FP-Agent, Behind Bars, LLMs Have Rhythm and ML privacy channels [59], [66], [67], [24], [38], [47]. The latter locate guardrail, access-control, interference, VPN or filter-origin boundaries and are excluded from mitigation denominators. The 32 map-only families comprise 22 without a scoped mitigation evaluation, five diagnostics, four unbound quantitative assertions and SplitDecoy's boundary case. Table 9 groups the evaluated mechanisms.

Table 11 defines the nine evidence items for a requested reuse. Component quality qualifies as utility for a component claim; an adaptation comparison requires the corresponding attacker-fit evidence. We derive each family cell from the record ledger, retaining B if any record supplies it, otherwise P, ?, U, A, then inapplicable. This availability projection keeps the record identities and source locators. Different columns can refer to different configurations, so configuration-level combinations are checked within a record.

5.2. Inference is reported more consistently than utility or matched adaptation

Table 10 puts the reporting asymmetry first. Inference is explicitly bound in 16/17 applicable families, compared with 4/17 for utility and 2/11 for a matched clean/protected-fitting contrast on defended tests. Cost is bound in 11/19. These item-wise fractions have different applicability populations. Restricting to the same 14 families eligible for inference, utility and cost gives 13, three and six explicit bindings, respectively; the cost count is twice the utility count in that common population.

TABLE 9. EIGHT INTERVENTION FAMILIES ORGANIZED BY THREAT POLARITY (TYPE I: CONFIDENTIALITY PROTECTION; TYPE II: PROVENANCE/IDENTITY ATTRIBUTION; TYPE III: BOT-DETECTION EVASION), RECEIVING VIEW AND SUPPORTED DEPLOYMENT JUDGMENT. BECAUSE OBJECTIVES AND RISK MAGNITUDES DIFFER, THESE ROWS ARE NOT A COMMON EFFICACY RANKING.

Intervention & polarity	Evaluated view, target and source	Established result and applicable boundary
Padding / batching (Type I: confidentiality)	Network packet lengths and groups; topic inference. NetEcho and Whisper Leak [40], [43].	Collected padding traces retain inference signal; packet-group batching is evaluated by replay. A protected upstream length view does not cover downstream re-emission or residual grouping. Live incremental cost remains U.
Constant-rate / transport cover (Type I: confidentiality)	Network workflow fingerprint; Dangol's OHTTP testbed [19].	The co-located loopback setting yields 0.30 recovery at about 13× bytes and 11× latency. Transformed-capture results are separate; neither result supplies objective agent-task utility.
Concurrent query decoys (Type I: confidentiality)	Destination/domain observer; private prompt and trait inference. WRA [6].	Real/decoy research prompts execute in parallel as read-only queries whose decoy answers are discarded locally; interleaved-domain evaluation reduces one-decoy E_{func} from 0.770 to 0.610. Terminal decoy utility/cost are U.
Action humanization (Type III: bot-evasion)	Mobile GUI touch traces; evading human–agent bot detection. AHB [17].	Unlike WRA's read-only privacy decoys, AHB's `Fake` mode dispatches live mobile GUI touches to evade platform bot detectors, dropping Trip Planning success from 0.75 to 0.15. Non-Fake `History+Long` (0.7143 success) avoids GUI mutation.
Cache isolation / keying (Type I: confidentiality)	Co-tenant prefix timing; SafeKV, KVGov and CachePrune [46], [37], [7].	SafeKV reports post-trigger AUC near 0.5 and separate TTFT/throughput comparisons. KVGov's HMAC median 1.3 μs prices key computation, excluding lost reuse. CachePrune jointly binds marked-token protection, QA/summarization utility and serving performance.
Cache timing equalization (Type I: confidentiality)	Co-tenant cache-fusion timing; SpliceDefense [51].	Constant-Time Boundary Fusion (CTBF) suppresses the tested signal over 15 tasks/75 steps. Quantized Chunk Padding (QCP) adds mean 7.5 tokens and $<$ 5ms/$<$ 1% prefill time; this does not price combined deployment.
Timestamp perturbation (Type II: attribution)	Recorded website action metadata; model/agent attribution. UI Traces [5].	Fitting on delayed traces (D_2) recovers substantial attribution performance under offline timestamp edits that retain action types/order. Live execution delay and task outcomes are not measured by the trace transformation.
Oblivious / unshared memory (Type I: confidentiality)	Host or local resource observations; TDXRay, Gao et al. and OTRO [30], [23], [32].	TDXRay measures oblivious-map lookup cost; Gao measures loading/memory cost after disabling zero-copy. OTRO reduces conditional access-count information (3.94 to 0.04 bits) with separate unpadded TTFT (4.59%) and epoch-padding (2.8ms) bounds.

Authorized operator diagnostics (Type IV: ClawGuard, cycle detection, LumiMAS, LLM-serving RCA and MIRAGE) are baseline telemetry rather than privacy attacks or evasion defenses, and are cataloged separately.

TABLE 10. REPORTING ASYMMETRY ACROSS THE 20 MITIGATION FAMILIES. EACH DENOMINATOR CONTAINS FAMILIES TO WHICH THAT ITEM APPLIES; B IS EXPLICIT, P PARTIAL. THE COMMON-POPULATION ROW REMOVES DENOMINATOR DIFFERENCES.

Reported binding	B / applicable	(B+P) / applicable
Inference outcome (R7)	16/17	17/17
Cost (R9)	11/19	15/19
Utility (R8)	4/17	6/17
Protected fitting (R4)	4/11	9/11
Matched adaptation (R5)	2/11	2/11
Same 14 families, inference / cost / utility:		
Explicit B	13/14 6/14 3/14	
B or P	14/14 10/14 5/14	

Expanding R4/R5 applicability to all 20 gives 4/20 and 2/20 (B), or 9/20 and 2/20 (B+P). Source judgments stay fixed. Removal checks and member identities accompany the census (main Appendix F; S38).

Sensitivity distinguishes two adaptation questions. Accepting P raises inference to 17/17, utility to 6/17 and cost to 15/19, while matched adaptation remains 2/11. Protected fitting rises sharply from 4/11 to 9/11. Thus fitting is often partially described; the larger gap is the explicit link to the protected condition and the matched adaptation contrast. In the common 14-family population, accepting P gives 14 inference, five utility and ten cost bindings. Expanding adaptation applicability to all 20 families, treating previously inapplicable entries as U only for this stress check, gives 4/20 protected fitting and 2/20 matched adaptation, or 9/20 and 2/20 with P. Removing CachePrune or all six final-search records preserves the inference–utility contrast (Appendix F).

5.3. Agent privacy evidence lacks an explicit task-utility binding

We select agent privacy/attribution families by P5 generation and Type I or II in the core map. All four members are visible: UI Traces (R8 U), WRA (P), Dangol (U), and the cloud traffic benchmark (P). WRA reports separately judged usefulness; the cloud benchmark asserts semantic preservation. Neither provides the same-condition task score required by R8 B. Hence the subset has 0/4 explicit and 2/4 explicit-or-partial utility bindings. This pattern survives both removal checks. It is a gap in *explicit task-utility binding*, not absence of all functional evidence.

The four R8 B families show what existing utility evidence can support. TDXRay binds vocabulary coverage to lookup behavior; CachePrune binds generation quality to its serving configuration; SparLeak pairs recovery and ROUGE by intensity [30], [7], [28]. AHB binds executed task accuracy to its wrapper, but its goal is detector evasion, not the privacy/attribution objective of the four-family subset

TABLE 11. CLAIM-CONDITIONED REUSE RULES. EACH ROW NAMES THE REQUESTED INFERENCE AND THE EVIDENCE NEEDED TO SUPPORT IT.

Item	Reuse query → required binding (**B**)
R1 Target	Same protected interest → explicit hidden property, inference goal and authorization relation.
R2 Receiver	Same observer protection → permissions, released view and auxiliary inputs.
R3 Intervention	Same executed mechanism → placement/mode: stored data, wire emission or actions.
R4 Protected fit	Resistance to protected-data fitting → identified protected training/calibration and testing.
R5 Matched adaptation	Effect of adaptation → clean/protected fits tested on the same defended population.
R6 Workload	Changed workload/system → tasks, configuration and the relevant holdout axis.
R7 Inference	Quantified protection → metric, denominator and scoped defended evaluation condition.
R8 Utility	Preserved task/component function → performance at the same protected state/operating point.
R9 Cost	Resource-budget feasibility → units, baseline and cost object at that operating point.

B: bound; **P**: partial binding; U: unreported in inspected scope; A: source unavailable; ?: binding unverified; –: inapplicable. First determine applicability, then assign evidence status. Family **B** means at least one scoped record satisfies the item, not that every result does.

TABLE 12. RECORD-DERIVED EVIDENCE AVAILABILITY FOR ALL 20 EVALUATED-MITIGATION FAMILIES. R1–R9 ARE SCOPED REUSE ITEMS, NOT OVERALL PAPER SCORES.

Family	R1	R2	R3	R4	R5	R6	R7	R8	R9
UI Traces [5]	**B**	**P**	**B**	**B**	**B**	**B**	**B**	U	–
WRA [6]	**B**	**P**	**B**	–	–	**P**	**B**	**P**	U
NetEcho [40]	**B**	**B**	**B**	**B**	?	**B**	**B**	U	**P**
Early Bird [3]	**B**	**B**	**B**	**B**	?	**B**	**B**	U	**B**
Whisper Leak [43]	**B**	**B**	**B**	**P**	U	**B**	**B**	U	**P**
SpliceLeak [51]	**B**	**B**	**B**	–	–	**B**	**B**	U	**B**
Time Will Tell [52]	**B**	**B**	**B**	**P**	U	**B**	**B**	U	**B**
I Know What You Said [23]	**B**	**B**	**B**	–	–	**P**	–	–	**B**
Speculative decoding [42]	**B**	**B**	**P**	**P**	U	**B**	**B**	–	**B**
Dangol workflow [19]	**B**	**B**	**B**	**B**	**B**	**B**	**B**	U	**B**
Cloud traffic benchmark [55]	**B**	**P**	**B**	U	U	**P**	**B**	**P**	U
Remote timing [39]	**B**	**B**	**P**	–	–	**B**	–	U	**B**
TDXRay [30]	**B**	**B**	**B**	–	–	**B**	–	**B**	**B**
SafeKV [46]	**B**	**B**	**B**	–	–	**B**	**B**	U	**P**
AHB [17]	**B**	**B**	**B**	**P**	U	**B**	**B**	**B**	U
KVGov [37]	**B**	**B**	**B**	–	–	**B**	**B**	U	**B**
CachePrune [7]	**B**	**B**	**B**	–	–	**B**	**B**	**B**	**B**
OTRO [32]	**B**	**B**	**B**	–	–	**B**	**B**	–	**B**
SparLeak [28]	**B**	**B**	**B**	**P**	U	**B**	**B**	**B**	U
Shape and Substance [29]	**B**	**B**	**P**	U	U	**B**	**P**	U	**P**
Bound	20	17	17	4	2	17	16	4	11
Applicable	20	20	20	11	11	20	17	17	19

B: bound; **P**: partial; U: unreported in inspected scope; ?: unverified binding; –: inapplicable. Table 11 defines R1–R9. Each cell is generated from records: B if any is B, otherwise P, ?, U, A, then –. Different columns can refer to different configurations. Table 25 lists every inference/cost projection and its record identities. The 13 prior family-display corrections changed aggregation, not the underlying record judgments (S36).

[17]. This distinction identifies the task metric needed when a privacy cover method is moved into an acting planner.

5.4. Family-level inference and cost often describe different records

Eight families have explicit inference and explicit cost somewhere in their evaluations. In five—Early Bird, SpliceLeak, Time Will Tell, KVGov and OTRO—no individual record carries both. The distinctions are concrete: prefix policy versus semantic processing; background-load mitigation versus QCP; a length instruction versus a padding estimate; partial flushing versus HMAC; and padded versus unpadded ORAM. SPEC, Dangol and CachePrune supply the five same-record inference/cost results listed in Table 13. Appendix F lists all eight families' exact supporting records.

This separation remains five of seven families after removing CachePrune and four of six after removing the final-search additions. Accepting partial cost and inference evidence reduces it to three of 12; a partial cost estimate can reconnect an operating point without supplying the fully bound cost. The strict majority is therefore specific to *explicit* bindings. It does not justify calling all remaining cost evidence a microbenchmark or an end-to-end measurement.

Positive combinations for deployment. The joint queries start from 33 mitigation records and exclude a record only if a required item is inapplicable. Table 13 names every positive result; Appendix F lists all query members and excluded entries. CachePrune's two target outcomes share one quality/cost package. SparLeak provides intensity-specific inference/quality pairs. These positive examples support feasibility analysis at their measured configurations; other combinations identify the exact measurement to request.

5.5. Execution evidence is distributed across measurement modes

Among the 33 mitigation records, 14 are executed collections/prototypes, six trace transformations, five simulations/derived estimates, six component/calibration measurements and two unresolved placements. The frequently cited 20/39 executed total includes all six contextual controls. Restricting instead to network receivers (P3 observation) yields 12 mitigation records: six executed, four transformed and two with unresolved placement. The data therefore support a mixture of execution and transformation, not a predominance of network replay. Excluding final-search records leaves 11/27 executed mitigation records. Even an executed prototype supplies only its measured local scope; the six component/calibration records include a hardware timing baseline as well as component-cost measurements.

Comparison permissions as a supporting index. The declared question groups generate 133 comparisons, of which 17 retain a descriptive comparison. Early Bird's prefix oracle and CachePrune's direct-recovery result can inform the same token-recovery shortlist while keeping their costs separate. Record splitting and question grouping determine this inventory; it guides a specific comparison rather than

TABLE 13. POSITIVE JOINT RESULTS AND EXACT RECORD IDENTITIES. ALL ROWS START FROM 33 MITIGATION RECORDS; A REQUIRED INAPPLICABLE ITEM EXCLUDES A RECORD. EVERY OUTCOME RETAINS THE SOURCE LOCATOR IN APPENDIX B.

Requested evidence	Bound record IDs	Interpretation at the source configuration
Inference + utility (4/22)	AHB-online-executed; CACHEPRUNE-selective; CACHEPRUNE-context; SPARLEAK-replacement	Executed task outcomes or component quality; the two CachePrune targets share one package.
Inference + cost (5/22)	SPEC-padding-REST; DANGOL-capture; DANGOL-livebinding; both CACHEPRUNE records	Payload overhead, transformed-capture cost, loopback cost and serving cost remain distinct.
Adaptation + inference (2/17)	UI-delay-offline; DANGOL-capture	Matched clean/protected fitting on a defended test; no live-utility inference.
Inference + utility + cost (2/20)	CACHEPRUNE-selective; CACHEPRUNE-context	One family, one default configuration, two target outcomes.
All five items (0/14)	None in this extraction	A specified query has no bound record; no estimate of all published or unpublished evidence.

TABLE 14. THREE EVIDENCE OBJECTS FOR RQ4. THE SAME FAMILY MAY SUPPLY SEVERAL OBJECTS IN SEPARATE EXPERIMENTS.

Object	Required link	Permitted reuse
Attack generalization	Training/profiling source → held-out inputs or systems; attack outcome.	An attacker or recognition method covers the stated target domain.
Defended condition	Named mechanism, local fitting, defended test and outcome.	Retain that operating point; multi-model testing alone is not held-out transfer.
Defense transfer	Source-selected policy/calibration held fixed → target; local-calibration comparator and attacker fit specified.	Transport the evaluated policy only over the tested direction and scope.

measuring a field-wide defect rate. Complete decisions and coarser-granularity projections are retained in Appendix F. The reporting findings above are recomputed from cells and records, independently of pair counts.

SoK Takeaway 2 (Bind function and cost to protection). Inference reporting is widespread in the inspected mitigation sources; explicit utility and matched adaptation are much less common. Check the measured operating point before combining protection and cost, and request a task outcome when a privacy intervention changes agent execution.

6. Which Evidence Supports Defense Transfer?

6.1. Separate the evaluated object before counting holdouts

RQ4 distinguishes three questions: does an *attack* recognize new inputs or systems; is a *defense condition* evaluated; and does a *fixed attacker or defense calibration* carry from a specified source to a target? Table 14 makes the distinction operational. Repeating a defense on several models can establish condition-specific outcomes without establishing a held-out model or transferring a fitted calibration.

The 82-family matrix in Appendix D is a discovery index for these questions. Its six axes are task/input/user, model, harness/topology, hardware, network/load and time. H marks reported separation, V condition variation/recalibration, X a scoped restriction, and ? unresolved extraction. The 13 task/input H entries include only four mitigation families; all three model H entries concern attack/identification generalization in families outside the 20-family mitigation audit. The two topology H entries occur in mitigation families, but their held-out objects are workflow recognizers. The five model/topology H entries thus locate attack-generalization evidence, with no source-frozen defense-calibration transfer in those entries. The 13 task entries require their individual defended-condition linkage; they are not uniformly attack-only. S34 lists all H objects, while S38 retains their member identities. This evidence-object reading is stable to treating condition variation as additional coverage: varying or locally refitting a defense still differs from transporting a source calibration.

6.2. A changed generator and a tested defense are different bindings

LLM-Fin, ArchSnoop and LeakyLMs evaluate recognition of unseen model variants or architectures [79], [27], [50]. This supplies attack-side model generalization. In the traffic benchmark, graph-to-star and star-to-graph classification explicitly move the fitted attacker across topologies; the prompt/temperature countermeasures are separate evaluations [55, V-B–C]. Dangol's leave-composition-out test likewise examines recognition of a novel workflow shape, separately from its protected-capture frontier [19, 9.8–9.12]. The reusable result in both topology entries is the recognizer's tested generalization.

There is nevertheless useful *defense-side variation* evidence. Time Will Tell applies the same fixed-length instruction to two models and reports different residual classification ASRs; this tests a common rule under changed generators [52, VII-B]. Whisper Leak evaluates mitigation conditions across provider-models, while its prompt holdout describes the classifier protocol [43, 3.2,6]. Preserve those per-condition results; a shared rule, a refitted local attack, a source-trained attack and a source-calibrated defense are four distinct comparisons.

Deployment decision. For the proposed change, name the source, target, held-out unit, quantities held fixed

or refitted, and protected outcome. Carry forward only the corresponding binding. An attack-generalization result justifies including that attacker at the target; a defended-condition result supplies its local outcome; a source-to-target defense-calibration result supports only the stated transfer. This prevents both overclaiming portability and discarding useful local evidence. Unresolved hardware/time entries in the discovery matrix identify source follow-up, not absent experiments.

7. Synthesis: From Evidence to Reuse Decisions

The census (§5) covers all 39 records and the holdout matrix (§6) all 82 families. We now select three examples by the *comparison decision* they expose, not by effect size: a changed receiving path, cover with different state effects, and an adaptation/cost mismatch within a named mechanism. Their value is the decision they change: retain an aligned result, separate a wrongly joined package, or require a result at a changed execution boundary.

7.1. Case 1: bind the final target-bearing emitter

NetEcho's upstream packet-size shaping operates on a different projection from the tool-call sequence. Its plaintext-aware frontend then emits a new size trace (grouped token counts or character lengths), reopening a size channel after upstream protection [40]. SplitDecoy is a separate return-path record: the cloud sees nonzero gradient row support on every tested frame across nine seeds [36, Sections 5.1–5.3]. Both change the receiver boundary and therefore require a separately bound inference result.

The data processing inequality requires $S \to Z \to Y$: a downstream process that receives only Z cannot increase $I(S;Z)$ [80, Theorem 2.8.1]. A frontend also receiving content X instead forms $Y = f(Z, X)$. The applicable bound is $I(S;Y) \leq I(S;Z,X)$; it need not be bounded by $I(S;Z)$. For a constructed binary example, let S be uniform, the protected upstream length Z constant, and the downstream length $Y = 100 + S$ bytes. Then $I(S;Z) = 0$ but $I(S;Y) = 1$ bit and optimal downstream classification is exact. This calculation illustrates the missing conditional-independence premise; it is not a measured leakage estimate for NetEcho.

7.2. Case 2: resolve the WRA versus AHB cover asymmetry

WRA's one-decoy E_{func} drops from 0.770 to 0.610, whereas AHB's Trip Planning success falls from 0.75 to 0.15 under `Online+Fake` humanization [6], [17]. Grouping both as "cover traffic" conceals different objectives and executions.

WRA discards decoy reads for Type I confidentiality; AHB executes GUI events for Type III evasion, potentially navigating or submitting forms. Their utility outcomes are not contradictory, and AHB's missing condition-specific denominators preclude a matched causal effect. A shared workflow setting does not establish shared state effects. Terminal-state reward and dispatch-time authorization also answer different questions: an unconfirmed write can still reach the intended final state. Reuse requires an explicit state/authorization outcome at the intervention being carried forward.

7.3. Case 3: bind adaptive retraining and cost objects

A matched adaptation contrast compares clean-trained D_0 and defense-trained D_2 on identical defended tests. UI Traces and Dangol raise the record-binding question this leaves open: whether a refit, a measured delay and a utility outcome refer to the same executed condition [5], [19]. Cost objects show the same pattern. Early Bird's 6ms/4% GPTCache cost does not price K-prefix isolation; KV-Gov's $1.3\,\mu$s HMAC microcost omits lost reuse; SpliceDefense's QCP cost does not price CTBF+QCP [3], [37], [51]. CachePrune binds protection, QA/summarization utility and throughput, while OTRO separates conditional leakage from TTFT/padding costs [7], [32]. The reuse decision follows the evaluated configuration: retain each source's bound outcome, and request the missing adaptation, execution or cost measurement before extending it.

7.4. A positive reuse path: from comparison to a candidate configuration

The Early Bird K-prefix record and CachePrune's direct-recovery record form one of the 17 descriptive matches: both address sensitive-token recovery through cache-reuse probes [3, V-A], [7, 5.1–5.6]. This permits an analyst to retain both mechanisms as candidates for that threat, while keeping the oracle and executed-serving measurements separate. The joint audit then changes the next step. CachePrune supplies a default configuration package linking recovery, generation quality and serving cost; the analyst can use that package for a source-scoped feasibility assessment. Early Bird's semantic-cache overhead cannot price its K-prefix candidate, so that candidate still needs a serving-cost measurement. Neither a pooled recovery ranking nor a new-hardware guarantee follows from this comparison.

What the three findings add. Record-level comparison turns a mechanism-level shortlist into correctly scoped alternatives. Configuration-level joint binding distinguishes a reusable package from borrowed utility or cost. Deployment decisions then identify the changed object and the measurement it needs. Compared with established WF practice on adaptive attacks, shifts and online costs [8], [9], the LLM/agent synthesis makes three additional joins explicit: prefix versus semantic-cache primitives, discarded read outputs versus executed state changes, and component generation quality versus planner outcomes. The source-backed examples identify which measurements a particular

TABLE 15. CLAIM-CONDITIONED REUSE DECISIONS DERIVED FROM THE RECORD AND HOLDOUT AUDITS. EACH CHANGED BOUNDARY IDENTIFIES THE EVIDENCE THAT MUST BE BOUND BEFORE CARRYING A RESULT FORWARD.

Reuse question	Evidence decision	Retain or measure next
Emitter / representation	An upstream or buffered capture can differ from the final receiving trace and delivery schedule.	Bind inference to the final receiver; retain the emitter's overflow, byte and delivery-delay scope (R2/R3/R9; §7).
Read cover → writes	Discarding a reply does not establish unchanged commits, authorization or workflow outcomes.	Bind the protected state and utility to the executed intervention; include cover calls, quotas and failures (R3/R8/R9).
Customer / model calibration	A reported split can separate inputs without changing the generator or calibration population.	State held-out units and calibration provenance; separate local fitting variability from source/domain changes; evaluate the requested transfer axis (R5/R6; §6).
Attacker observation	An outcome for one projection or fitted attacker does not supply a result for a richer receiving view.	Evaluate comparable defended tests under the intended observation and adaptation contract; retain receiver-specific inference metrics (R2/R5/R7).
Stored trace → online path	A trace transformation or simulated span does not measure executed queues, completion or online latency.	Bind the receiver, harness and resource denominator to the online condition; establish the relevant topology/network holdout (R3/R6/R9).
Bound cache-reuse package	Early Bird prefix recovery and CachePrune direct recovery permit a scoped candidate comparison; CachePrune binds a shared utility/cost package.	Retain that package for its measured configuration; measure prefix-policy serving cost separately, and re-evaluate any changed workload/hardware.
Component → workflow utility	Component quality or logical state conformance does not determine a planner's response to scheduling and deadlines.	Report utility at the requested workflow boundary and the same intervention/configuration as protection and cost (R3/R8/R9).

deployment query may combine and which it must keep separate.

7.5. A deployment query determines the evidence to retain

Start with the intended receiver, protected outcome and intervention. Record-level comparison identifies whether the cited result measures that object. Configuration-level joint binding identifies which inference, utility and cost outcomes share its configuration. Deployment decisions use the holdout matrix to identify evidence for the proposed change and specify the remaining measurement. Table 15 makes these outputs operational: a new emitter needs a receiver-specific capture, writes need state/authorization checks, source-calibration comparisons need a local fitting-variability control, and an online scheduler needs its own planner utility and latency measurement. The calibration control follows from the comparison design: changing the source can also change the finite training sample used to fit the defense.

The following worked applications show how the procedure represents external reevaluations and internal examples. Appendix E retains the detailed controlled case study.

7.6. External worked applications: retain, separate, remeasure

We apply the procedure to a fixed external frame: both defenses in Deep Fingerprinting, all three in WFDefProxy, and all nine in Mathews et al.'s reevaluation, with DynaFlow's live implementation kept as a separate record [81], [82], [9]. The resulting 15 measurement records represent 14 defense–study pairs in three studies. The ledger retains favorable outcomes alongside weaknesses, fixes the Standard setting for the nine-defense table, and records a locator and a concrete measurement decision for every row (S35).

The worked output separates decisions that a mechanism name conflates (Table 16). Stronger protected-data fitting requires revisiting WTF-PAD's attack result, while Walkie-Talkie's symmetric-pairing result remains a valid scoped outcome. FRONT's live protection improves even as its bandwidth cost changes. A record-level audit therefore preserves the security observation and replaces the cost estimate separately. DynaFlow's live result needs a separate execution/workload record, preventing a simulated cost from being attached to a live attack result.

The complete frame produces 11 protected-fit/view measurement requirements and four execution/cost requirements. Four favorable controls remain usable in their measured scopes: symmetric Walkie-Talkie, live FRONT, live Tamaraw and one-split TrafficSliver. The actionable output is a scoped retain, separate or remeasure decision for each record. The examples faithfully encode established WF practice. Outcomes were visible during extraction, so agreement is an illustration of the representation rather than evidence of prediction accuracy or an improvement over ordinary expert reading. These three WF studies do not establish LLM/agent effectiveness; the cache and tool examples provide the domain-specific application.

7.7. Seven internal illustrative judgments

The existing controlled study supplies illustrative applications: seven written judgments specify which apparent protection claims should survive a changed target, receiver or attacker. The protocol-only repository state predates the

TABLE 16. EXTERNAL APPLICATIONS: RETAIN A SCOPED RESULT, SEPARATE AN OUTCOME PACKAGE, OR REQUEST A SPECIFIC MEASUREMENT. FULL 15-RECORD LEDGER: S35. ACCURACY AND TPR RETAIN THEIR ORIGINAL POPULATIONS.

Published contrast	Evidence-reuse decision
WTF-PAD [81]	Refit: DF accuracy 90.7%; prior attackers in the same table reach at most 69.0%.
Walkie-Talkie [81]	Retain symmetric-pairing scope: DF accuracy 49.7%; do not inherit it for asymmetric pairing.
FRONT [82]	Separate outcomes: live DF TPR improves, 72.06% to 42.79%, while data overhead rises, 52% to 75%.
Tamaraw [82]	Retain low live TPR, 6.43%; use live bandwidth overhead 179%, rather than simulated 207/224%.
DynaFlow [9]	Separate live record: 379% mean per-site time overhead; changed URLs/crawler prevent a matched simulation–live effect.

TABLE 17. CURRENT OUTCOMES OF ALL SEVEN TIME-ORDERED INTERNAL JUDGMENTS. SUPPORTED REFERS TO THE SPECIFIED COUNTEREXAMPLE, NOT A FRAMEWORK-ACCURACY ESTIMATE.

ID	Reuse claim tested	Finding and next measurement
P1	Buffered protection/cost carries to a stream	Supported: verbosity accuracy 100%; capture the stream emitter.
P2	Low clean-trained accuracy establishes protection	Supported: D_0 50%, D_2 100%; refit on protected data.
P3	Padded events hide response identity	Supported: top-1 58–62%, chance 2.1%; test the whole trace.
P4	Equal target-verb count hides cancellation	Supported: count BA 50%, suite BA 76.9%; expand the view.
P5	Protected dispatch implies a protected tool link	Source premise unmet; excluded from favorable evidence. Establish dispatch protection first.
P6	Framing alone hides intent	Partial: bucketed return leaks; resolve the remaining target cells.
P7	Joint cover/framing protects both targets	Not established: retain the joint screen before accepting protection.

first archived result commit on 21 October 2026; S35 preserves the ordering receipt, original labels and current reassessment. Table 17 brings all seven into the main paper.

Four judgments have supporting counterexamples; one is partly supported. P5 has a failed source premise and is excluded from favorable evidence for either transfer failure or framework success. P7's proposed joint protection remains unestablished. The examples specify the receiving view and protected-data fit used in each calculation. Repository ordering establishes chronology only; author-designed examples do not measure the accuracy of the procedure. The decisions differ by failure mode: P1–P4 require the changed emitter, attacker or trace view to be measured; P5 requires establishing source protection first; P6/P7 require evidence for the still-open claim. Appendix E retains the measurement scopes and statistical qualifications.

TABLE 18. REUSABLE REPORTING CARD. FILL EACH EVIDENCE CELL WITH A VERSION AND LOCATOR OR "UNKNOWN"; THE FINAL ACTION IS CONDITIONAL ON THE PROPOSED DEPLOYMENT.

Field	What to enter
Receiver/target	Hidden property, accessible signal and permissions.
Intervention	Primitive, execution mode, state and trust boundary.
Attacker/test	Training or calibration condition; defended test population.
Function	Component quality or task outcome; shared configuration.
Cost	Resource, units, denominator, baseline and included operations.
Extension	Changed workload/system axis; available holdout; retain, separate or measure next.

8. Implications and Ethics

8.1. Deployment and reporting implications

The decision table distinguishes receiving permissions from intervention capabilities. An access-network operator, VPC monitor and receiving tool service expose different views; remote paths can reveal destinations even with Encrypted ClientHello [83]. A result for one ingress projection therefore needs a separately bound claim before it is applied to host instrumentation or downstream commit logs.

Likewise, read queries, provider-isolated execution and gateway-routed cover have different state contracts. Kubernetes dry-run, Stripe sandboxes and transactional rollback offer isolation mechanisms with their own limits [84], [85], [86]; idempotence alone is insufficient, and sequence changes can survive rollback [87]. Evidence reuse requires the trusted boundary, authorization policy, resource quotas and protected outcome to accompany the mechanism name. The corpus's component outcomes remain useful when their execution and cost scopes match the intended query.

A usable reporting card. Table 18 is a reusable six-field template. Its last column asks for a deployment decision after recording the source evidence. Table 19 fills it for two venue-linked sources and the running cache example, preserving their different scopes. The editable artifact template includes an explicit unknown state and source-locator slot for every field.

8.2. Ethics and artifact scope

The audit uses source-linked public research records. Stakeholders include tenants, shared-capacity providers and users affected by observable behavior or state changes. Bot-detection evasion and research-intent privacy have different authorization relations [17], [6]. The appendix case study uses cloned public benchmark databases and controlled sockets; model interactions use public tasks and simulated users. No production account is modified.

AI tools assisted candidate-field extraction, implementation, source reconciliation and prose revision. The author-review coverage is stated in §2. Reliability controls retain

TABLE 19. FILLED REPORTING CARDS FROM SOURCE RESULTS. EARLY BIRD IS VENUE-LINKED (INSPECTED AUTHOR VERSION); TDXRAY IS AN S&P PAPER; CACHEPRUNE IS A CITED PREPRINT AND SERVES AS THE RUNNING CONFIGURATION EXAMPLE. NO NEW BENCHMARK IS IMPLIED.

Field	Early Bird [3, V-A–B]	TDXRay [30, 6.1, Fig. 8, Table 3]	CachePrune [7, 5.1–5.6]
Receiver/target	Timing oracle; sensitive prefix recovery.	Memory-access view; oblivious lookup defense component.	Cache-reuse probes; direct and contextual token recovery.
Intervention	Prefix-granularity oracle; semantic-cache prototype separately.	Oblivious mapping microbenchmark.	Selective cache sharing at default configuration.
Attacker/test	Granularity-specific calibration; matched adaptation unresolved.	Component cost, no bound protected end-to-end inference row.	Recovery evaluation; default quality/cost package; no matched adaptation contrast.
Function	No bound prefix-policy utility.	No bound full-inference utility.	QA/summarization quality at stated configuration.
Cost	Semantic-cache overhead belongs to a different primitive.	Lookup microcost; not complete serving cost.	Serving throughput/cost; preserve default package.
Extension/action	Retain prefix candidate; measure its own serving cost.	Retain component cost; measure protected service and its function.	Retain scoped package; test changed workload/hardware.

source versions and locators, separate unavailable from unreported evidence, preserve corrections, and recompute aggregation and joins from the record cells before compilation. These controls detect inconsistent accounting and support source inspection; they do not estimate an AI extraction-error rate. Authors retain responsibility for source interpretation. Case-study observations and protocols remain separate from the literature atlas.

9. Conclusion

The inspected LLM/agent mitigation sources report inference outcomes more consistently than utility or matched adaptive evaluation. This pattern survives accepting partial bindings and restricting outcomes to a common family population. Agent privacy/attribution studies provide partial usefulness evidence but no explicit task-utility binding in the four-family subset. Five of eight families with explicit inference and cost place the two in different records; the operating point determines whether they can be combined.

The SoK connects these reporting patterns to deployment decisions through scoped result records. Positive packages such as CachePrune and SparLeak show what can be retained, while the holdout synthesis identifies whether a changed system has attack-side or defense-side evidence. The released record identities, sensitivity checks and filled reporting cards make the findings inspectable and turn the remaining bindings into concrete measurement requirements.

References

[1] R. Weiss, D. Ayzenshteyn, G. Amit, and Y. Mirsky, "What was your prompt? a remote keylogging attack on AI assistants," in *33rd USENIX Security Symposium*. USENIX Association, 2024, pp. 3367–3384. [Online]. Available: https://www.usenix.org/conference/usenixsecurity24/presentation/weiss

[2] G. Wu, Z. Zhang, Y. Zhang, W. Wang, J. Niu, Y. Wu, and Y. Zhang, "I know what you asked: Prompt leakage via KV-cache sharing in multi-tenant LLM serving," in *Network and Distributed System Security Symposium (NDSS)*, 2025. [Online]. Available: https://www.ndss-symposium.org/ndss-paper/i-know-what-you-asked-prompt-leakage-via-kv-cache-sharing-in-multi-tenant-llm-serving/

[3] L. Song, Z. Pang, W. Wang, Z. Wang, X. Wang, H. Chen, W. Song, Y. Jin, D. Meng, and R. Hou, "The early bird catches the leak: Unveiling timing side channels in LLM serving systems," *IEEE Transactions on Information Forensics and Security*, vol. 20, pp. 11 431–11 446, 2025, consulted arXiv:2409.20002v5; archived access 18 October 2026. [Online]. Available: https://ieeexplore.ieee.org/document/11206380/

[4] Y. Zhang, X. Deng, Z. Gu, Y. Chen, K. Xu, Q. Li, and J. Wu, "Exposing LLM user privacy via traffic fingerprint analysis: A study of privacy risks in LLM agent interactions," *arXiv preprint arXiv:2510.07176*, 2025. [Online]. Available: https://arxiv.org/abs/2510.07176

[5] W. Lugoloobi, S. Marro, J. Magomere, J. Wright, and C. Russell, "Known by their actions: Fingerprinting LLM browser agents via UI traces," *arXiv preprint arXiv:2605.14786*, 2026, consulted arXiv:2605.14786v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2605.14786v1

[6] H. Jeong, M. Teymoorianfard, A. Kumar, A. Houmansadr, and E. Bagdasarian, "Network-level prompt and trait leakage in local research agents," in *35th USENIX Security Symposium*. USENIX Association, 2026, pp. 6485–6504. [Online]. Available: https://www.usenix.org/conference/usenixsecurity26/presentation/jeong

[7] G. Wu, Z. Li, Y. Zhang, Z. Zhang, J. Niu, Y. Wu, and Y. Zhang, "Cacheprune: Privacy-aware and fine-grained KV cache sharing for efficient LLM inference," *arXiv preprint arXiv:2605.23640v1*, 2026, consulted 2605.23640v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2605.23640v1

[8] M. Juarez, S. Afroz, G. Acar, C. Diaz, and R. Greenstadt, "A critical evaluation of website fingerprinting attacks," in *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, 2014, pp. 263–274.

[9] N. Mathews, J. K. Holland, S. E. Oh, M. S. Rahman, N. Hopper, and M. Wright, "SoK: A critical evaluation of efficient website fingerprinting defenses," in *IEEE Symposium on Security and Privacy*, 2023.

[10] F. Tramèr, N. Carlini, W. Brendel, and A. Madry, "On adaptive attacks to adversarial example defenses," in *Advances in Neural Information Processing Systems*, vol. 33, 2020. [Online]. Available: https://arxiv.org/abs/2002.08347v2

[11] N. Carlini, A. Athalye, N. Papernot, W. Brendel, J. Rauber, D. Tsipras, I. Goodfellow, A. Madry, and A. Kurakin, "On evaluating adversarial robustness," *arXiv preprint arXiv:1902.06705*, 2019. [Online]. Available: https://arxiv.org/abs/1902.06705v2

[12] T. A. Ghaleb, "Fingerprinting AI coding agents on GitHub," *arXiv preprint arXiv:2601.17406v1*, 2026, consulted arXiv:2601.17406v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2601.17406v1

[13] W. Wang, Z. Liu, H. Kou, X. Liu, and J. Shen, "The chatbot knows it's you: Dialogue attribution in unauthenticated human–LLM sessions," in *Proceedings of the ACM Web Conference 2026*, 2026, pp. 8797–8807. [Online]. Available: https://dl.acm.org/doi/10.1145/3774904.3793048

[14] S. Shao, Y. Li, H. Yao, Y. Chen, Y. Yang, and Z. Qin, "Reading between the lines: Towards reliable black-box LLM fingerprinting via zeroth-order gradient estimation," 2026, version 2, 21 January 2026. [Online]. Available: https://arxiv.org/html/2510.06605v2

[15] J. Leshin, M. Shah, I. Timmis, and D. Kang, "Behavioral fingerprints for LLM endpoint stability and identity," 2026, version 1, 19 March 2026. [Online]. Available: https://arxiv.org/html/2603.19022v1

[16] T. Zhu, Y. Wu, W. Jin, Y. Li, and G. Huang, "StepFinder: A temporal semantic framework for failure attribution in multi-agent systems," 2026, version 1, 2 June 2026. [Online]. Available: https://arxiv.org/html/2606.03467v1

[17] J. Zhu, L. Yang, R. Shan, C. Zheng, Z. Zheng, W. Liu, Y. Yu, W. Zhang, and J. Lin, "Turing test on screen: A benchmark for mobile GUI agent humanization," *arXiv preprint arXiv:2604.09574v1*, 2026, consulted arXiv:2604.09574v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/html/2604.09574v1

[18] L. L. Gan, J. Wu, L. Ge, L. Yang, Y. Song, J. Zhang, H. Jin, W. Wang, and G. Xue, "ClawGuard: Out-of-band detection of LLM agent workflow hijacking via EM side channel," *arXiv preprint arXiv:2605.06205v1*, 2026. [Online]. Available: https://arxiv.org/abs/2605.06205v1

[19] B. Dangol, "From privacy to workflow integrity: Communication-graph metadata in autonomous agent interoperability," *arXiv preprint arXiv:2606.07150v3*, 2026, consulted arXiv:2606.07150v3; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2606.07150v3

[20] N. Alhalawani, G. Acar, Z. Liu, and L. Batina, "Physical side-channel attacks on edge LLMs: Extracting output tokens and user prompts," *arXiv preprint arXiv:2610.14215v1*, 2026, consulted 2610.14215v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.14215v1

[21] A. Chaudhuri, S. Shukla, S. Bhattacharya, and D. Mukhopadhyay, "Energon: Unveiling transformers from GPU power and thermal side-channels," in *IEEE/ACM International Conference on Computer-Aided Design*, 2025. [Online]. Available: https://arxiv.org/abs/2508.01768v1

[22] R. Ding, T. Xu, X. Shen, A. A. Ding, and Y. Fei, "MoEcho: Exploiting side-channel attacks to compromise user privacy in mixture-of-experts LLMs," in *ACM SIGSAC Conference on Computer and Communications Security*, 2025, pp. 2159–2173. [Online]. Available: https://arxiv.org/abs/2508.15036v1

[23] Z. Gao, J. Hu, F. Guo, Y. Zhang, Y. Han, S. Liu, H. Li, and Z. Lv, "I know what you said: Unveiling hardware cache side-channels in local large language model inference," in *34th USENIX Security Symposium*, 2025, pp. 1649–1668. [Online]. Available: https://www.usenix.org/conference/usenixsecurity25/presentation/gao-zibo

[24] C. Gu, R. Levine, Z. Zhang, T. Sorensen, and Y. Guo, "Behind bars: A side-channel attack on NVIDIA MIG cache partitioning using memory barriers," in *35th USENIX Security Symposium*. USENIX Association, 2026, pp. 5373–5392. [Online]. Available: https://www.usenix.org/conference/usenixsecurity26/presentation/gu-cheng

[25] R. Nazaraliyev, Y. Zhang, and N. Abu-Ghazaleh, "Page404: Breaking GPU isolation through hardware-assisted page migration," *arXiv preprint arXiv:2610.20962v1*, 2026, consulted 2610.20962v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.20962v1

[26] A. Adiletta and B. Sunar, "Spill the beans: Exploiting CPU cache side-channels to leak tokens from large language models," *arXiv preprint arXiv:2505.00817*, 2025. [Online]. Available: https://arxiv.org/abs/2505.00817

[27] H. Weng, R. Zhou, Y. Qu, X. Ji, and W. Xu, "ArchSnoop: LLM architecture snooping via electromagnetic side-channel on edge devices," in *Proceedings of the 19th ACM Conference on Security and Privacy in Wireless and Mobile Networks*, 2026, pp. 51–62. [Online]. Available: https://dl.acm.org/doi/10.1145/3765613.3811673

[28] F. Chen, L. Du, J. Zhou, P. Li, and Z. Su, "Sparleak: Privacy leakage from sparse attention in LLM inference on shared gpus," *arXiv preprint arXiv:2609.38830v1*, 2026, consulted 2609.38830v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.38830v1

[29] E. Hadad and M. Guri, "Shape and substance: Dual-layer side-channel attacks on local vision-language models," *arXiv preprint arXiv:2603.25403v1*, 2026, consulted 2603.25403v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2603.25403v1

[30] T. Hornetz, H. Yavarzadeh, A. Cheu, A. Gascon, L. Gerlach, D. Moghimi, P. Schoppmann, M. Schwarz, and R. Zhang, "Tdxray: Microarchitectural side-channel analysis of Intel TDX for real-world workloads," in *IEEE Symposium on Security and Privacy*, 2026. [Online]. Available: https://tdxray.cpusec.org/assets/tdxray_sp26.pdf

[31] Y. Jo, J. Park, E. Lee, and D. Song, "Sparseety: Extracting tokens from sparsity-exploiting LLM serving systems via deterministic side channels," *arXiv preprint arXiv:2608.02995v1*, 2026, consulted 2608.02995v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2608.02995v1

[32] J. Lee, Y. Wang, R. Rajat, D. Wong, M. Li, and M. Annavaram, "Otro: Oblivious tokenization path with square-root oram," *arXiv preprint arXiv:2606.17358v1*, 2026, consulted 2606.17358v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.17358v1

[33] R. Weiss, B. Konstantinov, E. Sheetrit, T. Simon, and Y. Mirsky, "Detokenization leaks: Reconstructing local LLM outputs from cache traces," *arXiv preprint arXiv:2609.06674v1*, 2026. [Online]. Available: https://arxiv.org/abs/2609.06674v1

[34] J.-J. Jiao, H. Yang, and R. Wen, "Llm-wfin: A fine-grained large language model (LLM)-oriented website fingerprinting attack via fusing interrupt trace and network traffic," *Electronics*, vol. 14, no. 7, p. 1263, 2025, consulted publisher version 2025-03-23; source-scoped access 21 October 2026. [Online]. Available: https://doi.org/10.3390/electronics14071263

[35] D. Scheinert, A. Acker, T. Wittkopp, S. Becker, H. Yous, K. Reddy, I. Farhat, H. Hacid, and O. Kao, "Beyond microservices: Testing web-scale RCA methods on GPU-driven LLM workloads," 2026, version 1, 2 March 2026. [Online]. Available: https://arxiv.org/html/2603.02057v1

[36] G. Politis and E. Pappas, "Privacy failure in split-LLM training, the returned gradient nullifies the decoys," *arXiv preprint arXiv:2609.04382v1*, 2026, consulted 2609.04382v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.04382v1

[37] T. C. Addagada, "Governing the KV cache: Preventing timing side-channel leakage in multi-tenant LLM inference," *arXiv preprint arXiv:2608.09225*, 2026, consulted arXiv:2608.09225v1; archived access 19 October 2026. [Online]. Available: https://arxiv.org/abs/2608.09225v1

[38] S. Alhazbi, A. Hussain, G. Oligeri, and P. Papadimitratos, "LLMs have rhythm: Fingerprinting large language models using inter-token times and network traffic analysis," *IEEE Open Journal of the Communications Society*, vol. 6, pp. 5050–5071, 2025. [Online]. Available: https://arxiv.org/abs/2502.20589v1

[39] N. Carlini and M. Nasr, "Remote timing attacks on efficient language model inference," *arXiv preprint arXiv:2410.17175v1*, 2024. [Online]. Available: https://arxiv.org/abs/2410.17175v1

[40] Z. Zhang, G. Wu, S. Deng, S. Wang, and Y. Zhang, "Netecho: From real-world streaming side-channels to full LLM conversation recovery," *arXiv preprint arXiv:2510.25472*, 2025, consulted arXiv:2510.25472v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2510.25472v1

[41] S. Li, T. Cui, M. Chen, X. Lin, Z. Gu, X. Deng, K. Xu, and Q. Li, "From length to content: Token-length side-channel attacks on LLM api merged outputs," in *35th USENIX Security Symposium*, 2026, pp. 2187–2206. [Online]. Available: https://www.usenix.org/conference/usenixsecurity26/presentation/li-sijia

[42] J. Wei, A. Abdulrazzag, T. Zhang, A. Muursepp, and G. Saileshwar, "When speculation spills secrets: Side channels via speculative decoding in llms," *arXiv preprint arXiv:2411.01076*, 2024, consulted arXiv:2411.01076v4; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2411.01076v4

[43] G. McDonald and J. B. Or, "Whisper leak: a side-channel attack on large language models," *arXiv preprint arXiv:2511.03675*, 2025, consulted arXiv:2511.03675v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2511.03675v1

[44] R. A. Bakar, "Characterizing contention-induced reliability collapse in KV-cache timing side channels for multi-tenant LLM serving," *arXiv preprint arXiv:2609.06853v1*, 2026. [Online]. Available: https://arxiv.org/abs/2609.06853v1

[45] T.-R. Chen, "Sensitive-topic leakage through LLM routing metadata: Measurement and mitigation," *arXiv preprint arXiv:2610.09981v1*, 2026, consulted 2610.09981v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.09981v1

[46] K. Chu, Z. Lin, D. Xiang, Z. Shen, J. Su, C. Chu, Y. Yang, W. Zhang, W. Wu, and W. Zhang, "Selective KV-cache sharing to mitigate timing side-channels in LLM inference," *arXiv preprint arXiv:2508.08438v2*, 2026, consulted arXiv:2508.08438v2; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2508.08438v2

[47] E. Debenedetti, G. Severi, N. Carlini, C. A. Choquette-Choo, M. Jagielski, M. Nasr, E. Wallace, and F. Tramèr, "Privacy side channels in machine learning systems," in *33rd USENIX Security Symposium*. USENIX Association, 2024. [Online]. Available: https://www.usenix.org/conference/usenixsecurity24/presentation/debenedetti

[48] C. Gu, X. L. Li, R. Kuditipudi, P. Liang, and T. Hashimoto, "Auditing prompt caching in language model apis," in *International Conference on Machine Learning*, vol. 267, 2025, pp. 20 477–20 496. [Online]. Available: https://proceedings.mlr.press/v267/gu25b.html

[49] X. Zheng, H. Han, S. Shi, Q. Fang, Z. Du, X. Hu, and Q. Guo, "Inputsnatch: Stealing input in LLM services via cache-sharing timing side-channel attacks," in *Advanced Parallel Processing Technologies (APPT 2026)*, 2027, pp. 483–497, first online August 24, 2026; publisher citation year 2027. [Online]. Available: https://link.springer.com/chapter/10.1007/978-981-92-4805-6_32

[50] S. Majidi, N. Mireshghallah, and K. Taram, "Leaky language models: Stealing architecture and inference optimizations via per-token timing," *arXiv preprint arXiv:2607.20723*, 2026, consulted arXiv:2607.20723v1; archived access 19 October 2026. [Online]. Available: https://arxiv.org/abs/2607.20723v1

[51] H. Sun, S. Liu, S. Ma, J. Li, M. Xiao, and W. Jiang, "Agent-assisted side-channel attacks on non-prefix KV cache in rag," *arXiv preprint arXiv:2606.21842*, 2026, consulted arXiv:2606.21842v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2606.21842v1

[52] T. Zhang, G. Saileshwar, and D. Lie, "Time will tell: Timing side channels via output token count in large language models," *arXiv preprint arXiv:2412.15431*, 2024, consulted arXiv:2412.15431v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2412.15431v1

[53] I. Yona, I. Shumailov, J. Hayes, and N. Carlini, "Stealing user prompts from mixture of experts," *arXiv preprint arXiv:2410.22884v1*, 2024. [Online]. Available: https://arxiv.org/abs/2410.22884v1

[54] P. Revankar, K. Chauhan, J. Kim, S. N. Nur, V. Siu, and C.-G. Wang, "Mirage: A polarity-flipping encoding subspace in LLM agents," *arXiv preprint arXiv:2606.10304v1*, 2026, consulted 2606.10304v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.10304v1

[55] S. Pouryousef, J. Lopez, S. R. Nowmi, M. M. Kamol, M. Hossain, M. Tran, and M. S. Rahman, "A large-scale benchmark and risk assessment of traffic analysis attacks on cloud LLM services," *arXiv preprint arXiv:2609.31877v1*, 2026, consulted arXiv:2609.31877v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2609.31877v1

[56] I. Fayolle, S. Bouhenniche, S. Pélissier, P. Laperdrix, C. Maurice, and W. Rudametkin, "On the internet, nobody knows you're an LLM bot: Unmasking web agents with multi-layer fingerprinting," *arXiv preprint arXiv:2606.30119v1*, 2026, consulted 2606.30119v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.30119v1

[57] R. Solomon, Y. Y. Levi, L. Vaknin, E. Aizikovich *et al.*, "LumiMAS: A comprehensive framework for real-time monitoring and enhanced observability in multi-agent systems," *arXiv preprint arXiv:2508.12412v2*, 2026. [Online]. Available: https://arxiv.org/abs/2508.12412v2

[58] F. George, H. Kumar, D. Pathak, K. Ray, M. Verma, and P. Moogi, "Unsupervised cycle detection in agentic applications," *arXiv preprint arXiv:2511.10650v1*, 2025. [Online]. Available: https://arxiv.org/abs/2511.10650v1

[59] J. Geng, R. He, Z. Fei, B. Yi, R. Wang, Z. Liu, X. Hu, X. Wu, and Q. Zeng, "Agent skills matter: Inferring proprietary skills from execution trajectories," *arXiv preprint arXiv:2607.25560v1*, 2026, consulted arXiv:2607.25560v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2607.25560v1

[60] H. Oderinwale, "Agent trajectories as programs: fingerprinting and programming coding-agent behavior," *arXiv preprint arXiv:2606.16988v1*, 2026. [Online]. Available: https://arxiv.org/abs/2606.16988v1

[61] X. Wang, B. Zhao, M. Backes, F. Boenisch, and A. Dziedzic, "Agentprov: Auditing agentic LLM api providers via tool-use policy probes," *arXiv preprint arXiv:2609.00052v1*, 2026, consulted 2609.00052v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.00052v1

[62] C.-Y. Wang, X.-H. Xie, T.-Z. Wang, F.-C. Luo, and Y. Cui, "Who is behind the harness? fingerprinting LLMs through agentic behavior," *arXiv preprint arXiv:2609.28559v1*, 2026, consulted 2609.28559v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.28559v1

[63] S.-J. Park, M. Choi, J. Kim, and M.-S. Kang, "Cross-agent campaign attribution: Linking asynchronous attacks across LLM agents," *arXiv preprint arXiv:2607.18826v1*, 2026, consulted 2607.18826v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2607.18826v1

[64] C. Song, H. Wang, and R. Martin, "Personafingerprint: Measuring persona inference on modern websites with LLM-driven browsing," *arXiv preprint arXiv:2605.15962*, 2026. [Online]. Available: https://arxiv.org/abs/2605.15962

[65] D. Kang, H. Jeong, J. Sheffey, P. Datta, and A. Houmansadr, "Whose agent are you? multi-layer fingerprinting and attribution of autonomous web agents," *arXiv preprint arXiv:2606.20910v1*, 2026. [Online]. Available: https://arxiv.org/abs/2606.20910v1

[66] A. Shahriar, M. N. Rahman, S. Ahmed, F. Sadeque, and M. R. Parvez, "AgentTell: Behavioural side-channel leakage in browser-use agents," *arXiv preprint arXiv:2609.32915v1*, 2026, consulted arXiv:2609.32915v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2609.32915v1

[67] E. Wang, Z. Shafiq, and Y. Vekaria, "Fp-agent: Fingerprinting AI browsing agents," *arXiv preprint arXiv:2605.01247*, 2026, consulted arXiv:2605.01247v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2605.01247v1

[68] S. Banerjee, P. Sahu, M. Luo, A. Vahldiek-Oberwagner, N. J. Yadwadkar, and M. Tiwari, "SoK: A systems perspective on compound AI threats and countermeasures," *arXiv preprint arXiv:2411.13459v1*, 2024, consulted arXiv:2411.13459v1; inspected 20 October 2026. [Online]. Available: https://arxiv.org/abs/2411.13459v1

[69] J. Forough, M. Kogias, and H. Haddadi, "When agents handle secrets: A survey of confidential computing for agentic AI," *arXiv preprint arXiv:2605.03213v2*, 2026, consulted arXiv:2605.03213v2; inspected 20 October 2026. [Online]. Available: https://arxiv.org/abs/2605.03213v2

[70] O. Kotevska, S. Jha, A. Bellet, R. Hu, N. D. Bastian, R. F. da Silva, R. Madduri, and K. Kim, "Privacy foundations for multi-institutional scientific artificial intelligence," *arXiv preprint arXiv:2609.39787v1*, 2026, consulted 2609.39787v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.39787v1

[71] X. Cai, R. Nithyanand, T. Wang, R. Johnson, and I. Goldberg, "A systematic approach to developing and evaluating website fingerprinting defenses," in *ACM CCS*, 2014, pp. 227–238. [Online]. Available: https://doi.org/10.1145/2660267.2660362

[72] M. Juarez, M. Imani, M. Perry, C. Diaz, and M. Wright, "Toward an efficient website fingerprinting defense," in *ESORICS*, 2016, pp. 27–46. [Online]. Available: https://arxiv.org/abs/1512.00524

[73] T. Wang and I. Goldberg, "Walkie-talkie: An efficient defense against passive website fingerprinting attacks," in *26th USENIX Security Symposium*, 2017. [Online]. Available: https://www.usenix.org/conference/usenixsecurity17/technical-sessions/presentation/wang-tao

[74] F. Croce, M. Andriushchenko, V. Sehwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, and M. Hein, "RobustBench: a standardized adversarial robustness benchmark," in *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, vol. 1, 2021. [Online]. Available: https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/a3c65c2974270fd093ee8a9bf8ae7d0b-Abstract-round2.html

[75] C. Dwork and A. Roth, "The algorithmic foundations of differential privacy," *Foundations and Trends in Theoretical Computer Science*, vol. 9, no. 3–4, pp. 211–407, 2014. [Online]. Available: https://www.cis.upenn.edu/~aaroth/Papers/privacybook.pdf

[76] G. Smith, "On the foundations of quantitative information flow," in *Foundations of Software Science and Computational Structures*, ser. LNCS, vol. 5504, 2009, pp. 288–302. [Online]. Available: https://link.springer.com/chapter/10.1007/978-3-642-00596-1_21

[77] J. B. Almeida, M. Barbosa, G. Barthe, F. Dupressoir, and M. Emmi, "Verifying Constant-Time implementations," in *25th USENIX Security Symposium*, 2016, pp. 53–70. [Online]. Available: https://www.usenix.org/conference/usenixsecurity16/technical-sessions/presentation/almeida

[78] E. Stefanov, M. van Dijk, E. Shi, C. Fletcher, L. Ren, X. Yu, and S. Devadas, "Path ORAM: An extremely simple oblivious RAM protocol," in *ACM SIGSAC Conference on Computer and Communications Security*, 2013, pp. 299–310. [Online]. Available: https://doi.org/10.1145/2508859.2516660

[79] N. Nazari, F. Xiang, C. Fang, H. M. Makrani, A. Puri, K. Patwari, H. Sayadi, S. Rafatirad, C.-N. Chuah, and H. Homayoun, "LLM-FIN: Large language models fingerprinting attack on edge devices," in *International Symposium on Quality Electronic Design*, 2024. [Online]. Available: https://web.cs.ucdavis.edu/~srafatir/pages/documents/llm_fingerprint.pdf

[80] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed. Wiley-Interscience, 2006, theorem 2.8.1, Data-processing inequality. [Online]. Available: https://onlinelibrary.wiley.com/doi/book/10.1002/047174882X

[81] P. Sirinam, M. Imani, M. Juarez, and M. Wright, "Deep fingerprinting: Undermining website fingerprinting defenses with deep learning," in *ACM CCS*, 2018, pp. 1928–1943.

[82] J. Gong, W. Zhang, C. Zhang, and T. Wang, "WFDefProxy: Modularly implementing and empirically evaluating website fingerprinting defenses," arXiv:2111.12629v1, 2021. [Online]. Available: https://arxiv.org/abs/2111.12629v1

[83] E. Rescorla, K. Oku, N. Sullivan, and C. A. Wood, "TLS Encrypted Client Hello," RFC 9849, Mar. 2026. [Online]. Available: https://www.rfc-editor.org/rfc/rfc9849

[84] Kubernetes Authors, "Kubernetes API Concepts: Dry-run," https://kubernetes.io/docs/reference/using-api/api-concepts/#dry-run, accessed October 24, 2026.

[85] Stripe, "Sandboxes," https://docs.stripe.com/sandboxes, accessed October 24, 2026.

[86] PostgreSQL Global Development Group, "PostgreSQL 18 Documentation: ROLLBACK," https://www.postgresql.org/docs/18/sql-rollback.html, accessed October 24, 2026.

[87] ——, "PostgreSQL 18 Documentation: Sequence Manipulation Functions," https://www.postgresql.org/docs/18/functions-sequence.html, accessed October 24, 2026.

[88] R. H. Anwar, J. O. Weiss, T. Alves, S. Kundu, and M. T. Raza, "InferNet: Exploiting aggregate GPU profiles as side-channel for DNN architecture inference," *ACM Transactions on AI Security and Privacy*, 2026. [Online]. Available: https://rhasnainanwar.me/assets/pdf/2304.03388v2.pdf

[89] W. Cui, H. Zhang, Y. Huang, Q. Li *et al.*, "SymDiag: Explainable diagnosis for LLM reasoning via neuro-symbolic verification," in *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, 2026, consulted arXiv:2608.08786v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2608.08786v1

[90] N. Lahjouji and A. G. Colaco, "Agents that know too much: A data-centric survey of privacy in LLM agents," *arXiv preprint arXiv:2606.26627*, 2026, consulted arXiv:2606.26627v1; archived access 19 October 2026. [Online]. Available: https://arxiv.org/abs/2606.26627v1

[91] J. Roh, E. Bagdasarian, H. Haddadi, and A. S. Shamsabadi, "SPILLage: Agentic oversharing on the web," *arXiv preprint arXiv:2602.13516v1*, 2026. [Online]. Available: https://arxiv.org/abs/2602.13516v1

[92] G. Yu, Q. Wang, R. Lang, S. Su, and X. Wang, "PlanTwin: Privacy-preserving planning abstractions for cloud-assisted LLM agents," *arXiv preprint arXiv:2603.18377v2*, 2026. [Online]. Available: https://arxiv.org/abs/2603.18377v2

[93] J. Xu, L. Fan, Z. Wang, X. Li, and H. Liu, "Selection-invariant communication compilers for privacy-aware multi-agent LLM workflows," *arXiv preprint arXiv:2609.26076v1*, 2026. [Online]. Available: https://arxiv.org/abs/2609.26076v1

[94] Y. Qiao, D. Liu, H. Yang, W. Zhou, and S. Hu, "Agent tools orchestration leaks more: Dataset, benchmark, and mitigation," *arXiv preprint arXiv:2512.16310v4*, 2026, version 4. [Online]. Available: https://arxiv.org/abs/2512.16310v4

[95] M. Umar, A. P. Marathe, M. D. Gupta, S. J. Ghosh, G. E. Suh, and W. Xiong, "Efficient memory side-channel protection for embedding generation in machine learning," in *IEEE International Symposium on High-Performance Computer Architecture*, 2025. [Online]. Available: https://bearhw.ece.vt.edu/content/dam/bearhw_ece_vt_edu/publications/2025_HPCA_SecEmb.pdf

[96] L. Wu, C. Wang, T. Liu, Y. Zhao, and H. Wang, "From assistants to adversaries: Exploring the security risks of mobile LLM agents," *arXiv preprint arXiv:2505.12981v1*, 2025, version 1. [Online]. Available: https://arxiv.org/abs/2505.12981

[97] J. W. Kim, B. Min, K. Edemacu, Y. Choi, S.-H. Cho, and B. Jang, "Dp-memview: A memory interface for attribute-level transcript privacy in long-term LLM agents," *arXiv preprint arXiv:2608.03130v1*, 2026, consulted 2608.03130v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2608.03130v1

[98] M. Mudryi, M. Chaklosh, and G. M. Wójcik, "The hidden dangers of browsing AI agents," *arXiv preprint arXiv:2505.13076v1*, 2025, version 1. [Online]. Available: https://arxiv.org/abs/2505.13076

[99] Y.-M. Wang, Y.-X. Li, S.-L. Fan, Y.-H. Zheng, and T.-H. Chang, "Ask without telling: Local slms consult cloud LLMs without revealing task intent," *arXiv preprint arXiv:2609.32642v1*, 2026, consulted 2609.32642v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.32642v1

[100] A. Ukani, H. Haddadi, A. S. Shamsabadi, and P. Snyder, "Privacy practices of browser agents," *arXiv preprint arXiv:2512.07725*, 2025. [Online]. Available: https://arxiv.org/abs/2512.07725

[101] Z. Liu, S. van Hoek, P. Horváth, D. Lauret, X. Xu, and L. Batina, "Real-world edge neural network implementations leak private interactions through physical side channel," *arXiv preprint arXiv:2501.14512v1*, 2025, consulted 2501.14512v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2501.14512v1

[102] P. G. Pennas, K. Papaioannou, M. Guarnieri, and T. D. Doudali, "Cachesolidarity: Preventing prefix caching side channels in multi-tenant LLM serving systems," *arXiv preprint arXiv:2603.10726v1*, 2026, consulted 2603.10726v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2603.10726v1

[103] C.-Y. Zhou, Q.-L. Jiang, S.-N. Wu, and X.-H. Zhou, "The vision encoder as a privacy boundary: Visual-token side channels in encoder-free vision-language models," *arXiv preprint arXiv:2606.14783v1*, 2026, consulted 2606.14783v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.14783v1

[104] G. Politis and E. Pappas, "What gradients add to text leakage in split language models, counted per token and per document," *arXiv preprint arXiv:2610.04128v1*, 2026, consulted 2610.04128v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.04128v1

Appendix A. Audit Trail

The artifact preserves the 82-family map, 39-record ledger, 180 family cells and 351 record cells with source versions and locators. S33 publishes query membership and sensitivity checks. Case-study protocols, observations and uncertainty calculations retain their original cohorts and provenance.

Appendix B. Result-record index

Table 20 indexes all 39 records by measurement mode and source locator. Marks distinguish six contextual controls (‡) and six final-search additions (†); all enter the pair inventory. The ledger retains full locators, versions and hashes.

Appendix C. Case-Study Artifact Index

Table 21 maps the single appendix case to historical artifact keys and supplementary methods. These keys preserve original task/customer memberships; S30 retains the corrected permutation/Holm tests and calibration resampling.

Appendix D. Family-Level Holdout Evidence

This is a family-level discovery index, not a count of defense-transfer tests. The three model H entries concern attack/identification generalization; even an H in a mitigation family must be linked to its defended condition before reuse. The main RQ4 table separates those evidence objects, and S34 identifies every H entry. Table 22 covers 82 families and 492 cells. Axes: T task/input/user; M model; C harness/topology; H hardware; N network/load; D time. Codes: H reported holdout; V condition comparison/recalibration; X scoped restriction; ? unresolved; – conceptual/inapplicable. Scope C/A denotes core/adjacent. H records separation, not transfer success; ? does not establish an absent experiment. S28 retains source locators and qualified interpretations.

Appendix E. Illustrative Case Study and Reconstruction Scope

This appendix explains the role of the archived implementation examples. They instantiate the same distinctions used in the literature atlas: final receiver versus upstream trace, executed state versus discarded responses, and online planner utility versus reused trajectories. They are illustrative applications, not validation data for the comparison rules. The complete numerical grids, negative outcomes, intervals and corrected tests remain unchanged in S11–S32.

The transport example has 48 saved responses on 24 topics (12 train/12 test). The public *retail* setting has 115 tasks from 53 customers; its frozen split contains 46 test tasks from 24 customers, including 5 cancellation-positive and 7 return-positive customers. The separate live planner example has 100 episodes: 37/50 baseline and 33/50 defended successes, a difference of -8.0 points with customer interval $[-23.9, 6.5]$. All eight calibration contrasts have Holm-adjusted $p = 1.000$; the original absolute protection screen remains unmet. Neither these unresolved contrasts nor resampling the same customers validates the framework.

The artifact separates *analysis reconstruction* from *new model collection*. Archived requests, responses, traces, calibration hashes and predictions permit the reported calculations to be reconstructed. Fresh calls to mutable Sol/Luna/MiniMax/GLM aliases, with documented unequal budgets, cannot reproduce an identical generator. No fixed-weight replication claim is made. Historical protocols and

model identities remain explicit in S14–S32; the literature synthesis does not depend on obtaining a favorable result from these examples.

Appendix F. Reconstructing the Audit in This PDF

Table 23 gives the reporting-pattern sensitivity to removing CachePrune and final-search additions. The latter removes six records from four families (the † marker in the record index); six contextual ‡ controls are excluded throughout. Removing both CachePrune and all final-search additions equals the latter row. This preserves the inference–utility and matched-adaptation patterns, and the agent subset remains 0/4 B or 2/4 B+P. S38 gives exact identities and the common-denominator calculation.

Table 24 provides all 33 mitigation records' query memberships, including inapplicability; its column totals reconstruct the positive counts and eligible denominators in Table 13. Table 25 explains every family inference/cost conjunction. The 39 record IDs, measurement modes and source locators are in Appendix B.

Pair inventory and granularity. Table 26 covers all 133 within-question pair decisions. The 105 different-object pairs and eight of 15 within-family pairs reflect the declared split and grouping rules; a normal difference in research scope is not a source-paper defect. Table 27 projects the unchanged judgments onto coarser units. Mixed edges retain all constituent statuses, and within-unit edges disappear from the index without being resolved. Family-only indexing yields 44 edges, eight mixed, and hides all 15 within-family pairs. The pair CSV preserves the locators and individual reasons. These indexing sensitivities support audit reconstruction; the reporting findings use the record cells directly.

TABLE 20. COMPLETE INDEX OF THE 39 RESULT RECORDS (26 FAMILIES). MODE IS THE PRIMARY MEASUREMENT MODE: EX EXECUTED COLLECTION/PROTOTYPE (20), TR TRACE TRANSFORMATION (6), SI SIMULATION/DERIVED ESTIMATE (5), CO COMPONENT/CALIBRATION MEASUREMENT (6), UN PLACEMENT UNRESOLVED (2). ‡: CONTEXTUAL RECORD (6); †: ADDED BY THE FINAL SEARCH (6). T=TABLE, F=FIGURE, APP.=APPENDIX; LOCATORS REFER TO THE CITED SOURCE.

Record ID / family	Mode	Boundary; source locator
UI-delay-offline [5]	Tr	Offline delay; no live task outcome. 6.1 F5; A.5
WRA-decoy-execution [6]	Ex	Executed decoys; usefulness proxy only. 8.1–8.3 T8–11
NET-upstream-local [40]	Ex	Protected deployment/local serving; topology scoped. IV-C–E TIII–IV
NET-dummy-padding-simulation [40]	Tr	Synthetic insert/pad; no live execution. VII TIX
EARLY-K-calibration-oracle [3]	Si	K-specific fit then oracle; not live KV trial. V-A TVI; IV-A
EARLY-semantic-deidentification [3]	Ex	GPTCache prototype cost; not KV-granularity cost. V-B
WHISPER-randompadding-capture [43]	Ex	New padded captures; rounded medians. 6.1 T3–4
WHISPER-batching-replay [43]	Tr	Offline batching; model-dependent result. 6.2 F7
WHISPER-injection-replay [43]	Tr	Synthetic packets; no live zero-delay claim. 6.3 T5
SPLICE-CTBF-signal [51]	Ex	Prototype background load; tested signal only. VII-C TII
SPLICE-QCP-cost [51]	Co	QCP component cost; excludes combined CTBF. VII-C TII
TIME-length-instruction [52]	Ex	Length instruction; unequal outputs remain. 7.2
TIME-padding-estimate [52]	Si	Analytic plaintext cost; not network bytes. 7.3
SPEC-padding-REST [42]	Un	Padding/aggregation; payload cost, placement unverified. 6.1 F6–7 T2
DANGOL-capture [19]	Tr	Capture transform; distinct from live population. 9.8–9.9 T6
DANGOL-livebinding [19]	Ex	Live loopback binding; do not join capture accuracy. 9.7, OHTTP paragraph after T5
SIG-common-guardrail‡ [59]	Ex	Common instruction; not a causal guardrail effect. 4.1–4.2; App. C
SAFE-post-trigger [46]	Ex	Post-trigger probes; detector cost not service TTFT. 7.2.4 F12; 7.3–7.4
AHB-online-executed [17]	Ex	Executed UI edits; task/state changes possible. 5.1.2 T1; E.1–E.3
AHB-offline [17]	Tr	Offline touch edits; no online task damage claim. T1 offline rows; E.1–E.3
KVGOV-partialflush-sim [37]	Si	Deterministic simulation; not live protected recovery. 5.1–5.3
KVGOV-hardware-gap [37]	Co	Hardware gap calibration; not protected ASR. 5.7–5.8
KVGOV-HMAC-cost [37]	Co	HMAC microcost; not cache reuse/task loss. 5.9
TRAFFIC-prompt-reformulation [55]	Ex	Collected prompt rewrites; task utility unmeasured. V-C1 F9a
TRAFFIC-temperature [55]	Ex	Collected decoding changes; no monotone benefit. V-C1 F9a
CARLINI-fixed-rate [39]	Un	Rate/replayed cost; per-token, not request delay. VII F11–12
GAO-zero-copy-off [23]	Co	Loader component costs; not end-to-end TTFT. 7.1
TDXRAY-oblivious-map [30]	Co	Lookup microbenchmark; not full inference cost. 6.1 F8 T3
AGENTTELL-fixed-instruction‡ [66]	Ex	Live local browser tasks; loaded/cold context contrast. 4.3.1; 5 Finding 1
FP-CLOUDFLARE‡ [67]	Ex	Live access control; not user confidentiality. 6.4 T5
BAR-co-running-noise‡ [24]	Ex	MIG noise test; not an LLM defense trial. 7.2–7.3 T8
RHYTHM-VPN‡ [38]	Ex	Live VPN shift; baseline-trained attacker only. VI-B; VII
DEB-filter-origin‡ [47]	Ex	Filter creates attack view; not mitigation of it. Threat model; 4.2
CACHEPRUNE-selective† [7]	Ex	Selective cache sharing; direct token recovery/quality. 3.1; 5.1–5.6; T1; App. A
CACHEPRUNE-context† [7]	Ex	Contextual guessing; distinct from direct recovery. 3.1; 5.1–5.6; T1; App. A
OTRO-unpadded† [32]	Co	ORAM TTFT/memory; access-count padding excluded. VII-C–E; TIII
OTRO-epochpad† [32]	Si	Epoch-padding estimate; not measured padded TTFT. VII-C–E; TIII
SPARLEAK-replacement† [28]	Ex	Sparse replacement/quality; attack cost is separate. 3.1; 4–7; T4
SHAPE-padding-estimate† [29]	Si	Constant-work proposal/estimate; no executed pair. IX-B1

TABLE 21. ARTIFACT INDEX FOR THE APPENDIX CASE STUDY. COMPONENT NAMES DESCRIBE THE ARCHIVED MEASUREMENT BOUNDARY; THE RIGHT COLUMN LOCATES IMPLEMENTATION AND FULL-GRID DETAIL.

Case-study component	Artifact keys	Observation/execution boundary and comparison	Detailed source
Transport emission	E1, E3, E5	Buffered two-link re-emission and token-stream emitters; upstream framing, final-link framing and matched attacker refitting are distinct comparisons.	S9/S11; S12
State and link	E2, E6, E7	Prescribed stateful plans and randomized dispatch cover, then agent–tool TLS composition; dispatch-visible calls and encrypted-link records are separate views.	S9/S11; S12
Trajectories and localization	E9, E10, E12	Saved autonomous Sol trajectories; original pre-test calibration, exploratory customer cross-fitting and oracle operation-category localization retain their different fit/test boundaries.	S14; S15/S21
Timing and WF controls	E13, E14	Local live tool-loop timing/full tunnels, then stronger learners and specified WF adaptations; measured clocks and simulations are distinguished, with matched suites for joint comparisons.	S19; S20
Extensions	E8	Stand-in-planner pipeline check.	S13
	E11	Retrospective training-fold operating-point selection with the historical suite.	S15
	E15	Additional same-provider Luna cohort.	S22
	E16	Same-operation semantic-attribute diagnostic on reused Sol data.	S23.6
	R11/R12	Two served model endpoints: native trajectories, frozen training calibration and TLS write localization.	S24
	R14	Attacker/calibration transfer matrices, non-oracle write detection and write-only framing with the endpoint suite.	S25
	R15	Fresh paired live-planner episodes under isolated prefix cover and injected tool-response deadlines.	S26

TABLE 22. COMPLETE HOLDOUT-CODE MATRIX; CODES AND AXES ARE DEFINED IN APPENDIX D. SOURCE-QUALIFIED ROW NOTES ARE IN SUPPLEMENT S28.

Family	Scope	T	M	C	H	N	D	Family	Scope	T	M	C	H	N	D
F01 AgentPrint [4]	C	?	?	?	?	?	?	F42 ML privacy channels [47]	C	?	V	?	?	?	?
F02 UI Traces [5]	C	V	V	V	?	?	?	F43 MoE stealing [53]	C	?	?	?	?	?	?
F03 WRA [6]	C	H	?	?	?	?	?	F44 Energon [21]	C	?	V	?	?	?	?
F04 FP-Agent [67]	C	H	?	X	?	?	?	F45 SafeKV [46]	C	?	V	?	?	?	?
F05 Prompt length [1]	C	H	?	V	?	V	V	F46 InferNet [88]	A	?	V	?	V	?	?
F06 PromptPeek [2]	C	V	?	?	?	V	?	F47 ArchSnoop [27]	C	?	H	?	V	?	?
F07 NetEcho [40]	C	V	?	V	?	?	?	F48 Coding fingerprints [12]	A	X	?	X	?	?	X
F08 Early Bird [3]	C	?	V	?	?	?	?	F49 AHB [17]	C	?	?	X	?	?	?
F09 InputSnatch [49]	C	H	?	?	?	?	?	F50 WildAuth [13]	A	H	?	X	?	?	?
F10 Whisper Leak [43]	C	H	V	?	?	?	?	F51 Cycles [58]	C	X	?	X	?	?	X
F11 Spill the Beans [26]	C	V	?	?	?	?	?	F52 LumiMAS [57]	C	?	V	V	?	?	?
F12 SpliceLeak [51]	C	V	?	?	?	?	?	F53 SymDiag [89]	A	V	V	?	?	?	?
F13 Time Will Tell [52]	C	?	?	?	?	V	V	F54 LeakyLMs [50]	C	?	H	?	V	?	?
F14 Prompt Echo [41]	C	H	V	?	?	?	?	F55 KVGov [37]	C	?	?	?	V	?	?
F15 I Know What You Said [23]	C	H	V	V	V	?	?	F56 Data-centric survey [90]	A	-	-	-	-	-	-
F16 Speculation [42]	C	H	?	?	?	V	?	F57 Serving RCA [35]	C	?	?	X	?	?	?
F17 Behind Bars [24]	C	?	V	?	V	?	?	F58 ZeroPrint [14]	A	?	V	?	?	?	?
F18 Prompt caching [48]	C	?	V	?	?	?	V	F59 Endpoint fingerprints [15]	A	?	?	V	?	?	V
F19 Whose Agent [65]	C	?	?	?	?	?	?	F60 StepFinder [16]	A	H	?	?	?	?	?
F20 Workflow traffic [19]	C	?	?	H	?	?	?	F61 Web-agent detection [56]	C	?	V	?	?	?	?
F21 GPU contention [44]	C	?	?	V	?	V	?	F62 Cross-campaign [63]	C	V	?	V	?	?	?
F22 AgentTell [66]	C	V	?	?	?	V	?	F63 AgentProv [61]	C	?	V	?	?	?	?
F23 Traffic benchmark [55]	C	?	?	H	?	?	?	F64 LiDAR [62]	C	?	V	V	?	?	?
F24 Remote timing [39]	C	?	V	?	?	?	V	F65 LLM-WFIN [34]	C	V	?	?	?	?	?
F25 Persona fingerprint [64]	C	?	?	X	?	?	?	F66 MIRAGE [54]	C	V	V	?	?	?	?
F26 TDXRay [30]	C	?	?	?	?	?	?	F67 Routing metadata [45]	C	V	?	V	?	?	?
F27 Detokenization [33]	C	?	?	X	?	?	?	F68 Shape and Substance [29]	C	V	V	V	V	?	?
F28 Spillage [91]	A	V	?	?	?	?	?	F69 SplitDecoy [36]	C	?	?	?	?	?	?
F29 PlanTwin [92]	A	V	?	?	?	?	?	F70 PAGE404 [25]	C	V	V	?	?	?	?
F30 SICC [93]	A	?	?	?	?	?	?	F71 Edge EM [20]	C	H	?	?	?	?	?
F31 TOPR [94]	A	V	?	?	?	?	?	F72 SparLeak [28]	C	V	V	V	?	?	?
F32 ClawGuard [18]	C	?	?	V	?	?	V	F73 Sparseety [31]	C	V	V	V	?	?	?
F33 SIGLEAK [59]	C	V	?	?	?	?	?	F74 OTRO [32]	C	?	?	?	?	?	?
F34 ProcGrep [60]	C	H	?	V	?	?	?	F75 CachePrune [7]	C	V	V	?	?	?	?
F35 SecEmb [95]	A	?	X	?	?	?	?	F76 Privacy foundations [70]	A	-	-	-	-	-	-
F36 Mobile agents [96]	A	?	?	?	?	?	?	F77 DP-MemView [97]	A	?	?	?	?	?	?
F37 Hidden risks [98]	A	?	?	?	?	?	?	F78 PriCon [99]	A	?	?	?	?	?	?
F38 Browser privacy [100]	A	?	?	V	?	?	?	F79 SCAAR [101]	A	?	?	?	?	?	?
F39 LLMs Have Rhythm [38]	C	?	V	?	?	H	V	F80 CacheSolidarity [102]	A	?	?	?	?	?	?
F40 MoEcho [22]	C	H	V	?	?	V	?	F81 Vision encoder [103]	A	?	?	?	?	?	?
F41 LLM-Fin [79]	A	?	H	?	?	?	?	F82 Split gradients [104]	A	?	?	?	?	?	?

TABLE 23. SENSITIVITY OF REPORTING PATTERNS. FRACTIONS GIVE EXPLICIT / APPLICABLE FAMILY COUNTS; BRACKETED NUMERATORS ALSO ACCEPT P. REMOVING FINAL-SEARCH ADDITIONS EXCLUDES SIX RECORDS FROM FOUR FAMILIES; CONTEXTUAL CONTROLS ARE EXCLUDED IN EVERY ROW. IC SEPARATE MEANS NO SINGLE RECORD CARRIES BOTH EXPLICIT INFERENCE AND COST.

Inventory	R7	R9	R8	R4	R5	Agent R8	IC separate
All mitigation	16/17 [17]	11/19 [15]	4/17 [6]	4/11 [9]	2/11 [2]	0/4 [2]	5/8
Without CachePrune	15/16 [16]	10/18 [14]	3/16 [5]	4/11 [9]	2/11 [2]	0/4 [2]	5/7
Without final-search records	13/13 [13]	9/15 [12]	2/14 [4]	4/9 [8]	2/9 [2]	0/4 [2]	4/6

TABLE 24. ALL 33 MITIGATION RECORDS: QUERY MEMBERSHIP. B: ELIGIBLE AND BOUND; E: ELIGIBLE WITH AT LEAST ONE INCOMPLETE BINDING; X: REQUIRED ITEM INAPPLICABLE. IU=INFERENCE/UTILITY; IC=INFERENCE/COST; AI=ADAPTATION/INFERENCE; IUC=TRIPLE; ALL=R4/R5/R7/R8/R9. FULL R1–R9 SOURCE RATIONALES REMAIN IN S28.

Record / cited family	IU	IC	AI	IUC	All
UI-delay-offline [5]	E	x	B	x	x
WRA-decoy-execution [6]	E	E	x	E	x
NET-upstream-local [40]	E	E	E	E	E
NET-dummy-padding-simulation [40]	E	E	E	E	E
EARLY-K-calibration-oracle [3]	E	E	E	E	E
EARLY-semantic-deidentification [3]	x	x	x	x	x
WHISPER-randompadding-capture [43]	E	E	E	E	E
WHISPER-batching-replay [43]	E	E	E	E	E
WHISPER-injection-replay [43]	E	E	E	E	E
SPLICE-CTBF-signal [51]	E	E	x	E	x
SPLICE-QCP-cost [51]	x	x	x	x	x
TIME-length-instruction [52]	E	E	E	E	E
TIME-padding-estimate [52]	x	x	x	x	x
SPEC-padding-REST [42]	x	B	E	x	x
DANGOL-capture [19]	E	B	B	E	E
DANGOL-livebinding [19]	E	B	E	E	E
SAFE-post-trigger [46]	E	E	x	E	x
AHB-online-executed [17]	B	E	E	E	E
AHB-offline [17]	E	x	E	x	x
KVGOV-partialflush-sim [37]	E	E	x	E	x
KVGOV-hardware-gap [37]	x	x	x	x	x
KVGOV-HMAC-cost [37]	x	x	x	x	x
TRAFFIC-prompt-reformulation [55]	E	E	E	E	E
TRAFFIC-temperature [55]	E	E	E	E	E
CARLINI-fixed-rate [39]	x	x	x	x	x
GAO-zero-copy-off [23]	x	x	x	x	x
TDXRAY-oblivious-map [30]	x	x	x	x	x
CACHEPRUNE-selective [7]	B	B	x	B	x
CACHEPRUNE-context [7]	B	B	x	B	x
OTRO-unpadded [32]	x	x	x	x	x
OTRO-epochpad [32]	x	E	x	x	x
SPARLEAK-replacement [28]	B	E	E	E	E
SHAPE-padding-estimate [29]	E	E	E	E	E

TABLE 25. EVERY FAMILY WITH B IN BOTH INFERENCE AND COST COLUMNS OF TABLE 12. SEPARATE COLUMN WITNESSES DO NOT AUTOMATICALLY FORM ONE ROW-LEVEL RESULT. IDS EXPAND IN APPENDIX B.

Family	Inference record(s)	Cost record(s)	Same-record inference + cost
[3]	EARLY-K-calibration-oracle	EARLY-semantic-deidentification	–
[51]	SPLICE-CTBF-signal	SPLICE-QCP-cost	–
[52]	TIME-length-instruction	TIME-padding-estimate	–
[42]	SPEC-padding-REST	SPEC-padding-REST	SPEC-padding-REST
[19]	DANGOL-capture; DANGOL-livebinding	DANGOL-capture; DANGOL-livebinding	DANGOL-capture; DANGOL-livebinding
[37]	KVGOV-partialflush-sim	KVGOV-HMAC-cost	–
[7]	CACHEPRUNE-selective; CACHEPRUNE-context	CACHEPRUNE-selective; CACHEPRUNE-context	CACHEPRUNE-selective; CACHEPRUNE-context
[32]	OTRO-epochpad	OTRO-unpadded	–

TABLE 26. ALL 133 WITHIN-QUESTION DECISIONS; 608 OTHER PAIRS HAVE NO DECLARED COMPARISON. COUNTS DESCRIBE THIS INVENTORY AND ITS GROUPING RULES.

Question	Match	Different object	Unresolved
behavioral humanization	1	2	0
cache service	2	37	6
fixed instruction	0	0	1
generator fingerprint	3	0	0
local memory	0	14	1
traffic envelope	11	52	3
Total	17	105	11

TABLE 27. SENSITIVITY TO THE COMPARISON UNIT. C/M/U: UNANIMOUS DESCRIPTIVE MATCH/OBJECT MISMATCH/UNRESOLVED; MIX: CONSTITUENT JUDGMENTS DIFFER. HIDDEN: ORIGINAL WITHIN-QUESTION PAIRS ABSORBED INSIDE A UNIT. NO NEW PAIR IS ADJUDICATED.

Unit	Units	C	M	U	Mix	Hidden
Scoped result records	39	17	105	11	0	0
Shared configuration package	38	16	98	9	1	1
Family and measurement mode	35	10	83	9	1	4
Family only	26	2	31	3	8	15