SoK: Transfer Boundaries of LLM and Agent Side-Channel Defenses

***Abstract*—Side-channel defenses for LLM serving and agent systems are evaluated at one operating point and reused at another. This SoK contributes record-level comparison, configuration-level joint binding and deployment decisions for reusing defense evidence. We organize 82 study families and audit 39 result records from 26 families, binding each outcome to its target, receiver, intervention, attacker fit, workload and cost object. First, a complete pair inventory distinguishes source-qualified comparisons, known object mismatches and unresolved bindings. Comparisons within the same paper expose changes of primitive, target and cost that mechanism-level summaries conceal. Second, configuration-linked joins distinguish outcomes that can be interpreted together. CachePrune illustrates this distinction: two recorded outcomes share one inference–utility–cost package at a single configuration. Third, a six-axis holdout matrix connects reported input, model, topology, hardware, network and temporal changes to deployment decisions. The resulting procedure identifies which comparison is supported, which protection, utility and cost outcomes belong together, and which changed boundary requires further measurement. It turns the literature audit into an actionable specification of the evidence needed to reuse a defense claim.**

1. Introduction

Encryption conceals message content, yet LLM engines and agents expose packet sizes, cache timing, visited domains and tool sequences [1], [2], [3], [4], [5], [6]. Defenses follow: padding, batching, cache isolation, decoy queries and action humanization. Each is evaluated at one receiver, execution boundary and workload, yet its headline number is cited as a property of the mechanism. A padding result measured on an upstream link says nothing about a frontend that re-emits decoded text; decoys whose read responses are discarded say nothing about cover that executes writes; and the overhead of a semantic cache does not price prefix isolation reported in the same paper.

We systematize defense evidence around three operations: *compare result records, join outcomes at one configuration, and decide what evidence a deployment change requires*. A result record binds one outcome to its target, receiver projection, intervention boundary, attacker fit, workload and cost object (§3). The 82-family map locates these records in the field; the nine-item audit and six-axis holdout matrix establish their evidence scope. Together they answer one question: *which defense claims does the literature support when a deployment changes?* Transfer boundaries here delimit the scope of an evaluated result and the evidence needed to extend it.

Contributions. The SoK contributes three connected results, each grounded in the complete literature audit:

- **Record-level comparison (§§3–5).** The complete inventory of 741 record pairs separates supported descriptive comparisons from object mismatches and unresolved evidence. Within-family contrasts expose substitutions of primitive, target and cost that paper-level summaries conceal. Cohort accounting and family-deletion checks make the dependence on grouping and record multiplicity inspectable.
- **Configuration-level joint binding (§5).** The 351 record–item cells identify which protection, utility and cost outcomes describe the same evaluated intervention. CachePrune's two records supply one inference–utility–cost package at a shared configuration. Deduplicating that package keeps record counts, family coverage and jointly supported operating points distinct.
- **Deployment decisions (§§6–7).** The 82-family, six-axis holdout map identifies which deployment changes the sources evaluate. Table 11 combines that map with comparison and joint-binding judgments to specify which evidence can be retained and which measurement a proposed extension requires.

WF practice establishes adaptive attacks, distribution shifts and online-cost evaluation [7], [8]. The SoK carries this discipline across LLM receiving views, stateful tool interventions and configuration-linked utility/cost evidence. Table 1 connects each research question to the complete audit that answers it.

Scope of the synthesis. Counts describe source-bound evidence in the indexed collection. Comparison and joint-binding judgments are conditioned on the intended reuse query: they identify which measured outcomes can be carried forward and which deployment change requires new evidence.

2. Scope and Systematization Method

2.1. Research questions, eligibility and threat polarity

We organize the systematization around four questions: which hidden targets, threat polarities and auxiliary observables appear across LLM and agent systems (RQ1); how

TABLE 1. RESEARCH QUESTIONS SUPPORTING THE THREE SOK CONTRIBUTIONS: RECORD-LEVEL COMPARISON, CONFIGURATION-LEVEL JOINT BINDING AND DEPLOYMENT DECISIONS. RQ1 SUPPLIES THEIR COMMON FIELD CONTEXT.

Question	Systematized finding	Contribution and literature basis
RQ1: targets and observables	Signal generation clusters at serving (32 families) and orchestration (19), while receiving views cluster at application, hardware and transport. Receiver and threat polarity determine the protection being compared.	Common context: 82-family index, 58-core-family map and result model (§§3–4).
RQ2: comparable outcomes	Six of the 11 families with multiple records contain a primitive, target or cost mismatch. The full pair inventory separates descriptive matches, mismatches and unresolved bindings; shared mechanism names do not bind execution or state effects.	Record-level comparison: all 741 pairs, with source locators and comparison reasons (§5).
RQ3: transfer axes	Reported task/input holdouts (13 families) and model holdouts (three) address different changes. Hardware/time entries are mostly unresolved, identifying the source scope that must be established before reuse.	Deployment decisions: 82-family, six-axis matrix and claim-conditioned decision table (§6, Table 11).
RQ4: joint evidence	CachePrune supplies one configuration-linked inference–utility–cost package through two records. Separate positive family cells do not establish a joint operating point.	Configuration-level joint binding: 351 record–item cells and configuration-linked joins (§5).

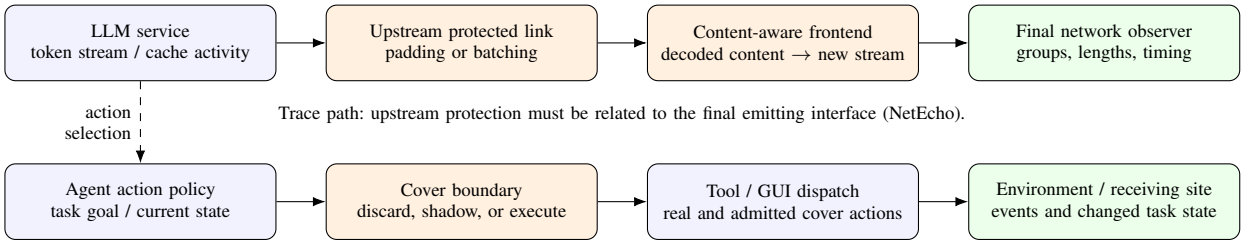


Figure 1. Receiving views and execution boundaries in the literature. A content-aware frontend can create a new trace, while admitted cover actions can change the protected state; each requires its own evidence binding.

signal generation origin relates to receiver observation permissions (RQ2); which holdouts establish portability across tasks, configurations and hardware (RQ3); and which evidence binds workflow protection, logical state preservation and resource cost, distinguishing these from full end-to-end planner utility (RQ4).

Core inclusion requires five conjunctive criteria: (C1) LLM or agent execution, (C2) a hidden target property, (C3) an execution-dependent auxiliary observable rather than ordinary output semantics alone, (C4) specified observer permissions, and (C5) an empirical inference evaluation. Mixed-feature papers qualify only if they isolate an auxiliary-only C3 component; thus coding fingerprints, WildAuth, ZeroPrint, endpoint-content fingerprints and StepFinder remain adjacent boundary comparators [9], [10], [11], [12], [13]. Table 2 states the observable eligibility rules.

We use four narrative groupings of objectives and observer roles, not mutually exclusive classes. An unauthorized identity inference can be both attribution and a confidentiality concern. Each record separately specifies the inference target, observer, authorization relation and protected interest: **Type I** concerns confidentiality/privacy; **Type II** concerns model, harness or identity attribution; **Type III** concerns evasion of a platform's human–agent detector; **Type IV** concerns authorized operator diagnostics. Types can overlap. WRA protects research intent (I); UI Traces attributes agents (II); AHB evades a defensive detector (III); ClawGuard diagnoses loops (IV) [6], [5], [14], [15]. These roles are mapped per core family, alongside generating and receiving layers, in Table 6.

TABLE 2. OBSERVABLE-BASED ELIGIBILITY. A HIDDEN TARGET ALONE DOES NOT MAKE A SEMANTIC ATTACK A SIDE CHANNEL.

Observed input	Decision and required distinction
Length, timing, cache response	Core when measured execution dependence supports hidden-target inference.
Own-site actions	Core attribution with site permission; no automatic access to other sites or private tool logs.
Reply states the secret	Direct disclosure; adjacent, not auxiliary-channel evidence.
Reply style/content identifies a user	Adjacent unless a separately evaluated auxiliary mechanism is established. WildAuth stays a mixed-view comparator.
Output-filter membership oracle	Can be core: the filter's execution, not ordinary semantic output inference, mediates the signal.
Privileged loop/fault logs	Separate diagnostic component when C1–C5 hold; does not increase confidentiality evidence.

2.2. Corpus construction and reconciliation

Forty seeds, venue-targeted proceedings screening and one-hop citation snowballing produced an initial map of 60

families (43 core, 17 adjacent). A supplementary database and bulk keyword search returned 615 occurrences and 565 unique records; auditing its 204 exclusions rerouted four records, and the search added 15 core and seven adjacent families. The final collection has 58 core and 24 adjacent families. The completed search frames, closed on October 21, 2026, combine Semantic Scholar and native arXiv queries with venue-targeted screening of S&P, USENIX Security, NDSS, CCS and PMLR. C1–C5 screening and version deduplication operate on records before family reconciliation; the supplement retains queries, dispositions and excluded incomplete search frames.

2.3. Source-bound result records

Each result record binds its target, observation boundary, operating point and source section/table/figure/paragraph locator. The original records and family audit entries were author-checked against these sources. The record-local extension reuses these bindings and documents targeted source-follow-up changes. The pair inventory fixes the mechanism-question cohorts; joint queries start from the same 33 mitigation records and publish their eligibility masks. S33 reports status-rule, unresolved-binding and family-deletion sensitivity without changing archived labels. Appendix B indexes all 39 records with their measurement modes; Appendix E describes the supporting artifact.

SoK Takeaway 1 (Map evidence to the receiving deployment). A family-level label does not identify which signal, observer permission or protected objective a result supports. The complete core map separates signal origin from receiving access; the result ledger then binds each comparison to a source version and operating point.

3. Result Records and Comparability Rules

The atomic unit of our systematization is one *evaluated result record*, rather than a whole paper or a single headline score. A result record binds a source version and locator to seven structured fields: Source, Target (and threat polarity Types I–IV), Receiver projection, Intervention boundary, Attacker training/calibration (D_0: clean-condition fit; D_2: protected-condition fit), Workload/holdout, and Outcomes (inference metric, task/component utility, and cost object; Table 3). A study family contributes multiple records whenever it evaluates distinct execution modes (e.g., stored capture vs. live loopback in Dangol [16]), distinct mechanisms (e.g., K-prefix KV simulation vs. GPTCache semantic de-identification in Early Bird [3]), or separately scoped performance benchmarks (e.g., SafeKV [17] and OTRO [18]).

Indexing schema and comparison contracts. The mnemonic SCOPE (Secret/target, Channel, Observer, Pipeline, Exploitation goal) names the target-and-channel columns of our index, paired with evaluation provenance. Generating layers span hardware (P1), serving (P2), transport (P3), application/API (P4), orchestration (P5), tools/MCP (P6), environment (P7) and multi-agent communication (P8). We apply three comparison contracts (Table 4): a comparison is *compatible* when required fields align; *incompatible* when a positive mechanism or boundary conflict is established (such as substituting token-payload ratios for wire bytes or offline timestamp edits for live action delays); and *unresolved* when a required binding is unreported (U), unavailable (A) or unverified (?). Table 5 illustrates a complete record-to-judgment lineage.

From contracts to an audit. Section 5 operationalizes these contracts as nine reuse items with a shared evidence-status alphabet (Table 8) and applies them to every mitigation family (Table 9).

TABLE 3. A SINGLE RESULT RECORD. ONLY FIELDS RELEVANT TO THE REQUESTED COMPARISON MUST ALIGN; UNAVAILABLE REQUIRED FIELDS REMAIN UNRESOLVED.

Record field	Required meaning
Source	Family/result ID, version, source locator, reading scope and hash.
Target	Hidden property, inference goal and protection relation: confidentiality, attribution, human/agent distinction or authorized diagnosis.
Receiver	Actual permission and released representation, including session/document linkage and any auxiliary inputs.
Intervention	Where protection ends; whether edits affect stored traces, transmission or executed actions.
Attacker	Fit/calibration condition, test population, training access and adaptation budget.
Workload	Task, system, topology/hardware and grouping/holdout.
Outcomes	Inference metric / denominator; objective or proxy utility; applicable cost object; uncertainty and field states.

TABLE 4. COMPARISON RULES. A KNOWN MISMATCH AND A MISSING BINDING HAVE DIFFERENT CONSEQUENCES.

Requirement	Condition and resulting license
Descriptive comparison	Align target, receiver permission, metric meaning and cost object. Different systems may be described together; token payload and wire bytes cannot be ranked as the same cost.
Paired effect	Compare the same defended test units and outcome. A known test-unit mismatch makes that subtraction incompatible; an unverified fit binding leaves the effect unresolved.
Deployment transfer	Establish a bridge for the changed receiver, intervention, workload, topology or hardware. A demonstrated violation blocks the proposed transfer; an untested bridge leaves it unresolved. Permit only the range and endpoints actually covered.

4. Corpus Map Across 82 Study Families

4.1. Generating layers, receiving projections and threat polarities

Table 6 maps all 58 core families by signal-generating layers, receiving projections and threat polarity. Multi-layer families retain every reported layer in their generating and observing sets.

TABLE 5. FILLED RECORD AND ITS LICENSED JUDGMENT. VALUES ARE SOURCE-REPORTED; NO CROSS-MODE EFFECT IS COMPUTED.

Join key / field	Entry
Family / result	Dangol workflow / `DANGOL-livebinding`.
Version / locator	arXiv:2606.07150v3, §9.7, live-protocol paragraph after Table 5 [16].
Target / view	Workflow classification from live Oblivious HTTP (OHTTP)/Agent-to-Agent (A2A) traffic; not GUI-state observation.
Intervention	Live fixed-rate cover with fresh per-session identifiers, distinct from transformed stored captures.
Workload / testbed	Workflow classes with client, relay, gateway and agent co-located on one loopback host.
Outcome / cost	Accuracy 0.30; about $13\times$ bytes and $11\times$ latency in that condition.
Unresolved fields	Objective terminal task utility U; precise paired live-refit comparison ?.
Licensed judgment	These costs describe this live protocol condition. They neither price dummy GUI actions nor prove task preservation. A matched objective endpoint is needed for the latter.

Cross-layer synthesis. A remote network observer, a same-site script and a privileged diagnostic hook expose different views even when they share an attribution target. WRA's read-only privacy decoys, AHB's stateful evasion and SplitDecoy's exposed return path illustrate why receiving permissions and protected objectives must accompany the generating layer [6], [14], [34].

Layer incidence highlights where the systematization's comparison problem arises. In the 58-family map, signal-generation labels cluster at P2 serving/runtime (32 families) and P5 orchestration (19); receiving views cluster at P4 application/API (18), P1 hardware (16), and P3 transport (14). Generation at P3/P4 is rare (one family each), despite frequent observation at those layers. This source/receiver asymmetry makes the transfer question concrete: protection at a frequent source layer cannot be presumed to cover a different receiving projection. The map records layer sets, so it does not assign a causal edge to each source/receiver pair.

Purpose labels further show why a single "agent side channel" category is insufficient: Type I appears 41 times, Type II 13, Type III 2, and Type IV 5. The 32 map-only families extend coverage beyond mitigation comparisons: 15 list serving origins, 11 hardware and 11 orchestration; 19 concern confidentiality, nine attribution, one evasion and all five authorized-diagnostic families. Mapping these families locates privacy targets, operator capabilities and candidate validation environments even when no scoped mitigation record enters the ledger. Counts are overlapping family-label incidences; the audit then applies its reuse rules to the 39 scoped result records.

4.2. Relationship to prior systematizations

Prior surveys map compound-system threats, TEE boundaries and privacy contracts [66], [67], [68]. WF already supplies feature leakage analysis, cost trade-offs and environment-shift tests [69], [7], [70], [71]. Mathews et al. re-evaluate nine defenses with stronger attacks and an on-network DynaFlow implementation [8]. Adaptive attacks and online-cost evaluation are established foundations.

What this synthesis adds. Our unit is a *scoped result*, so execution mode, protected state and cost remain attached when one paper reports several mechanisms. The 39-record census exposes these within-paper substitutions; the 82-family holdout matrix separates evidence for each transfer axis. This connects evidence that earlier taxonomies organize by attack or mechanism: one family can contain several execution modes, a receiving projection can cross system layers, and utility/cost can belong to a different operating point from inference. The synthesis preserves these distinctions when translating a reported result into a deployment decision.

5. Comparability Audit Across All 58 Core Families

5.1. Unified 58-core disposition and 39-record ledger

The corpus partitions as $82 = 58$ core $+24$ adjacent families. Within the core, $58 = 20$ mitigation $+6$ contextual $+32$ map-only families. The ledger holds 39 records: 33 from the 20 mitigation families and one from each contextual family—SIGLEAK, AgentTell, FP-Agent, Behind Bars, LLMs Have Rhythm and ML privacy channels [57], [64], [65], [23], [36], [44]—which supply fixed-guardrail, access-control, interference, VPN or filter-origin contrasts rather than an audited mitigation claim. The 32 map-only families comprise 22 without a scoped mitigation evaluation, five diagnostics, four unbound quantitative assertions and SplitDecoy's boundary case. Table 7 groups mechanisms; its eight rows are mechanism groups, not sampled studies.

Whole-ledger synthesis. Each record has one primary measurement mode: 20 executed collections/prototypes, six trace transformations, five simulations/derived estimates, six component/calibration measurements and two unresolved placements (Table 12). Eleven of the 20 mitigation families contribute more than one record, and nine of them span more than one mode; this is why a mechanism name cannot carry utility or cost between records. Executed collection establishes its stated execution scope, not automatically full-workflow utility.

5.2. Within-family and within-cohort record comparisons

The completed audit covers all $\binom{39}{2} = 741$ pairs. The 608 pairs outside the declared mechanism-question cohorts are not comparison failures. The 133 related pairs comprise 66 traffic-envelope, 45 cache-service, 15 local-memory, three behavioral-humanization, three generator-fingerprint and one fixed-instruction pair; they include all 15 same-family pairs. Three findings follow:

TABLE 6. COMPLETE MAP OF THE 58 CORE STUDY FAMILIES. EACH FAMILY APPEARS ONCE; LAYER AND TYPE LABELS MAY BE MULTIPLE AND DO NOT IMPLY ONE-TO-ONE CAUSAL PATHS. THE BLOCKS GROUP SOURCE LAYERS, NOT EVIDENTIAL QUALITY.

Core family	$G \to O$	Type
hardware generation, P1 (8)		
Edge EM leakage [19]	$1 \to 1$	I
Energon [20]	$1 \to 1$	II
MoEcho [21]	$1 \to 1$	I
I Know What You Said [22]	$1 \to 1$	I
Behind Bars [23]	$1 \to 1$	I
PAGE404 [24]	$1 \to 1$	I
Spill the Beans [25]	$1 \to 1$	I
ArchSnoop [26]	$1 \to 1$	II
hardware and service, P1/P2 (8)		
SparLeak [27]	$1,2 \to 1$	I
Shape and Substance [28]	$1,2 \to 1$	I
TDXRay [29]	$1,2 \to 1$	I
SparSEEty [30]	$1,2 \to 1$	I
OTRO [18]	$1,2 \to 1$	I
Detokenization Leaks [31]	$1,2 \to 1$	I
LLM-WFIN [32]	$1,2 \to 1,3$	I
LLM-serving RCA [33]	$1,2 \to 2$	IV
service generation, P2 (13 of 22)		
SplitDecoy [34]	$2 \to 2$	I
KVGov [35]	$2 \to 2,4$	I
LLMs Have Rhythm [36]	$2 \to 3$	II
Remote timing [37]	$2 \to 3$	I
NetEcho [38]	$2 \to 3$	I
Prompt Echo [39]	$2 \to 3$	I
Speculative decoding [40]	$2 \to 3$	I
Remote keylogging [1]	$2 \to 3$	I
Whisper Leak [41]	$2 \to 3$	I
Cache contention [42]	$2 \to 4$	I
Routing metadata [43]	$2 \to 4$	I
SafeKV [17]	$2 \to 4$	I
ML privacy channels [44]	$2 \to 4$	I
service generation, P2 (9)		
Early Bird [3]	$2 \to 4$	I
Prompt-cache audit [45]	$2 \to 4$	I
InputSnatch [46]	$2 \to 4$	I
Leaky Language Models [47]	$2 \to 4$	I
PromptPeek [2]	$2 \to 4$	I
SpliceLeak [48]	$2 \to 4$	I
Time Will Tell [49]	$2 \to 4$	I
CachePrune [50]	$2 \to 4$	I
MoE prompt stealing [51]	$2 \to 4$	I
cross-layer sources (5)		
ClawGuard [15]	$1,5 \to 1$	IV
MIRAGE [52]	$2,5 \to 2$	IV
Cloud traffic benchmark [53]	$2,5,8 \to 3$	I, II
Web-agent unmasking [54]	$3,7 \to 3,7$	II, III
LumiMAS [55]	$4,5 \to 4,6$	IV
orchestration/workflow sources (15)		
AgentPrint [4]	$5 \to 3$	I, II
Cycle detection [56]	$5 \to 4,6$	IV
SIGLEAK [57]	$5,6 \to 4,6$	I
ProcGrep [58]	$5,6 \to 4,6$	II
AgentProv [59]	$5,6 \to 6$	II
Lidar [60]	$5,6 \to 6$	II
Dangol workflow [16]	$5,6,8 \to 3,6,8$	I
Campaign attribution [61]	$5,6,8 \to 6,8$	II
WRA [6]	$5,7 \to 3$	I
PersonaFingerprint [62]	$5,7 \to 3$	I
Whose Agent? [63]	$5,7 \to 7$	II
AgentTell [64]	$5,7 \to 7$	I
UI Traces [5]	$5,7 \to 7$	II
FP-Agent [65]	$5,7 \to 7$	II
AHB [14]	$5,7 \to 7$	III

G denotes signal-generation layers and O the receiving observation layers. Layer indices: 1 hardware; 2 serving/runtime; 3 transport; 4 application/API; 5 orchestration; 6 tools; 7 browser/environment; 8 multi-agent communication. Types: I confidentiality/privacy; II attribution; III human–agent distinction/evasion; IV authorized diagnostics. These are objectives and observer roles, not mutually exclusive security guarantees. Privileged operator receivers and physical receivers are marked by the layer they inspect. Source versions, locators and normalization of prose-only layer descriptions are in the artifact.

- **Same paper, different object.** Eight of the 15 same-family pairs substitute primitives, targets or cost objects, distributed across six of the 11 families with multiple records. CachePrune's direct recovery and contextual guessing differ despite sharing a defense configuration; OTRO's measured unpadded latency does not measure its epoch-padded operating point.
- **Aligned outcome, unpaired population.** Of the other seven same-family pairs, four have different collection/execution populations and three lack a verified matched-test binding. Dangol's stored-capture ablation and live loopback remain distinct deployment measurements.
- **Across records.** Of 133 related pairs, 17 support a source-qualified descriptive comparison, 105 have an established object/receiver/target mismatch, and 11 remain unresolved. Descriptive compatibility permits scoped comparison, not pooled effects or a ranking across heterogeneous costs.

The final inventory retains all 528 earlier pair judgments and adds 31 related judgments for the six expansion records, with source locators and reasons. Examples illustrate this census rather than define its sample. The 15 same-family pairs share records and families. Removing KVGov leaves five mismatches among 12 pairs; removing Whisper Leak leaves eight among 12. Across all 26 family deletions, related-pair compatibility ranges from 7.8% to 15.2%. Resolving all 11 uncertain pair judgments favorably would increase the compatible count from 17 to 28, with the 105 known mismatches unchanged. S33 reports each cohort and deletion. These are inventory sensitivities, not independent-sample prevalence estimates.

Why the record changes the reuse decision. The census distinguishes three outcomes for a deployment query. A known object mismatch requires a new measurement at the intended operating point; an unresolved binding requires source reconciliation; a descriptive match permits a scoped comparison while retaining each cost's denominator. These

TABLE 7. EIGHT INTERVENTION FAMILIES ORGANIZED BY THREAT POLARITY (TYPE I: CONFIDENTIALITY PROTECTION; TYPE II: PROVENANCE/IDENTITY ATTRIBUTION; TYPE III: BOT-DETECTION EVASION), RECEIVING VIEW AND SUPPORTED DEPLOYMENT JUDGMENT. BECAUSE OBJECTIVES AND RISK MAGNITUDES DIFFER, THESE ROWS ARE NOT A COMMON EFFICACY RANKING.

Intervention & polarity	Evaluated view, target and source	Established result and applicable boundary
Padding / batching (Type I: confidentiality)	Network packet lengths and groups; topic inference. NetEcho and Whisper Leak [38], [41].	Collected padding traces retain inference signal; packet-group batching is evaluated by replay. A protected upstream length view does not cover downstream re-emission or residual grouping. Live incremental cost remains U.
Constant-rate / transport cover (Type I: confidentiality)	Network workflow fingerprint; Dangol's OHTTP testbed [16].	The co-located loopback setting yields 0.30 recovery at about $13\times$ bytes and $11\times$ latency. Transformed-capture results are separate; neither result supplies objective agent-task utility.
Concurrent query decoys (Type I: confidentiality)	Destination/domain observer; private prompt and trait inference. WRA [6].	Real/decoy research prompts execute in parallel as read-only queries whose decoy answers are discarded locally; interleaved-domain evaluation reduces one-decoy E_{func} from 0.770 to 0.610. Terminal decoy utility/cost are U.
Action humanization (Type III: bot-evasion)	Mobile GUI touch traces; evading human–agent bot detection. AHB [14].	Unlike WRA's read-only privacy decoys, AHB's `Fake` mode dispatches live mobile GUI touches to evade platform bot detectors, dropping Trip Planning success from 0.75 to 0.15. Non-Fake `History+Long` (0.7143 success) avoids GUI mutation.
Cache isolation / keying (Type I: confidentiality)	Co-tenant prefix timing; SafeKV, KVGov and CachePrune [17], [35], [50].	SafeKV reports post-trigger AUC near 0.5 and separate TTFT/throughput comparisons. KVGov's HMAC median 1.3 μs prices key computation, excluding lost reuse. CachePrune jointly binds marked-token protection, QA/summarization utility and serving performance.
Cache timing equalization (Type I: confidentiality)	Co-tenant cache-fusion timing; SpliceDefense [48].	Constant-Time Boundary Fusion (CTBF) suppresses the tested signal over 15 tasks/75 steps. Quantized Chunk Padding (QCP) adds mean 7.5 tokens and < 5ms/$< 1\%$ prefill time; this does not price combined deployment.
Timestamp perturbation (Type II: attribution)	Recorded website action metadata; model/agent attribution. UI Traces [5].	Fitting on delayed traces (D_2) recovers substantial attribution performance under offline timestamp edits that retain action types/order. Live execution delay and task outcomes are not measured by the trace transformation.
Oblivious / unshared memory (Type I: confidentiality)	Host or local resource observations; TDXRay, Gao et al. and OTRO [29], [22], [18].	TDXRay measures oblivious-map lookup cost; Gao measures loading/memory cost after disabling zero-copy. OTRO reduces conditional access-count information (3.94 to 0.04 bits) with separate unpadded TTFT (4.59%) and epoch-padding (2.8ms) bounds.

Authorized operator diagnostics (Type IV: ClawGuard, cycle detection, LumiMAS, LLM-serving RCA and MIRAGE) are baseline telemetry rather than privacy attacks or evasion defenses, and are cataloged separately.

outcomes prescribe different work. For example, aligning two packet-size accuracies cannot supply the missing online execution cost, and measuring a component's throughput cannot establish a planner's success under response deadlines. The 17 descriptive matches are useful precisely because their permissions remain explicit.

5.3. Nine-item reuse audit across 20 mitigation families

R1–R9 are author-defined reuse conditions, not a quality score. Table 8 links requested conclusions to necessary evidence. Missing evidence can block an analyst's deployment extension while leaving the source's own claim intact. Component utility suffices for component claims; R5 applies to adaptation comparisons. U and ? retain distinct source states: the 11 unverified cells are five R4 and six R5 bindings, which are distinguished from an attack failure.

Join evidence at the record, not the family. A family-level B means that at least one record supplies an item. It does not imply that its other B cells concern the same intervention. We therefore code all 351 record–item cells and join only B bindings within the same record and configuration (Table 10). All joint queries begin with 33 mitigation records. A query excludes a record only when a required item is inapplicable; U, ?, and P remain eligible but unbound. Table 10 separates excluded, eligible and bound counts, with deduplicated families/configurations. The six contextual controls are separate. CachePrune supplies the triple join: two outcomes share one quality/cost package. The denominator identifies eligible records for our query, not papers obliged to answer it. SparLeak's matched ASR/ROUGE pairs qualify as component utility under R8, correcting its former partial code.

Which evidence packages are available. Record-local joins resolve progressively stronger deployment queries: inference with utility is bound in 4 of 22 eligible records (three families/configurations), inference with cost in 5 of 22 (three families, four configurations), and matched adaptation with inference in 2 of 17. Their intersection cannot be reconstructed by adding these counts: only CachePrune's two outcomes share one inference–utility–cost package. A deployment analyst can therefore reuse the bound component outcome, price a missing cost, or measure workflow utility according to the actual join. This is more informative than treating every family-level B as a complete deployment result. Accepting P and ? as if bound leaves the triple join at two records, but changes matched adaptation–inference from two to 16 of 17. Thus uncertainty concentrates differently by query. Removing CachePrune removes the sole triple package; no other deletion does. These checks locate which additional source bindings would change a decision, rather

TABLE 8. CLAIM-CONDITIONED REUSE RULES. EACH ROW NAMES THE REQUESTED INFERENCE AND THE EVIDENCE NEEDED TO SUPPORT IT.

Item	Reuse query → required binding (**B**)
R1 Target	Same protected interest → explicit hidden property, inference goal and authorization relation.
R2 Receiver	Same observer protection → permissions, released view and auxiliary inputs.
R3 Intervention	Same executed mechanism → placement/mode: stored data, wire emission or actions.
R4 Protected fit	Resistance to protected-data fitting → identified protected training/calibration and testing.
R5 Matched adaptation	Effect of adaptation → clean/protected fits tested on the same defended population.
R6 Workload	Changed workload/system → tasks, configuration and the relevant holdout axis.
R7 Inference	Quantified protection → metric, denominator and scoped defended evaluation condition.
R8 Utility	Preserved task/component function → performance at the same protected state/operating point.
R9 Cost	Resource-budget feasibility → units, baseline and cost object at that operating point.

B: bound; **P**: partial binding; U: unreported in inspected scope; A: source unavailable; ?: binding unverified; –: inapplicable. First determine applicability, then assign evidence status. Family **B** means at least one scoped record satisfies the item, not that every result does.

than turn missing reports into evidence of failure.

SoK Takeaway 2 (Join the operating point). CachePrune's two outcomes identify one reusable inference–utility–cost package. Request evidence for the intended deployment query rather than combine a family's best individual cells. These counts inventory evidence; they do not rank paper quality.

6. Systematic Holdout Evidence for Transfer

6.1. Six axes across all 82 families

Appendix D gives the complete family-by-axis matrix, including the 24 adjacent families. It codes task/input/user, model, harness/topology, hardware, network/load and time separately from the archived source-bound extraction. H denotes an explicitly reported held-out unit/domain; V a condition comparison or recalibration; X a scoped restriction; ? unresolved extraction. H does not certify every grouping detail or successful transfer, and ? is never counted as absence. Conceptual families receive an inapplicable dash.

The matrix contains H entries for 13 families on task/input/user, three on model, two on harness/topology and one on network; none is resolved as H for hardware or time. The latter axes contain seven V entries each: calibration, condition sweeps and temporal observations already exist in the literature. With 73 hardware and 71 temporal entries unresolved, this is a map of *documented evidence*, not a claim that those papers omitted such tests. Multiple H entries are counted once per family/axis, not pooled as independent experiments.

TABLE 9. PER-FAMILY EVIDENCE BINDINGS FOR ALL 20 EVALUATED-MITIGATION FAMILIES. R1–R9 ARE SCOPED REUSE ITEMS, NOT OVERALL PAPER SCORES.

Family	R1	R2	R3	R4	R5	R6	R7	R8	R9
UI Traces [5]	**B**	**P**	**B**	**B**	**B**	**B**	**B**	U	–
WRA [6]	**B**	**P**	**B**	–	–	**P**	**B**	**P**	U
NetEcho [38]	**B**	**B**	**B**	**B**	**P**	**B**	**B**	U	U
Early Bird [3]	**B**	**B**	**B**	**B**	**P**	**B**	**B**	U	**B**
Whisper Leak [41]	**B**	**B**	**P**	?	?	**B**	**B**	U	**B**
SpliceLeak [48]	**B**	**B**	**B**	–	–	**B**	**B**	U	**B**
Time Will Tell [49]	**B**	**B**	**B**	?	?	**B**	**B**	U	**B**
I Know What You Said [22]	**B**	**B**	**B**	–	–	**B**	U	–	**B**
Speculative decoding [40]	**B**	**B**	**P**	?	?	**B**	**B**	–	**B**
Dangol workflow [16]	**B**	**B**	**B**	**B**	**B**	**B**	**B**	U	**B**
Cloud traffic benchmark [53]	**B**	**P**	**B**	?	?	**P**	**B**	**P**	U
Remote timing [37]	**B**	**B**	**B**	–	–	**B**	**B**	U	**P**
TDXRay [29]	**B**	**B**	**B**	–	–	**B**	U	**B**	**B**
SafeKV [17]	**B**	**B**	**B**	–	–	**B**	**B**	U	**B**
AHB [14]	**B**	**B**	**B**	**P**	?	**B**	**B**	**B**	U
KVGov [35]	**B**	**B**	**B**	–	–	**B**	**B**	U	**B**
CachePrune [50]	**B**	**B**	**B**	–	–	**B**	**B**	**B**	**B**
OTRO [18]	**B**	**B**	**B**	–	–	**B**	**B**	–	**P**
SparLeak [27]	**B**	**B**	**B**	?	?	**B**	**B**	**B**	U
Shape and Substance [28]	**B**	**B**	**P**	U	U	**B**	**P**	U	**P**
Bound	20	17	17	4	2	18	17	4	11
Applicable	20	20	20	11	11	20	20	17	19

B: bound; **P**: partial; U: unreported in inspected scope; ?: unverified binding (11 cells); –: inapplicable. Table 8 defines R1–R9 and the shared status alphabet. A bound cell is supported by at least one scoped record. Dangol R5 concerns stored capture; its live-refit binding remains unverified. CachePrune jointly binds R7–R9 at one operating point.

TABLE 10. QUERY ACCOUNTING FROM THE SAME 33 MITIGATION RECORDS. OUT: AT LEAST ONE REQUIRED ITEM IS INAPPLICABLE; ELIGIBLE: REMAINING RECORDS; BOUND: ALL REQUIRED ITEMS ARE B.

Query	Out	Eligible	Bound	Fam./Cfg.
Inference + utility	11	22	4	3/3
Inference + cost	11	22	5	3/4
Matched adaptation + inference	16	17	2	2/2
Inference + utility + cost	13	20	2	1/1
All five bindings	19	14	0	0/0

Required items by row: R7/R8; R7/R9; R4/R5/R7; R7/R8/R9; all five. The two sets of 22 eligible records differ. Six contextual controls are outside the starting 33. Fam./Cfg. deduplicates bound records. CachePrune supplies the triple join: two outcomes at one configuration. Joins preserve cost scope and do not imply paired effects. S33 lists every exclusion and status.

6.2. What each holdout licenses

Whisper Leak holds out target prompts; FP-Agent has a separate task holdout but unresolved visitor/host separation. Dangol's leave-composition-out and the traffic benchmark's topology shift change a different generator [41], [65], [16], [53]. LLM-Fin and ArchSnoop test unseen variants/configurations within known model families; LeakyLMs conditions architecture inference on known GPU type and calibration [72], [26], [47]. Thus a positive model holdout does not establish unknown-hardware transfer. UI Traces' recorded question overlap further shows why task labels and actual held-out units must remain distinct [5].

6.3. Operational consequence

Select the axis changed by deployment, require a source-bound separation on that axis, then check whether fitting or calibration crossed the split. Otherwise scope the result to its observed condition. Feature ablations and additional models alone supply V evidence, not a bridge to an unseen generator. Within-model refitting, source-trained attacker transfer and source-calibration transfer answer different questions and therefore require separate evidence bindings. The systematic result is the mismatch between the requested transfer axis and the available evidence, not a ranking by the number of holdouts.

From holdout counts to a transfer obligation. Each axis names a distinct source of variation. A task holdout tests new inputs under its stated generator; a model holdout changes the generator; topology and network holdouts change the receiving path; hardware and time can change the implementation or scheduling process. The 13 task/input H entries and three model H entries therefore answer different questions, rather than form a single robustness score. Where extraction remains unresolved, the next action is to establish the source's scope before demanding a new experiment. Where a deployment changes several axes, evidence must cover their intended combination: separate model and network tests do not automatically establish their joint operating point.

7. Synthesis: From Evidence to Reuse Decisions

The census (§5) covers all 39 records and the holdout matrix (§6) all 82 families. We now select three examples by the *comparison decision* they expose, not by effect size: a changed receiving path, cover with different state effects, and an adaptation/cost mismatch within a named mechanism. Each exposes a different reason why a deployment conclusion cannot be read directly from a mechanism name.

7.1. Case 1: bind the final target-bearing emitter

NetEcho's upstream packet-size shaping operates on a different projection from the tool-call sequence. Its plaintext-aware frontend then emits a new size trace (grouped token counts or character lengths), reopening a size channel after upstream protection [38]. SplitDecoy is a separate return-path record: the cloud sees nonzero gradient row support on every tested frame across nine seeds [34, Sections 5.1–5.3]. Both change the receiver boundary and therefore require a separately bound inference result.

The data processing inequality requires $S \to Z \to Y$: target-oblivious postprocessing cannot increase $I(S;Z)$ [73, Theorem 2.8.1]. Plaintext-aware re-emission or backward computation breaks that premise. The literature comparison must therefore bind both the final receiving projection and the emitter whose costs are measured.

7.2. Case 2: resolve the WRA versus AHB cover asymmetry

WRA's one-decoy E_{func} drops from 0.770 to 0.610, whereas AHB's Trip Planning success falls from 0.75 to 0.15 under `Online+Fake` humanization [6], [14]. Grouping both as "cover traffic" conceals different objectives and executions.

WRA discards decoy reads for Type I confidentiality; AHB executes GUI events for Type III evasion, potentially navigating or submitting forms. Their utility outcomes are not contradictory, and AHB's missing condition-specific denominators preclude a matched causal effect. A shared workflow setting does not establish shared state effects. Terminal-state reward and dispatch-time authorization also answer different questions: an unconfirmed write can still reach the intended final state. Reuse requires an explicit state/authorization outcome at the intervention being carried forward.

7.3. Case 3: bind adaptive retraining and cost objects

A matched adaptation contrast compares clean-trained D_0 and defense-trained D_2 on identical defended tests. UI Traces and Dangol raise the record-binding question this leaves open: whether a refit, a measured delay and a utility outcome refer to the same executed condition [5], [16]. Cost objects show the same pattern. Early Bird's 6ms/4% GPTCache cost does not price K-prefix isolation; KVGov's $1.3\,\mu$s HMAC microcost omits lost reuse; SpliceDefense's QCP cost does not price CTBF+QCP [3], [35], [48]. CachePrune binds protection, QA/summarization utility and throughput, while OTRO separates conditional leakage from TTFT/padding costs [50], [18]. The reuse decision follows the evaluated configuration: retain each source's bound outcome, and request the missing adaptation, execution or cost measurement before extending it.

7.4. A deployment query determines the evidence to retain

Start with the intended receiver, protected outcome and intervention. Record-level comparison identifies whether the cited result measures that object. Configuration-level joint binding identifies which inference, utility and cost outcomes share its configuration. Deployment decisions use the holdout matrix to identify evidence for the proposed change and specify the remaining measurement. Table 11 makes these outputs operational: a new emitter needs a receiver-specific capture, writes need state/authorization checks, source-calibration comparisons need a local fitting-variability control, and an online scheduler needs its own planner utility and latency measurement. The calibration control follows from the comparison design: changing the source can also change the finite training sample used to fit the defense.

TABLE 11. CLAIM-CONDITIONED REUSE DECISIONS DERIVED FROM THE RECORD AND HOLDOUT AUDITS. EACH CHANGED BOUNDARY IDENTIFIES THE EVIDENCE THAT MUST BE BOUND BEFORE CARRYING A RESULT FORWARD.

Changed boundary	What the old result leaves open	Minimum evidence before reuse
Emitter / representation	An upstream or buffered capture can differ from the final receiving trace and delivery schedule.	Bind inference to the final receiver; retain the emitter's overflow, byte and delivery-delay scope (R2/R3/R9; §7).
Read cover → writes	Discarding a reply does not establish unchanged commits, authorization or workflow outcomes.	Bind the protected state and utility to the executed intervention; include cover calls, quotas and failures (R3/R8/R9).
Customer / model calibration	A reported split can separate inputs without changing the generator or calibration population.	State held-out units and calibration provenance; separate local fitting variability from source/domain changes; evaluate the requested transfer axis (R5/R6; §6).
Attacker observation	An outcome for one projection or fitted attacker does not supply a result for a richer receiving view.	Evaluate comparable defended tests under the intended observation and adaptation contract; retain receiver-specific inference metrics (R2/R5/R7).
Stored trace → online path	A trace transformation or simulated span does not measure executed queues, completion or online latency.	Bind the receiver, harness and resource denominator to the online condition; establish the relevant topology/network holdout (R3/R6/R9).
Component → workflow utility	Component quality or logical state conformance does not determine a planner's response to scheduling and deadlines.	Report utility at the requested workflow boundary and the same intervention/configuration as protection and cost (R3/R8/R9).

Appendix E applies this procedure in one compact case study. The literature audit and the decision table establish the SoK contribution; the case study provides a worked application with separately scoped measurements.

8. Implications, Ethics and Limits

8.1. Deployment and reporting implications

The decision table distinguishes receiving permissions from intervention capabilities. An access-network operator, VPC monitor and receiving tool service expose different views; remote paths can reveal destinations even with Encrypted ClientHello [74]. A result for one ingress projection therefore needs a separately bound claim before it is applied to host instrumentation or downstream commit logs.

Likewise, read queries, provider-isolated execution and gateway-routed cover have different state contracts. Kubernetes dry-run, Stripe sandboxes and transactional rollback offer isolation mechanisms with their own limits [75], [76], [77]; idempotence alone is insufficient, and sequence changes can survive rollback [78]. Evidence reuse requires the trusted boundary, authorization policy, resource quotas and protected outcome to accompany the mechanism name. The corpus's component outcomes remain useful when their execution and cost scopes match the intended query.

8.2. Ethics and artifact scope

The audit uses source-linked public research records. Stakeholders include tenants, shared-capacity providers and users affected by observable behavior or state changes. Bot-detection evasion and research-intent privacy have different authorization relations [14], [6]. The appendix case study uses cloned public benchmark databases and controlled sockets; model interactions use public tasks and simulated users. No production account is modified.

AI tools assisted candidate-field extraction, implementation, source reconciliation and prose revision; the authors verified the original source-linked corpus and retain responsibility for the synthesis. The artifact retains source locators, coding decisions and executable checks, with case-study observations and protocols separated from the literature records.

8.3. Scope and threats to validity

Audit scope. The 82-family index and 39-record ledger define the collection and its quantitative denominators. R1–R9 specify evidence needed for a stated reuse query. Applicability depends on that query; unreported or unverified bindings do not establish a failure of the source's own claim. The inference–utility–cost join identifies one family/configuration with two recorded outcomes. Its status-rule stability is conditional on leaving unreported evidence unbound; its positive support depends on CachePrune. S33 exposes those dependencies. Executable counts and sensitivity checks do not measure agreement between independently coding researchers.

Source and comparison scope. Versions, locators and source availability can change what a record supports. The audit retains unresolved bindings and separates descriptive compatibility from matched effects or pooled estimates. Holdout codes record the reported axis and unit of separation; combined deployment changes require evidence for their intended combination. These distinctions preserve useful component findings while exposing the additional binding needed for an extension.

Worked application. The compact appendix case study illustrates use of the procedure in one retail workload and a saved-response transport setting. Its sampling, attacker, calibration and utility scopes accompany its results; the full supplementary record retains their uncertainty cal-

culations. The SoK findings above are computed from the literature audit.

9. Conclusion

This SoK organizes defense-evidence reuse into record-level comparison, configuration-level joint binding and deployment decisions. The complete pair audit exposes changes of primitive, target and cost within shared family names. Configuration-linked joins establish which protection, utility and cost outcomes can be interpreted together, retaining the distinction between two records and one shared evidence package. The six-axis holdout synthesis then connects those judgments to the deployment change being proposed. The resulting procedure gives readers three concrete outputs: a supported comparison, a configuration-bound set of outcomes, and a specification of the additional measurement needed for reuse. Published query membership, source-linked reasons and family/status sensitivity make these outputs inspectable before a defense claim is carried into a new deployment.

References

[1] R. Weiss, D. Ayzenshteyn, G. Amit, and Y. Mirsky, "What was your prompt? a remote keylogging attack on AI assistants," in *33rd USENIX Security Symposium*. USENIX Association, 2024, pp. 3367–3384. [Online]. Available: https://www.usenix.org/conference/usenixsecurity24/presentation/weiss

[2] G. Wu, Z. Zhang, Y. Zhang, W. Wang, J. Niu, Y. Wu, and Y. Zhang, "I know what you asked: Prompt leakage via KV-cache sharing in multi-tenant LLM serving," in *Network and Distributed System Security Symposium (NDSS)*, 2025. [Online]. Available: https://www.ndss-symposium.org/ndss-paper/i-know-what-you-asked-prompt-leakage-via-kv-cache-sharing-in-multi-tenant-llm-serving/

[3] L. Song, Z. Pang, W. Wang, Z. Wang, X. Wang, H. Chen, W. Song, Y. Jin, D. Meng, and R. Hou, "The early bird catches the leak: Unveiling timing side channels in LLM serving systems," *IEEE Transactions on Information Forensics and Security*, vol. 20, pp. 11 431–11 446, 2025, consulted arXiv:2409.20002v5; archived access 18 October 2026. [Online]. Available: https://ieeexplore.ieee.org/document/11206380/

[4] Y. Zhang, X. Deng, Z. Gu, Y. Chen, K. Xu, Q. Li, and J. Wu, "Exposing LLM user privacy via traffic fingerprint analysis: A study of privacy risks in LLM agent interactions," *arXiv preprint arXiv:2510.07176*, 2025. [Online]. Available: https://arxiv.org/abs/2510.07176

[5] W. Lugoloobi, S. Marro, J. Magomere, J. Wright, and C. Russell, "Known by their actions: Fingerprinting LLM browser agents via UI traces," *arXiv preprint arXiv:2605.14786*, 2026, consulted arXiv:2605.14786v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2605.14786v1

[6] H. Jeong, M. Teymoorianfard, A. Kumar, A. Houmansadr, and E. Bagdasarian, "Network-level prompt and trait leakage in local research agents," in *35th USENIX Security Symposium*. USENIX Association, 2026, pp. 6485–6504. [Online]. Available: https://www.usenix.org/conference/usenixsecurity26/presentation/jeong

[7] M. Juarez, S. Afroz, G. Acar, C. Diaz, and R. Greenstadt, "A critical evaluation of website fingerprinting attacks," in *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, 2014, pp. 263–274.

[8] N. Mathews, J. K. Holland, S. E. Oh, M. S. Rahman, N. Hopper, and M. Wright, "SoK: A critical evaluation of efficient website fingerprinting defenses," in *IEEE Symposium on Security and Privacy*, 2023.

[9] T. A. Ghaleb, "Fingerprinting AI coding agents on GitHub," *arXiv preprint arXiv:2601.17406v1*, 2026, consulted arXiv:2601.17406v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2601.17406v1

[10] W. Wang, Z. Liu, H. Kou, X. Liu, and J. Shen, "The chatbot knows it's you: Dialogue attribution in unauthenticated human–LLM sessions," in *Proceedings of the ACM Web Conference 2026*, 2026, pp. 8797–8807. [Online]. Available: https://dl.acm.org/doi/10.1145/3774904.3793048

[11] S. Shao, Y. Li, H. Yao, Y. Chen, Y. Yang, and Z. Qin, "Reading between the lines: Towards reliable black-box LLM fingerprinting via zeroth-order gradient estimation," 2026, version 2, 21 January 2026. [Online]. Available: https://arxiv.org/html/2510.06605v2

[12] J. Leshin, M. Shah, I. Timmis, and D. Kang, "Behavioral fingerprints for LLM endpoint stability and identity," 2026, version 1, 19 March 2026. [Online]. Available: https://arxiv.org/html/2603.19022v1

[13] T. Zhu, Y. Wu, W. Jin, Y. Li, and G. Huang, "StepFinder: A temporal semantic framework for failure attribution in multi-agent systems," 2026, version 1, 2 June 2026. [Online]. Available: https://arxiv.org/html/2606.03467v1

[14] J. Zhu, L. Yang, R. Shan, C. Zheng, Z. Zheng, W. Liu, Y. Yu, W. Zhang, and J. Lin, "Turing test on screen: A benchmark for mobile GUI agent humanization," *arXiv preprint arXiv:2604.09574v1*, 2026, consulted arXiv:2604.09574v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/html/2604.09574v1

[15] L. L. Gan, J. Wu, L. Ge, L. Yang, Y. Song, J. Zhang, H. Jin, W. Wang, and G. Xue, "ClawGuard: Out-of-band detection of LLM agent workflow hijacking via EM side channel," *arXiv preprint arXiv:2605.06205v1*, 2026. [Online]. Available: https://arxiv.org/abs/2605.06205v1

[16] B. Dangol, "From privacy to workflow integrity: Communication-graph metadata in autonomous agent interoperability," *arXiv preprint arXiv:2606.07150v3*, 2026, consulted arXiv:2606.07150v3; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2606.07150v3

[17] K. Chu, Z. Lin, D. Xiang, Z. Shen, J. Su, C. Chu, Y. Yang, W. Zhang, W. Wu, and W. Zhang, "Selective KV-cache sharing to mitigate timing side-channels in LLM inference," *arXiv preprint arXiv:2508.08438v2*, 2026, consulted arXiv:2508.08438v2; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2508.08438v2

[18] J. Lee, Y. Wang, R. Rajat, D. Wong, M. Li, and M. Annavaram, "Otro: Oblivious tokenization path with square-root oram," *arXiv preprint arXiv:2606.17358v1*, 2026, consulted 2606.17358v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.17358v1

[19] N. Alhalawani, G. Acar, Z. Liu, and L. Batina, "Physical side-channel attacks on edge LLMs: Extracting output tokens and user prompts," *arXiv preprint arXiv:2610.14215v1*, 2026, consulted 2610.14215v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.14215v1

[20] A. Chaudhuri, S. Shukla, S. Bhattacharya, and D. Mukhopadhyay, "Energon: Unveiling transformers from GPU power and thermal side-channels," in *IEEE/ACM International Conference on Computer-Aided Design*, 2025. [Online]. Available: https://arxiv.org/abs/2508.01768v1

[21] R. Ding, T. Xu, X. Shen, A. A. Ding, and Y. Fei, "MoEcho: Exploiting side-channel attacks to compromise user privacy in mixture-of-experts LLMs," in *ACM SIGSAC Conference on Computer and Communications Security*, 2025, pp. 2159–2173. [Online]. Available: https://arxiv.org/abs/2508.15036v1

[22] Z. Gao, J. Hu, F. Guo, Y. Zhang, Y. Han, S. Liu, H. Li, and Z. Lv, "I know what you said: Unveiling hardware cache side-channels in local large language model inference," in *34th USENIX Security Symposium*, 2025, pp. 1649–1668. [Online]. Available: https://www.usenix.org/conference/usenixsecurity25/presentation/gao-zibo

[23] C. Gu, R. Levine, Z. Zhang, T. Sorensen, and Y. Guo, "Behind bars: A side-channel attack on NVIDIA MIG cache partitioning using memory barriers," in *35th USENIX Security Symposium*. USENIX Association, 2026, pp. 5373–5392. [Online]. Available: https://www.usenix.org/conference/usenixsecurity26/presentation/gu-cheng

[24] R. Nazaraliyev, Y. Zhang, and N. Abu-Ghazaleh, "Page404: Breaking GPU isolation through hardware-assisted page migration," *arXiv preprint arXiv:2610.20962v1*, 2026, consulted 2610.20962v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.20962v1

[25] A. Adiletta and B. Sunar, "Spill the beans: Exploiting CPU cache side-channels to leak tokens from large language models," *arXiv preprint arXiv:2505.00817*, 2025. [Online]. Available: https://arxiv.org/abs/2505.00817

[26] H. Weng, R. Zhou, Y. Qu, X. Ji, and W. Xu, "ArchSnoop: LLM architecture snooping via electromagnetic side-channel on edge devices," in *Proceedings of the 19th ACM Conference on Security and Privacy in Wireless and Mobile Networks*, 2026, pp. 51–62. [Online]. Available: https://dl.acm.org/doi/10.1145/3765613.3811673

[27] F. Chen, L. Du, J. Zhou, P. Li, and Z. Su, "Sparleak: Privacy leakage from sparse attention in LLM inference on shared gpus," *arXiv preprint arXiv:2609.38830v1*, 2026, consulted 2609.38830v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.38830v1

[28] E. Hadad and M. Guri, "Shape and substance: Dual-layer side-channel attacks on local vision-language models," *arXiv preprint arXiv:2603.25403v1*, 2026, consulted 2603.25403v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2603.25403v1

[29] T. Hornetz, H. Yavarzadeh, A. Cheu, A. Gascon, L. Gerlach, D. Moghimi, P. Schoppmann, M. Schwarz, and R. Zhang, "Tdxray: Microarchitectural side-channel analysis of Intel TDX for real-world workloads," in *IEEE Symposium on Security and Privacy*, 2026. [Online]. Available: https://tdxray.cpusec.org/assets/tdxray_sp26.pdf

[30] Y. Jo, J. Park, E. Lee, and D. Song, "Sparseety: Extracting tokens from sparsity-exploiting LLM serving systems via deterministic side channels," *arXiv preprint arXiv:2608.02995v1*, 2026, consulted 2608.02995v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2608.02995v1

[31] R. Weiss, B. Konstantinov, E. Sheetrit, T. Simon, and Y. Mirsky, "Detokenization leaks: Reconstructing local LLM outputs from cache traces," *arXiv preprint arXiv:2609.06674v1*, 2026. [Online]. Available: https://arxiv.org/abs/2609.06674v1

[32] J.-J. Jiao, H. Yang, and R. Wen, "Llm-wfin: A fine-grained large language model (LLM)-oriented website fingerprinting attack via fusing interrupt trace and network traffic," *Electronics*, vol. 14, no. 7, p. 1263, 2025, consulted publisher version 2025-03-23; source-scoped access 21 October 2026. [Online]. Available: https://doi.org/10.3390/electronics14071263

[33] D. Scheinert, A. Acker, T. Wittkopp, S. Becker, H. Yous, K. Reddy, I. Farhat, H. Hacid, and O. Kao, "Beyond microservices: Testing web-scale RCA methods on GPU-driven LLM workloads," 2026, version 1, 2 March 2026. [Online]. Available: https://arxiv.org/html/2603.02057v1

[34] G. Politis and E. Pappas, "Privacy failure in split-LLM training, the returned gradient nullifies the decoys," *arXiv preprint arXiv:2609.04382v1*, 2026, consulted 2609.04382v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.04382v1

[35] T. C. Addagada, "Governing the KV cache: Preventing timing side-channel leakage in multi-tenant LLM inference," *arXiv preprint arXiv:2608.09225*, 2026, consulted arXiv:2608.09225v1; archived access 19 October 2026. [Online]. Available: https://arxiv.org/abs/2608.09225v1

[36] S. Alhazbi, A. Hussain, G. Oligeri, and P. Papadimitratos, "LLMs have rhythm: Fingerprinting large language models using inter-token times and network traffic analysis," *IEEE Open Journal of the Communications Society*, vol. 6, pp. 5050–5071, 2025. [Online]. Available: https://arxiv.org/abs/2502.20589v1

[37] N. Carlini and M. Nasr, "Remote timing attacks on efficient language model inference," *arXiv preprint arXiv:2410.17175v1*, 2024. [Online]. Available: https://arxiv.org/abs/2410.17175v1

[38] Z. Zhang, G. Wu, S. Deng, S. Wang, and Y. Zhang, "Netecho: From real-world streaming side-channels to full LLM conversation recovery," *arXiv preprint arXiv:2510.25472*, 2025, consulted arXiv:2510.25472v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2510.25472v1

[39] S. Li, T. Cui, M. Chen, X. Lin, Z. Gu, X. Deng, K. Xu, and Q. Li, "From length to content: Token-length side-channel attacks on LLM api merged outputs," in *35th USENIX Security Symposium*, 2026, pp. 2187–2206. [Online]. Available: https://www.usenix.org/conference/usenixsecurity26/presentation/li-sijia

[40] J. Wei, A. Abdulrazzag, T. Zhang, A. Muursepp, and G. Saileshwar, "When speculation spills secrets: Side channels via speculative decoding in llms," *arXiv preprint arXiv:2411.01076*, 2024, consulted arXiv:2411.01076v4; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2411.01076v4

[41] G. McDonald and J. B. Or, "Whisper leak: a side-channel attack on large language models," *arXiv preprint arXiv:2511.03675*, 2025, consulted arXiv:2511.03675v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2511.03675v1

[42] R. A. Bakar, "Characterizing contention-induced reliability collapse in KV-cache timing side channels for multi-tenant LLM serving," *arXiv preprint arXiv:2609.06853v1*, 2026. [Online]. Available: https://arxiv.org/abs/2609.06853v1

[43] T.-R. Chen, "Sensitive-topic leakage through LLM routing metadata: Measurement and mitigation," *arXiv preprint arXiv:2610.09981v1*, 2026, consulted 2610.09981v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.09981v1

[44] E. Debenedetti, G. Severi, N. Carlini, C. A. Choquette-Choo, M. Jagielski, M. Nasr, E. Wallace, and F. Tramèr, "Privacy side channels in machine learning systems," in *33rd USENIX Security Symposium*. USENIX Association, 2024. [Online]. Available: https://www.usenix.org/conference/usenixsecurity24/presentation/debenedetti

[45] C. Gu, X. L. Li, R. Kuditipudi, P. Liang, and T. Hashimoto, "Auditing prompt caching in language model apis," in *International Conference on Machine Learning*, vol. 267, 2025, pp. 20 477–20 496. [Online]. Available: https://proceedings.mlr.press/v267/gu25b.html

[46] X. Zheng, H. Han, S. Shi, Q. Fang, Z. Du, X. Hu, and Q. Guo, "Inputsnatch: Stealing input in LLM services via cache-sharing timing side-channel attacks," in *Advanced Parallel Processing Technologies (APPT 2026)*, 2027, pp. 483–497, first online August 24, 2026; publisher citation year 2027. [Online]. Available: https://link.springer.com/chapter/10.1007/978-981-92-4805-6_32

[47] S. Majidi, N. Mireshghallah, and K. Taram, "Leaky language models: Stealing architecture and inference optimizations via per-token timing," *arXiv preprint arXiv:2607.20723*, 2026, consulted arXiv:2607.20723v1; archived access 19 October 2026. [Online]. Available: https://arxiv.org/abs/2607.20723v1

[48] H. Sun, S. Liu, S. Ma, J. Li, M. Xiao, and W. Jiang, "Agent-assisted side-channel attacks on non-prefix KV cache in rag," *arXiv preprint arXiv:2606.21842*, 2026, consulted arXiv:2606.21842v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2606.21842v1

[49] T. Zhang, G. Saileshwar, and D. Lie, "Time will tell: Timing side channels via output token count in large language models," *arXiv preprint arXiv:2412.15431*, 2024, consulted arXiv:2412.15431v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2412.15431v1

[50] G. Wu, Z. Li, Y. Zhang, Z. Zhang, J. Niu, Y. Wu, and Y. Zhang, "Cacheprune: Privacy-aware and fine-grained KV cache sharing for efficient LLM inference," *arXiv preprint arXiv:2605.23640v1*, 2026, consulted 2605.23640v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2605.23640v1

[51] I. Yona, I. Shumailov, J. Hayes, and N. Carlini, "Stealing user prompts from mixture of experts," *arXiv preprint arXiv:2410.22884v1*, 2024. [Online]. Available: https://arxiv.org/abs/2410.22884v1

[52] P. Revankar, K. Chauhan, J. Kim, S. N. Nur, V. Siu, and C.-G. Wang, "Mirage: A polarity-flipping encoding subspace in LLM agents," *arXiv preprint arXiv:2606.10304v1*, 2026, consulted 2606.10304v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.10304v1

[53] S. Pouryousef, J. Lopez, S. R. Nowmi, M. M. Kamol, M. Hossain, M. Tran, and M. S. Rahman, "A large-scale benchmark and risk assessment of traffic analysis attacks on cloud LLM services," *arXiv preprint arXiv:2609.31877v1*, 2026, consulted arXiv:2609.31877v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2609.31877v1

[54] I. Fayolle, S. Bouhenniche, S. Pélissier, P. Laperdrix, C. Maurice, and W. Rudametkin, "On the internet, nobody knows you're an LLM bot: Unmasking web agents with multi-layer fingerprinting," *arXiv preprint arXiv:2606.30119v1*, 2026, consulted 2606.30119v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.30119v1

[55] R. Solomon, Y. Y. Levi, L. Vaknin, E. Aizikovich *et al.*, "LumiMAS: A comprehensive framework for real-time monitoring and enhanced observability in multi-agent systems," *arXiv preprint arXiv:2508.12412v2*, 2026. [Online]. Available: https://arxiv.org/abs/2508.12412v2

[56] F. George, H. Kumar, D. Pathak, K. Ray, M. Verma, and P. Moogi, "Unsupervised cycle detection in agentic applications," *arXiv preprint arXiv:2511.10650v1*, 2025. [Online]. Available: https://arxiv.org/abs/2511.10650v1

[57] J. Geng, R. He, Z. Fei, B. Yi, R. Wang, Z. Liu, X. Hu, X. Wu, and Q. Zeng, "Agent skills matter: Inferring proprietary skills from execution trajectories," *arXiv preprint arXiv:2607.25560v1*, 2026, consulted arXiv:2607.25560v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2607.25560v1

[58] H. Oderinwale, "Agent trajectories as programs: fingerprinting and programming coding-agent behavior," *arXiv preprint arXiv:2606.16988v1*, 2026. [Online]. Available: https://arxiv.org/abs/2606.16988v1

[59] X. Wang, B. Zhao, M. Backes, F. Boenisch, and A. Dziedzic, "Agentprov: Auditing agentic LLM api providers via tool-use policy probes," *arXiv preprint arXiv:2609.00052v1*, 2026, consulted 2609.00052v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.00052v1

[60] C.-Y. Wang, X.-H. Xie, T.-Z. Wang, F.-C. Luo, and Y. Cui, "Who is behind the harness? fingerprinting LLMs through agentic behavior," *arXiv preprint arXiv:2609.28559v1*, 2026, consulted 2609.28559v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.28559v1

[61] S.-J. Park, M. Choi, J. Kim, and M.-S. Kang, "Cross-agent campaign attribution: Linking asynchronous attacks across LLM agents," *arXiv preprint arXiv:2607.18826v1*, 2026, consulted 2607.18826v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2607.18826v1

[62] C. Song, H. Wang, and R. Martin, "Personafingerprint: Measuring persona inference on modern websites with LLM-driven browsing," *arXiv preprint arXiv:2605.15962*, 2026. [Online]. Available: https://arxiv.org/abs/2605.15962

[63] D. Kang, H. Jeong, J. Sheffey, P. Datta, and A. Houmansadr, "Whose agent are you? multi-layer fingerprinting and attribution of autonomous web agents," *arXiv preprint arXiv:2606.20910v1*, 2026. [Online]. Available: https://arxiv.org/abs/2606.20910v1

[64] A. Shahriar, M. N. Rahman, S. Ahmed, F. Sadeque, and M. R. Parvez, "AgentTell: Behavioural side-channel leakage in browser-use agents," *arXiv preprint arXiv:2609.32915v1*, 2026, consulted arXiv:2609.32915v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2609.32915v1

[65] E. Wang, Z. Shafiq, and Y. Vekaria, "Fp-agent: Fingerprinting AI browsing agents," *arXiv preprint arXiv:2605.01247*, 2026, consulted arXiv:2605.01247v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2605.01247v1

[66] S. Banerjee, P. Sahu, M. Luo, A. Vahldiek-Oberwagner, N. J. Yadwadkar, and M. Tiwari, "SoK: A systems perspective on compound AI threats and countermeasures," *arXiv preprint arXiv:2411.13459v1*, 2024, consulted arXiv:2411.13459v1; inspected 20 October 2026. [Online]. Available: https://arxiv.org/abs/2411.13459v1

[67] J. Forough, M. Kogias, and H. Haddadi, "When agents handle secrets: A survey of confidential computing for agentic AI," *arXiv preprint arXiv:2605.03213v2*, 2026, consulted arXiv:2605.03213v2; inspected 20 October 2026. [Online]. Available: https://arxiv.org/abs/2605.03213v2

[68] O. Kotevska, S. Jha, A. Bellet, R. Hu, N. D. Bastian, R. F. da Silva, R. Madduri, and K. Kim, "Privacy foundations for multi-institutional scientific artificial intelligence," *arXiv preprint arXiv:2609.39787v1*, 2026, consulted 2609.39787v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.39787v1

[69] X. Cai, R. Nithyanand, T. Wang, R. Johnson, and I. Goldberg, "A systematic approach to developing and evaluating website fingerprinting defenses," in *ACM CCS*, 2014, pp. 227–238. [Online]. Available: https://doi.org/10.1145/2660267.2660362

[70] M. Juarez, M. Imani, M. Perry, C. Diaz, and M. Wright, "Toward an efficient website fingerprinting defense," in *ESORICS*, 2016, pp. 27–46. [Online]. Available: https://arxiv.org/abs/1512.00524

[71] T. Wang and I. Goldberg, "Walkie-talkie: An efficient defense against passive website fingerprinting attacks," in *26th USENIX Security Symposium*, 2017. [Online]. Available: https://www.usenix.org/conference/usenixsecurity17/technical-sessions/presentation/wang-tao

[72] N. Nazari, F. Xiang, C. Fang, H. M. Makrani, A. Puri, K. Patwari, H. Sayadi, S. Rafatirad, C.-N. Chuah, and H. Homayoun, "LLM-FIN: Large language models fingerprinting attack on edge devices," in *International Symposium on Quality Electronic Design*, 2024. [Online]. Available: https://web.cs.ucdavis.edu/~srafatir/pages/documents/llm_fingerprint.pdf

[73] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed. Wiley-Interscience, 2006, theorem 2.8.1, Data-processing inequality. [Online]. Available: https://onlinelibrary.wiley.com/doi/book/10.1002/047174882X

[74] E. Rescorla, K. Oku, N. Sullivan, and C. A. Wood, "TLS Encrypted Client Hello," RFC 9849, Mar. 2026. [Online]. Available: https://www.rfc-editor.org/rfc/rfc9849

[75] Kubernetes Authors, "Kubernetes API Concepts: Dry-run," https://kubernetes.io/docs/reference/using-api/api-concepts/#dry-run, accessed October 24, 2026.

[76] Stripe, "Sandboxes," https://docs.stripe.com/sandboxes, accessed October 24, 2026.

[77] PostgreSQL Global Development Group, "PostgreSQL 18 Documentation: ROLLBACK," https://www.postgresql.org/docs/18/sql-rollback.html, accessed October 24, 2026.

[78] ——, "PostgreSQL 18 Documentation: Sequence Manipulation Functions," https://www.postgresql.org/docs/18/functions-sequence.html, accessed October 24, 2026.

[79] R. H. Anwar, J. O. Weiss, T. Alves, S. Kundu, and M. T. Raza, "InferNet: Exploiting aggregate GPU profiles as side-channel for DNN architecture inference," *ACM Transactions on AI Security and Privacy*, 2026. [Online]. Available: https://rhasnainanwar.me/assets/pdf/2304.03388v2.pdf

[80] W. Cui, H. Zhang, Y. Huang, Q. Li *et al.*, "SymDiag: Explainable diagnosis for LLM reasoning via neuro-symbolic verification," in *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, 2026, consulted arXiv:2608.08786v1; archived access 18 October 2026. [Online]. Available: https://arxiv.org/abs/2608.08786v1

[81] N. Lahjouji and A. G. Colaco, "Agents that know too much: A data-centric survey of privacy in LLM agents," *arXiv preprint arXiv:2606.26627*, 2026, consulted arXiv:2606.26627v1; archived access 19 October 2026. [Online]. Available: https://arxiv.org/abs/2606.26627v1

[82] J. Roh, E. Bagdasarian, H. Haddadi, and A. S. Shamsabadi, "SPILLage: Agentic oversharing on the web," *arXiv preprint arXiv:2602.13516v1*, 2026. [Online]. Available: https://arxiv.org/abs/2602.13516v1

[83] G. Yu, Q. Wang, R. Lang, S. Su, and X. Wang, "PlanTwin: Privacy-preserving planning abstractions for cloud-assisted LLM agents," *arXiv preprint arXiv:2603.18377v2*, 2026. [Online]. Available: https://arxiv.org/abs/2603.18377v2

[84] J. Xu, L. Fan, Z. Wang, X. Li, and H. Liu, "Selection-invariant communication compilers for privacy-aware multi-agent LLM workflows," *arXiv preprint arXiv:2609.26076v1*, 2026. [Online]. Available: https://arxiv.org/abs/2609.26076v1

[85] Y. Qiao, D. Liu, H. Yang, W. Zhou, and S. Hu, "Agent tools orchestration leaks more: Dataset, benchmark, and mitigation," *arXiv preprint arXiv:2512.16310v4*, 2026, version 4. [Online]. Available: https://arxiv.org/abs/2512.16310v4

[86] M. Umar, A. P. Marathe, M. D. Gupta, S. J. Ghosh, G. E. Suh, and W. Xiong, "Efficient memory side-channel protection for embedding generation in machine learning," in *IEEE International Symposium on High-Performance Computer Architecture*, 2025. [Online]. Available: https://bearhw.ece.vt.edu/content/dam/bearhw_ece_vt_edu/publications/2025_HPCA_SecEmb.pdf

[87] L. Wu, C. Wang, T. Liu, Y. Zhao, and H. Wang, "From assistants to adversaries: Exploring the security risks of mobile LLM agents," *arXiv preprint arXiv:2505.12981v1*, 2025, version 1. [Online]. Available: https://arxiv.org/abs/2505.12981

[88] J. W. Kim, B. Min, K. Edemacu, Y. Choi, S.-H. Cho, and B. Jang, "Dp-memview: A memory interface for attribute-level transcript privacy in long-term LLM agents," *arXiv preprint arXiv:2608.03130v1*, 2026, consulted 2608.03130v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2608.03130v1

[89] M. Mudryi, M. Chaklosh, and G. M. Wójcik, "The hidden dangers of browsing AI agents," *arXiv preprint arXiv:2505.13076v1*, 2025, version 1. [Online]. Available: https://arxiv.org/abs/2505.13076

[90] Y.-M. Wang, Y.-X. Li, S.-L. Fan, Y.-H. Zheng, and T.-H. Chang, "Ask without telling: Local slms consult cloud LLMs without revealing task intent," *arXiv preprint arXiv:2609.32642v1*, 2026, consulted 2609.32642v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2609.32642v1

[91] A. Ukani, H. Haddadi, A. S. Shamsabadi, and P. Snyder, "Privacy practices of browser agents," *arXiv preprint arXiv:2512.07725*, 2025. [Online]. Available: https://arxiv.org/abs/2512.07725

[92] Z. Liu, S. van Hoek, P. Horváth, D. Lauret, X. Xu, and L. Batina, "Real-world edge neural network implementations leak private interactions through physical side channel," *arXiv preprint arXiv:2501.14512v1*, 2025, consulted 2501.14512v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2501.14512v1

[93] P. G. Pennas, K. Papaioannou, M. Guarnieri, and T. D. Doudali, "Cachesolidarity: Preventing prefix caching side channels in multi-tenant LLM serving systems," *arXiv preprint arXiv:2603.10726v1*, 2026, consulted 2603.10726v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2603.10726v1

[94] C.-Y. Zhou, Q.-L. Jiang, S.-N. Wu, and X.-H. Zhou, "The vision encoder as a privacy boundary: Visual-token side channels in encoder-free vision-language models," *arXiv preprint arXiv:2606.14783v1*, 2026, consulted 2606.14783v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2606.14783v1

[95] G. Politis and E. Pappas, "What gradients add to text leakage in split language models, counted per token and per document," *arXiv preprint arXiv:2610.04128v1*, 2026, consulted 2610.04128v1; source-scoped access 21 October 2026. [Online]. Available: https://arxiv.org/html/2610.04128v1

Appendix A. Audit Trail

The artifact preserves the 82-family map, 39-record ledger, 180 family cells and 351 record cells with source versions and locators. S33 publishes query membership and sensitivity checks. Case-study protocols, observations and uncertainty calculations retain their original cohorts and provenance.

Appendix B. Result-record index

Table 12 indexes all 39 records by measurement mode and source locator. Marks distinguish six contextual controls (‡) and six final-search additions (†); all enter the pair inventory. The ledger retains full locators, versions and hashes.

Appendix C. Case-Study Artifact Index

Table 13 maps the single appendix case to historical artifact keys and supplementary methods. These keys preserve original task/customer memberships; S30 retains the corrected permutation/Holm tests and calibration resampling.

Appendix D. Family-Level Holdout Evidence

Table 14 covers 82 families and 492 cells. Axes: T task/input/user; M model; C harness/topology; H hardware; N network/load; D time. Codes: H reported holdout; V condition comparison/recalibration; X scoped restriction; ? unresolved; – conceptual/inapplicable. Scope C/A denotes core/adjacent. H records separation, not transfer success; ? does not establish an absent experiment. S28 retains source locators and qualified interpretations.

TABLE 12. COMPLETE INDEX OF THE 39 RESULT RECORDS (26 FAMILIES). MODE IS THE PRIMARY MEASUREMENT MODE: EX EXECUTED COLLECTION/PROTOTYPE (20), TR TRACE TRANSFORMATION (6), SI SIMULATION/DERIVED ESTIMATE (5), CO COMPONENT/CALIBRATION MEASUREMENT (6), UN PLACEMENT UNRESOLVED (2). ‡: CONTEXTUAL RECORD (6); †: ADDED BY THE FINAL SEARCH (6). T=TABLE, F=FIGURE, APP.=APPENDIX; LOCATORS REFER TO THE CITED SOURCE.

Record ID / family	Mode	Boundary; source locator
UI-delay-offline [5]	Tr	Offline delay; no live task outcome. 6.1 F5; A.5
WRA-decoy-execution [6]	Ex	Executed decoys; usefulness proxy only. 8.1–8.3 T8–11
NET-upstream-local [38]	Ex	Protected deployment/local serving; topology scoped. IV-C–E TIII–IV
NET-dummy-padding-simulation [38]	Tr	Synthetic insert/pad; no live execution. VII TIX
EARLY-K-calibration-oracle [3]	Si	K-specific fit then oracle; not live KV trial. V-A TVI; IV-A
EARLY-semantic-deidentification [3]	Ex	GPTCache prototype cost; not KV-granularity cost. V-B
WHISPER-randompadding-capture [41]	Ex	New padded captures; rounded medians. 6.1 T3–4
WHISPER-batching-replay [41]	Tr	Offline batching; model-dependent result. 6.2 F7
WHISPER-injection-replay [41]	Tr	Synthetic packets; no live zero-delay claim. 6.3 T5
SPLICE-CTBF-signal [48]	Ex	Prototype background load; tested signal only. VII-C TII
SPLICE-QCP-cost [48]	Co	QCP component cost; excludes combined CTBF. VII-C TII
TIME-length-instruction [49]	Ex	Length instruction; unequal outputs remain. 7.2
TIME-padding-estimate [49]	Si	Analytic plaintext cost; not network bytes. 7.3
SPEC-padding-REST [40]	Un	Padding/aggregation; payload cost, placement unverified. 6.1 F6–7 T2
DANGOL-capture [16]	Tr	Capture transform; distinct from live population. 9.8–9.9 T6
DANGOL-livebinding [16]	Ex	Live loopback binding; do not join capture accuracy. 9.7, OHTTP paragraph after T5
SIG-common-guardrail‡ [57]	Ex	Common instruction; not a causal guardrail effect. 4.1–4.2; App. C
SAFE-post-trigger [17]	Ex	Post-trigger probes; detector cost not service TTFT. 7.2.4 F12; 7.3–7.4
AHB-online-executed [14]	Ex	Executed UI edits; task/state changes possible. 5.1.2 T1; E.1–E.3
AHB-offline [14]	Tr	Offline touch edits; no online task damage claim. T1 offline rows; E.1–E.3
KVGOV-partialflush-sim [35]	Si	Deterministic simulation; not live protected recovery. 5.1–5.3
KVGOV-hardware-gap [35]	Co	Hardware gap calibration; not protected ASR. 5.7–5.8
KVGOV-HMAC-cost [35]	Co	HMAC microcost; not cache reuse/task loss. 5.9
TRAFFIC-prompt-reformulation [53]	Ex	Collected prompt rewrites; task utility unmeasured. V-C1 F9a
TRAFFIC-temperature [53]	Ex	Collected decoding changes; no monotone benefit. V-C1 F9a
CARLINI-fixed-rate [37]	Un	Rate/replayed cost; per-token, not request delay. VII F11–12
GAO-zero-copy-off [22]	Co	Loader component costs; not end-to-end TTFT. 7.1
TDXRAY-oblivious-map [29]	Co	Lookup microbenchmark; not full inference cost. 6.1 F8 T3
AGENTTELL-fixed-instruction‡ [64]	Ex	Live local browser tasks; loaded/cold context contrast. 4.3.1; 5 Finding 1
FP-CLOUDFLARE‡ [65]	Ex	Live access control; not user confidentiality. 6.4 T5
BAR-co-running-noise‡ [23]	Ex	MIG noise test; not an LLM defense trial. 7.2–7.3 T8
RHYTHM-VPN‡ [36]	Ex	Live VPN shift; baseline-trained attacker only. VI-B; VII
DEB-filter-origin‡ [44]	Ex	Filter creates attack view; not mitigation of it. Threat model; 4.2
CACHEPRUNE-selective† [50]	Ex	Selective cache sharing; direct token recovery/quality. 3.1; 5.1–5.6; T1; App. A
CACHEPRUNE-context† [50]	Ex	Contextual guessing; distinct from direct recovery. 3.1; 5.1–5.6; T1; App. A
OTRO-unpadded† [18]	Co	ORAM TTFT/memory; access-count padding excluded. VII-C–E; TIII
OTRO-epochpad† [18]	Si	Epoch-padding estimate; not measured padded TTFT. VII-C–E; TIII
SPARLEAK-replacement† [27]	Ex	Sparse replacement/quality; attack cost is separate. 3.1; 4–7; T4
SHAPE-padding-estimate† [28]	Si	Constant-work proposal/estimate; no executed pair. IX-B1

TABLE 13. ARTIFACT INDEX FOR THE APPENDIX CASE STUDY. COMPONENT NAMES DESCRIBE THE ARCHIVED MEASUREMENT BOUNDARY; THE RIGHT COLUMN LOCATES IMPLEMENTATION AND FULL-GRID DETAIL.

Case-study component	Artifact keys	Observation/execution boundary and comparison	Detailed source
Transport emission	E1, E3, E5	Buffered two-link re-emission and token-stream emitters; upstream framing, final-link framing and matched attacker refitting are distinct comparisons.	S9/S11; S12
State and link	E2, E6, E7	Prescribed stateful plans and randomized dispatch cover, then agent–tool TLS composition; dispatch-visible calls and encrypted-link records are separate views.	S9/S11; S12
Trajectories and localization	E9, E10, E12	Saved autonomous Sol trajectories; original pre-test calibration, exploratory customer cross-fitting and oracle operation-category localization retain their different fit/test boundaries.	S14; S15/S21
Timing and WF controls	E13, E14	Local live tool-loop timing/full tunnels, then stronger learners and specified WF adaptations; measured clocks and simulations are distinguished, with matched suites for joint comparisons.	S19; S20
Extensions	E8	Stand-in-planner pipeline check.	S13
	E11	Retrospective training-fold operating-point selection with the historical suite.	S15
	E15	Additional same-provider Luna cohort.	S22
	E16	Same-operation semantic-attribute diagnostic on reused Sol data.	S23.6
	R11/R12	Two served model endpoints: native trajectories, frozen training calibration and TLS write localization.	S24
	R14	Attacker/calibration transfer matrices, non-oracle write detection and write-only framing with the endpoint suite.	S25
	R15	Fresh paired live-planner episodes under isolated prefix cover and injected tool-response deadlines.	S26

TABLE 14. COMPLETE HOLDOUT-CODE MATRIX; CODES AND AXES ARE DEFINED IN APPENDIX D. SOURCE-QUALIFIED ROW NOTES ARE IN SUPPLEMENT S28.

Family	Scope	T	M	C	H	N	D	Family	Scope	T	M	C	H	N	D
F01 AgentPrint [4]	C	?	?	?	?	?	?	F42 ML privacy channels [44]	C	?	V	?	?	?	?
F02 UI Traces [5]	C	V	V	V	?	?	?	F43 MoE stealing [51]	C	?	?	?	?	?	?
F03 WRA [6]	C	H	?	?	?	?	?	F44 Energon [20]	C	?	V	?	?	?	?
F04 FP-Agent [65]	C	H	?	X	?	?	?	F45 SafeKV [17]	C	?	V	?	?	?	?
F05 Prompt length [1]	C	H	?	V	?	V	V	F46 InferNet [79]	A	?	V	?	V	?	?
F06 PromptPeek [2]	C	V	?	?	?	V	?	F47 ArchSnoop [26]	C	?	H	?	V	?	?
F07 NetEcho [38]	C	V	?	V	?	?	?	F48 Coding fingerprints [9]	A	X	?	X	?	?	X
F08 Early Bird [3]	C	?	V	?	?	?	?	F49 AHB [14]	C	?	?	X	?	?	?
F09 InputSnatch [46]	C	H	?	?	?	?	?	F50 WildAuth [10]	A	H	?	X	?	?	?
F10 Whisper Leak [41]	C	H	V	?	?	?	?	F51 Cycles [56]	C	X	?	X	?	?	X
F11 Spill the Beans [25]	C	V	?	?	?	?	?	F52 LumiMAS [55]	C	?	V	V	?	?	?
F12 SpliceLeak [48]	C	V	?	?	?	?	?	F53 SymDiag [80]	A	V	V	?	?	?	?
F13 Time Will Tell [49]	C	?	?	?	?	V	V	F54 LeakyLMs [47]	C	?	H	?	V	?	?
F14 Prompt Echo [39]	C	H	V	?	?	?	?	F55 KVGov [35]	C	?	?	?	V	?	?
F15 I Know What You Said [22]	C	H	V	V	V	?	?	F56 Data-centric survey [81]	A	-	-	-	-	-	-
F16 Speculation [40]	C	H	?	?	?	V	?	F57 Serving RCA [33]	C	?	?	X	?	?	?
F17 Behind Bars [23]	C	?	V	?	V	?	?	F58 ZeroPrint [11]	A	?	V	?	?	?	?
F18 Prompt caching [45]	C	?	V	?	?	?	V	F59 Endpoint fingerprints [12]	A	?	?	V	?	?	V
F19 Whose Agent [63]	C	?	?	?	?	?	?	F60 StepFinder [13]	A	H	?	?	?	?	?
F20 Workflow traffic [16]	C	?	?	H	?	?	?	F61 Web-agent detection [54]	C	?	V	?	?	?	?
F21 GPU contention [42]	C	?	?	V	?	V	?	F62 Cross-campaign [61]	C	V	?	V	?	?	?
F22 AgentTell [64]	C	V	?	?	?	V	?	F63 AgentProv [59]	C	?	V	?	?	?	?
F23 Traffic benchmark [53]	C	?	?	H	?	?	?	F64 LiDAR [60]	C	?	V	V	?	?	?
F24 Remote timing [37]	C	?	V	?	?	?	V	F65 LLM-WFIN [32]	C	V	?	?	?	?	?
F25 Persona fingerprint [62]	C	?	?	X	?	?	?	F66 MIRAGE [52]	C	V	V	?	?	?	?
F26 TDXRay [29]	C	?	?	?	?	?	?	F67 Routing metadata [43]	C	V	?	V	?	?	?
F27 Detokenization [31]	C	?	?	X	?	?	?	F68 Shape and Substance [28]	C	V	V	V	V	?	?
F28 Spillage [82]	A	V	?	?	?	?	?	F69 SplitDecoy [34]	C	?	?	?	?	?	?
F29 PlanTwin [83]	A	V	?	?	?	?	?	F70 PAGE404 [24]	C	V	V	?	?	?	?
F30 SICC [84]	A	?	?	?	?	?	?	F71 Edge EM [19]	C	H	?	?	?	?	?
F31 TOPR [85]	A	V	?	?	?	?	?	F72 SparLeak [27]	C	V	V	V	?	?	?
F32 ClawGuard [15]	C	?	?	V	?	?	V	F73 Sparseety [30]	C	V	V	V	?	?	?
F33 SIGLEAK [57]	C	V	?	?	?	?	?	F74 OTRO [18]	C	?	?	?	?	?	?
F34 ProcGrep [58]	C	H	?	V	?	?	?	F75 CachePrune [50]	C	V	V	?	?	?	?
F35 SecEmb [86]	A	?	X	?	?	?	?	F76 Privacy foundations [68]	A	-	-	-	-	-	-
F36 Mobile agents [87]	A	?	?	?	?	?	?	F77 DP-MemView [88]	A	?	?	?	?	?	?
F37 Hidden risks [89]	A	?	?	?	?	?	?	F78 PriCon [90]	A	?	?	?	?	?	?
F38 Browser privacy [91]	A	?	?	V	?	?	?	F79 SCAAR [92]	A	?	?	?	?	?	?
F39 LLMs Have Rhythm [36]	C	?	V	?	?	H	V	F80 CacheSolidarity [93]	A	?	?	?	?	?	?
F40 MoEcho [21]	C	H	V	?	?	V	?	F81 Vision encoder [94]	A	?	?	?	?	?	?
F41 LLM-Fin [72]	A	?	H	?	?	?	?	F82 Split gradients [95]	A	?	?	?	?	?	?

Appendix E.
Case Study: Applying the Evidence-Reuse Audit

This worked application follows the three SoK operations: compare records at a specified receiver, join configuration-linked outcomes, then identify the measurement required by a deployment change. Its four contrasts form one case study. Historical study keys and complete protocols remain in Appendix C and the supplement.

Setting and measurement scope. The transport component uses 48 saved responses on 24 topics (12 training/12 test). The agent component uses 115 public retail tasks from 53 customers: GPT-6.1 Sol/low supplies 690 repeated episodes, and served Luna, MiniMax and GLM aliases supply separately scoped extensions. The frozen retail split has 69 training/46 test tasks and 29/24 disjoint customers; test cancellation/return positives number 5/7 customers. The operation-category target identifies the requested cancellation/return action. It differs from the saved responses' verbosity target. TLS loopback captures, offline transforms, emulated WAN paths and fresh planner runs retain distinct execution/cost objects.

Follow the binding rather than the mechanism name. The first contrast binds cost to the final emitter. The second distinguishes unchanged observed tool categories from state preservation. For deterministic real tools, forwarding the same real actions, executing cover on disjoint state and suppressing cover replies preserves logical real-call responses and final state; this conformance property excludes wall-clock time. It does not measure the utility of a deadline-sensitive planner. The separate 100-episode prefix-cover study observes 37/50 baseline and 33/50 defended successes (difference -8.0 points, customer interval $[-23.9, 6.5]$), at that scheduler's scope.

Interpretation and reconstruction. The application specifies the evidence needed to extend each workload-specific result. Reused customers, mutable aliases, unequal inference budgets and exploratory analyses define its scope. The original screen (upper 95% BA < 0.60 and AUROC < 0.65 for both targets) remains unmet; relative attack reductions do not establish that absolute criterion. S11–S26 retain full protocols and results, S30 retains all corrected tests and sensitivity intervals, and S31–S32 collect the detailed grids. The SoK findings derive from the literature audit.

Recorded contrasts for the worked case. BA is maximum balanced accuracy over each named fitted suite; labels and populations are row-specific.

Boundary	Recorded contrast and scope
Emitter	Buffered cap: 50.0% verbosity BA at 2.19× bytes; eight-token streaming batches: 100.0% defended-refit BA at 0.21×. Same 48 responses, TLS 1.3, 12 test topics.
State	A shadow/live cancellation decoy gives the same refitted 100.0/100.0% cancel/return BA, but database equality is 100%/0%. All 115 reference programs, three seeds, task-folded suite.
Calibration	MiniMax-to-GLM swap: $+20.8/-17.3$ cancel/return BA points; bytes 15.95× → 18.65×. Local return BA: 50.0–67.9% over 100 resamples (95% sensitivity spread). Endpoint suite; 46 test tasks/24 customers.
Observation	Matched serial-shadow captures: the expanded timing/flush union raises cancellation BA 50.0%→61.4% ($+11.4$ points, interval $[3.7, 19.6]$). Standard prompt, 138 test episodes/24 customers.

Local recalibration variation is comparable to the original source swap, and all eight customer-paired contrasts have Holm-adjusted $p = 1.000$. The tests leave source effects unresolved; their nonsignificance establishes neither equivalence nor audit-framework validity. The 100 resamples measure sensitivity within the observed customers, adding no independent customers. Finally, timing changes the fitted attack result on the same captures. Offline WT's 2.04× bytes and measured shadow's 15.11× retain different attack and latency scopes. The latter's measured RPC p95 is 26.2 ms, excluding LLM generation.